

NCAA MBB — Exhaustive Visual-First EDA (Local Parquet)

Working locally (API down). Files are in the same folder as this notebook:

- `games.parquet`
- `teams.parquet`
- `team_stats.parquet`
- `elo.parquet`

Notebook goals:

1. Load + standardize schemas (dirty-data tolerant)
2. Build core game table with targets (home_win, margin)
3. Merge in ELO + team/team_stats when possible
4. Exhaustive, visual-first EDA (distributions, comparisons, drift, correlation, PCA, clustering)

```
In [1]: # =====  
# ALL IMPORTS (single cell)  
# =====  
  
# Standard library  
import os  
import re  
import json  
import math  
import time  
import random  
import warnings  
from pathlib import Path  
from datetime import datetime  
  
# Core DS  
import numpy as np  
import pandas as pd  
  
# Viz  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# Stats  
from scipy import stats  
from scipy.stats import ttest_ind, ks_2samp  
  
# Sklearn (EDA structure)  
from sklearn.preprocessing import StandardScaler, MinMaxScaler  
from sklearn.decomposition import PCA, TruncatedSVD  
from sklearn.manifold import TSNE
```

```

from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from sklearn.metrics import silhouette_score

# Quick separability sanity checks (EDA signal)
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score

warnings.filterwarnings("ignore")
sns.set_theme()
plt.style.use("seaborn-v0_8")

# Parquet engine check (required for pd.read_parquet)
ENGINE = None
try:
    import pyarrow # noqa: F401
    ENGINE = "pyarrow"
except Exception:
    try:
        import fastparquet # noqa: F401
        ENGINE = "fastparquet"
    except Exception as e:
        raise ImportError(
            "No parquet engine found. Install one:\n"
            "  pip install pyarrow    (recommended)\n"
            "or\n"
            "  pip install fastparquet\n"
        ) from e

print("Imports loaded. Parquet engine:", ENGINE)

```

Imports loaded. Parquet engine: pyarrow

0) Load Parquets (Local Folder) + Standardize Column Names

We assume the parquet files are in the same folder as the notebook (`Path.cwd()`).
Then we normalize all column names so downstream logic is stable.

```

In [2]: DATA_PATH = Path.cwd()

required_files = ["games.parquet", "teams.parquet", "team_stats.parquet", "elo.parquet"]
print("CWD:", DATA_PATH)
print("Parquet files in CWD:", [p.name for p in DATA_PATH.glob("*.parquet")])

missing = [f for f in required_files if not (DATA_PATH / f).exists()]
if missing:
    raise FileNotFoundError(f"Missing files in {DATA_PATH}: {missing}")

games = pd.read_parquet(DATA_PATH / "games.parquet", engine=ENGINE)
teams = pd.read_parquet(DATA_PATH / "teams.parquet", engine=ENGINE)
team_stats = pd.read_parquet(DATA_PATH / "team_stats.parquet", engine=ENGINE)
elo = pd.read_parquet(DATA_PATH / "elo.parquet", engine=ENGINE)

def clean_columns(df: pd.DataFrame) -> pd.DataFrame:

```

```

df = df.copy()
df.columns = (
    df.columns.astype(str)
    .str.strip()
    .str.lower()
    .str.replace(" ", "_")
    .str.replace("-", "_")
    .str.replace(r"[^a-z0-9_]", "", regex=True)
)
return df

games = clean_columns(games)
teams = clean_columns(teams)
team_stats = clean_columns(team_stats)
elo = clean_columns(elo)

# Optional season filter. Leave as None for all seasons.
TARGET_SEASONS = None # Example: [2024] or list(range(2018, 2025))

for df_name, df in (("games", games), ("team_stats", team_stats), ("elo", elo)):
    if "season" in df.columns:
        df["season"] = pd.to_numeric(df["season"], errors="coerce").astype("Int64")

if TARGET_SEASONS is not None:
    target = set(TARGET_SEASONS)
    if "season" in games.columns:
        games = games[games["season"].isin(target)].copy()
    if "season" in team_stats.columns:
        team_stats = team_stats[team_stats["season"].isin(target)].copy()
    if "season" in elo.columns:
        elo = elo[elo["season"].isin(target)].copy()

print("games:", games.shape)
print("teams:", teams.shape)
print("team_stats:", team_stats.shape)
print("elo:", elo.shape)

```

CWD: /Users/ceverson/Development/ml1_final_code

Parquet files in CWD: ['games.parquet', 'team_stats.parquet', 'wbb_team_stats.parquet', 'elo.parquet', 'wbb_elos.parquet', 'teams.parquet', 'wbb_games.parquet', 'wbb_teams.parquet']

games: (100003, 18)

teams: (368, 3)

team_stats: (6701, 30)

elo: (6690, 5)

Step 1 — Inspect `games` Table Structure

Before engineering targets or merging anything, we must:

- Identify home team column
- Identify away team column
- Identify home score column

- Identify away score column
- Check for season column
- Check for date column

Because the data is not clean, we will not assume column names. We will inspect schema and sample rows first.

```
In [3]: # =====
# Inspect games table
# =====

print("Shape:", games.shape)

print("\nColumns:")
for col in games.columns:
    print(" -", col)

print("\nDtypes:")
print(games.dtypes)

print("\nHead:")
display(games.head(5))

print("\nNumeric summary (coerce only likely-numeric columns):")

likely_numeric = ["score_vis", "score_home", "gameno", "venue_code", "id"]

games_numeric = games.copy()
for col in likely_numeric:
    if col in games_numeric.columns:
        games_numeric[col] = pd.to_numeric(games_numeric[col], errors="coerce")

# describe only the numeric columns we just coerced
numeric_only_df = games_numeric[[c for c in likely_numeric if c in games_numeric.columns]]

display(numeric_only_df.describe())
```


Shape: (100003, 18)

Columns:

- id
- visitor
- visitor_code
- score_vis
- home
- home_code
- score_home
- gamestatus
- overtime
- winner_code
- loser_code
- gameday
- gameno
- venue
- venue_code
- neutral
- league
- season

Dtypes:

id	object
visitor	object
visitor_code	object
score_vis	object
home	object
home_code	object
score_home	object
gamestatus	object
overtime	object
winner_code	object
loser_code	object
gameday	object
gameno	object
venue	object
venue_code	object
neutral	object
league	object
season	Int64

dtype: object

Head:

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	300941	Memphis Tigers	MEM	68	Kansas Jayhawks	KU	75
1	300939	North Carolina Tar Heels	UNC	66	Kansas Jayhawks	KU	84
2	300940	UCLA Bruins	UCLA	63	Memphis Tigers	MEM	78
3	502394	Bradley Braves	BRAD	64	Tulsa Golden Hurricane	TULS	70
4	502393	Ohio State Buckeyes	OSU	92	Massachusetts Minutemen	MASS	85

Numeric summary (coerce only likely-numeric columns):

	score_vis	score_home	gameno	venue_code	id
count	99812.000000	99812.000000	89441.000000	100003.000000	1.000030e+05
mean	67.478820	73.729572	1.000011	1438.569533	8.166069e+05
std	12.194329	13.230403	0.003344	15121.113035	4.134644e+05
min	14.000000	26.000000	1.000000	0.000000	3.008590e+05
25%	59.000000	65.000000	1.000000	359.000000	3.404215e+05
50%	67.000000	73.000000	1.000000	482.000000	7.936830e+05
75%	75.000000	82.000000	1.000000	845.000000	1.178260e+06
max	144.000000	177.000000	2.000000	518157.000000	1.504991e+06

Step 3 — Clean Types + Create Targets (home_win, margin, total_points)

We will:

- Convert `score_home`, `score_vis` to numeric (already validated)
- Convert `gameday` to datetime
- Normalize `neutral` to Y/N
- Filter to completed games (`gamestatus == "Final"` if present)
- Create:
 - `home_win` (target)
 - `margin`
 - `total_points`

Then we immediately visualize:

- Home vs Visitor score distributions
- Margin distribution (hist + box + violin)
- Home win rate overall and by neutral site

```
In [4]: # =====
# Clean types + create targets
# =====

g = games.copy()

# Scores -> numeric
g["score_home"] = pd.to_numeric(g["score_home"], errors="coerce")
g["score_vis"] = pd.to_numeric(g["score_vis"], errors="coerce")

# Date -> datetime
g["gameday"] = pd.to_datetime(g["gameday"], errors="coerce")

# Keep only rows with scores
g = g.dropna(subset=["score_home", "score_vis"])

# Optional: keep completed games only
if "gamestatus" in g.columns:
    # common values: Final, FINAL, etc.
    g["gamestatus"] = g["gamestatus"].astype(str).str.strip().str.lower()
    g = g[g["gamestatus"].isin(["final", "final/ot", "finalot", "final (ot)"])

# Normalize neutral flag to Y/N
def to_yn(x):
    if pd.isna(x):
        return np.nan
    s = str(x).strip().lower()
    if s in ["y", "yes", "true", "1", "t"]:
        return "Y"
    if s in ["n", "no", "false", "0", "f"]:
        return "N"
    return s.upper()

g["neutral"] = g["neutral"].map(to_yn)

# Targets
g["home_win"] = (g["score_home"] > g["score_vis"]).astype(int)
g["margin"] = g["score_home"] - g["score_vis"]
g["total_points"] = g["score_home"] + g["score_vis"]

print("Cleaned games shape:", g.shape)
print("Home win rate:", round(float(g["home_win"].mean()), 4))

print("\nNeutral distribution (counts):")
print(g["neutral"].value_counts(dropna=False))

print("\nMargin summary:")
print(g["margin"].describe())

# =====
```

```

# Visuals (first core block)
# =====

# Score distributions
plt.figure(figsize=(8, 4))
sns.kdeplot(g["score_home"], label="Home score", fill=True)
sns.kdeplot(g["score_vis"], label="Visitor score", fill=True)
plt.title("Score distributions: Home vs Visitor")
plt.legend()
plt.tight_layout()
plt.show()

# Margin distribution
plt.figure(figsize=(8, 4))
sns.histplot(g["margin"], bins=60, kde=True)
plt.axvline(0, color="red")
plt.title("Margin distribution (home - visitor)")
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 2.2))
sns.boxplot(x=g["margin"])
plt.title("Margin (boxplot)")
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 2.2))
sns.violinplot(x=g["margin"])
plt.title("Margin (violin)")
plt.tight_layout()
plt.show()

# Home win rate overall + by neutral site
plt.figure(figsize=(5, 4))
sns.barplot(x=["away_win", "home_win"], y=[1 - g["home_win"].mean(), g["home_win"].mean()])
plt.title("Home win rate (overall)")
plt.tight_layout()
plt.show()

neutral_wr = g.groupby("neutral")["home_win"].mean().dropna()
plt.figure(figsize=(5, 4))
neutral_wr.plot(kind="bar")
plt.title("Home win rate by neutral flag")
plt.ylabel("home win rate")
plt.tight_layout()
plt.show()

```

Cleaned games shape: (99803, 21)

Home win rate: 0.6486

Neutral distribution (counts):

neutral

NaN 88371

Y 11432

Name: count, dtype: int64

Margin summary:

count 99803.000000

mean 6.251566

std 15.852792

min -61.000000



25% -5.000000

50% 6.000000


75% 15.000000

max 108.000000


Name: margin, dtype: float64




Margin (boxplot)




Margin (violin)



Home win rate (overall)





Step 4 — Structural Effects

We now examine structural drivers before merging ELO or stats:

1. Margin distribution by neutral site
2. Home win rate by neutral
3. Overtime impact on margin
4. Top teams by win rate (dominance patterns)
5. Monthly trend of margin (if date is usable)

This helps us understand:

- Tournament dynamics
- Home court strength
- Whether dominance is stable or concentrated

```
In [5]: # =====
# Step 4 – Structural effects (visual-first)
# =====

# 1) Neutral: treat NaN as "N" for analysis (most likely non-neutral games)
g["neutral2"] = g["neutral"].fillna("N")

print("Neutral2 counts:")
print(g["neutral2"].value_counts())

# --- Home win rate by neutral2 ---
wr_by_neutral = g.groupby("neutral2")["home_win"].mean().sort_index()
```

```

plt.figure(figsize=(5,4))
wr_by_neutral.plot(kind="bar")
plt.title("Home win rate: Neutral vs Non-neutral")
plt.ylabel("Home win rate")
plt.tight_layout()
plt.show()

# --- Margin distribution by neutral2 ---
plt.figure(figsize=(8,4))
sns.kdeplot(g[g["neutral2"]=="N"]["margin"], label="Non-neutral (N)", fill=True)
sns.kdeplot(g[g["neutral2"]=="Y"]["margin"], label="Neutral (Y)", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution by neutral site")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure(figsize=(8,4))
sns.boxplot(x="neutral2", y="margin", data=g)
plt.title("Margin by neutral site (boxplot)")
plt.tight_layout()
plt.show()

# 2) Overtime impact (normalize overtime flag)
def to_ot(x):
    if pd.isna(x):
        return "N"
    s = str(x).strip().lower()
    if s in ["y", "yes", "true", "1", "ot", "t"]:
        return "Y"
    return "N"

g["ot_flag"] = g["overtime"].map(to_ot)

print("\nOT counts:")
print(g["ot_flag"].value_counts())

# --- Margin by OT ---
plt.figure(figsize=(8,4))
sns.kdeplot(g[g["ot_flag"]=="N"]["margin"], label="No OT", fill=True)
sns.kdeplot(g[g["ot_flag"]=="Y"]["margin"], label="OT", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: OT vs No OT")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure(figsize=(8,4))
sns.boxplot(x="ot_flag", y="margin", data=g)
plt.title("Margin: OT vs No OT (boxplot)")
plt.tight_layout()
plt.show()

# 3) Top teams by win rate (home team perspective + overall)
# Home-only win rate for teams listed as home

```



```

home_team_wr = g.groupby("home")["home_win"].mean().sort_values(ascending=False)

plt.figure(figsize=(8,5))
home_team_wr.sort_values().plot(kind="barh")
plt.title("Top 15 home teams by home-win rate")
plt.xlabel("Home win rate")
plt.tight_layout()
plt.show()

# Overall win rate using winner_code/loser_code (these exist, use for description)
# Build overall W/L counts per team code
wins = g["winner_code"].value_counts()
losses = g["loser_code"].value_counts()
wl = pd.DataFrame({"wins": wins, "losses": losses}).fillna(0)
wl["games"] = wl["wins"] + wl["losses"]
wl["win_rate"] = wl["wins"] / wl["games"]
wl = wl[wl["games"] >= 10].sort_values("win_rate", ascending=False)

plt.figure(figsize=(8,5))
wl.head(15)["win_rate"].sort_values().plot(kind="barh")
plt.title("Top 15 team codes by overall win rate (min 10 games)")
plt.xlabel("Win rate")
plt.tight_layout()
plt.show()

# 4) Monthly trend of margin (if gameday parsed)
if pd.api.types.is_datetime64_any_dtype(g["gameday"]):
    tmp = g.dropna(subset=["gameday"]).copy()
    tmp["month"] = tmp["gameday"].dt.to_period("M").dt.to_timestamp()

    monthly_margin = tmp.groupby("month")["margin"].mean()
    monthly_home_wr = tmp.groupby("month")["home_win"].mean()

    plt.figure(figsize=(10,4))
    monthly_margin.plot()
    plt.title("Monthly average margin")
    plt.ylabel("Avg margin")
    plt.tight_layout()
    plt.show()

    plt.figure(figsize=(10,4))
    monthly_home_wr.plot()
    plt.title("Monthly home win rate")
    plt.ylabel("Home win rate")
    plt.tight_layout()
    plt.show()
else:
    print("\n'gameday' is not datetime; monthly plots skipped.")

```



Neutral2 counts:


neutral2

N 88371


Y 11432



Name: count, dtype: int64








OT counts:
ot_flag
N 99132
Y 671
Name: count, dtype: int64







Visual Interpretation — Men's Structural Effects (Verbose)

Visually, men's basketball shows structure, but it is not cleanly deterministic.

The margin distribution is broad and centered. There is a noticeable density around close games, which implies that competitive outcomes are common, not rare edge cases. Blowouts exist, but they do not visually dominate the league-level distribution. Put differently: there is signal in team strength, but variance still "survives" across the full range of outcomes.

Neutral-site effects reinforce that structure matters, but also that it does not fully decide games. When games move to neutral floors, the distribution compresses somewhat (and the home win advantage should fall). That supports the idea that venue contributes a real structural effect. However, the remaining spread indicates that even after removing home advantage, men's games still contain meaningful dispersion. This is consistent with a league where matchups, shooting variance, and late-game dynamics can meaningfully swing results.

Overtime impacts appear more like variance amplification than a structural regime shift. OT games increase randomness and compress margins, but visually they do not redefine overall dominance patterns. They are more consistent with "the game was already close" rather than "a different kind of game."

The dominance patterns (top teams by win rate) are present, but the curve looks gradual rather than cliff-like. Elite programs exist, but the separation from the middle is not extreme. That implies the league has tiers, but not a rigid hierarchy. Many teams are plausibly competitive within their tier, and even top teams are not visually separated to the point where outcomes become nearly automatic.

Overall, based purely on visuals, men's basketball appears structured but variance-driven: outcomes correlate with strength, yet there remains substantial room for upset behavior and overlap across performance tiers.

Step 5 — Quantify Structural Effects

We now formally test:

1. Home win rate difference: neutral vs non-neutral
2. Margin difference: neutral vs non-neutral
3. Margin difference: OT vs non-OT

We compute:

- Mean differences
- Cohen's d
- Two-sample t-test

```

In [6]: # =====
# Step 5 – Quantify structural effects
# =====

def cohens_d(a, b):
    a = np.asarray(a)
    b = np.asarray(b)
    a = a[~np.isnan(a)]
    b = b[~np.isnan(b)]
    if len(a) < 2 or len(b) < 2:
        return np.nan
    s = np.sqrt(((a.std(ddof=1) ** 2) + (b.std(ddof=1) ** 2)) / 2)
    if s == 0:
        return np.nan
    return (a.mean() - b.mean()) / s

def ttest_report(a, b):
    a = np.asarray(a)
    b = np.asarray(b)
    a = a[~np.isnan(a)]
    b = b[~np.isnan(b)]
    if len(a) < 2 or len(b) < 2:
        return (np.nan, np.nan)
    t, p = ttest_ind(a, b, equal_var=False)
    return (t, p)

# --- 1) Home win rate: Neutral vs Non-neutral ---
wr_N = g[g["neutral2"] == "N"]["home_win"].mean()
wr_Y = g[g["neutral2"] == "Y"]["home_win"].mean()

print("Home win rate (Non-neutral N):", round(float(wr_N), 4))
print("Home win rate (Neutral Y):", round(float(wr_Y), 4))
print("Difference (N - Y):", round(float(wr_N - wr_Y), 4))

plt.figure(figsize=(5,4))
sns.barplot(x=["Non-neutral (N)", "Neutral (Y)"], y=[wr_N, wr_Y])
plt.title("Home win rate: Neutral vs Non-neutral")
plt.ylim(0, 1)
plt.tight_layout()
plt.show()

# --- 2) Margin: Neutral vs Non-neutral ---
m_N = g[g["neutral2"] == "N"]["margin"].values
m_Y = g[g["neutral2"] == "Y"]["margin"].values

d_margin = cohens_d(m_N, m_Y)
t_stat, p_val = ttest_report(m_N, m_Y)

print("\nMargin (mean) Non-neutral:", round(float(np.nanmean(m_N)), 3))
print("Margin (mean) Neutral:", round(float(np.nanmean(m_Y)), 3))
print("Mean diff (N - Y):", round(float(np.nanmean(m_N) - np.nanmean(m_Y)), 3))
print("Cohen's d (N - Y):", round(float(d_margin), 3))
print("Welch t-test p-value:", round(float(p_val), 3))

plt.figure(figsize=(8,4))

```

```

sns.kdeplot(m_N, label="Non-neutral (N)", fill=True)
sns.kdeplot(m_Y, label="Neutral (Y)", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: Neutral vs Non-neutral")
plt.legend()
plt.tight_layout()
plt.show()

# --- 3) Margin: OT vs No OT ---
m_no = g[g["ot_flag"] == "N"]["margin"].values
m_ot = g[g["ot_flag"] == "Y"]["margin"].values


d_ot = cohens_d(m_no, m_ot)
t_stat2, p_val2 = ttest_report(m_no, m_ot)

print("\nMargin (mean) No OT:", round(float(np.nanmean(m_no)), 3))
print("Margin (mean) OT: ", round(float(np.nanmean(m_ot)), 3))
print("Mean diff (NoOT - OT):", round(float(np.nanmean(m_no) - np.nanmean(m_
print("Cohen's d (NoOT - OT):", round(float(d_ot), 3))
print("Welch t-test p-value: ", p_val2)


plt.figure(figsize=(8,4))
sns.kdeplot(m_no, label="No OT", fill=True)
sns.kdeplot(m_ot, label="OT", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: OT vs No OT")
plt.legend()
plt.tight_layout()
plt.show()

```


Home win rate (Non-neutral N): 0.667
 Home win rate (Neutral Y): 0.506
 Difference (N - Y): 0.161



Margin (mean) Non-neutral: 6.962
Margin (mean) Neutral: 0.758
Mean diff (N - Y): 6.204
Cohen's d (N - Y): 0.416
Welch t-test p-value: 0.0



Margin (mean) No OT: 6.289
Margin (mean) OT: 0.726
Mean diff (NoOT - OT): 5.563
Cohen's d (NoOT - OT): 0.455
Welch t-test p-value: 5.003625881614498e-75



Step 6 — Inspect Feature Tables (team_stats, elo, teams)

Before merging anything, we need to understand:

- What is the primary key?
- Is there season?

- Are stats per game or aggregated?
- Does elo contain pregame ratings?

We inspect structure + sample rows.

```
In [7]: # =====
# Step 6 – Inspect feature tables (TRULY IOPub-safe)
# =====

def inspect_df_safe(name, df, top_missing=15, preview_rows=2, preview_cols=10):
    print("\n" + "="*90)
    print(f"{name}")
    print("Shape:", df.shape)
    print("Num columns:", df.shape[1])

    # ----- Column printing (capped) -----
    print("\nColumns:")
    total_cols = df.shape[1]

    if total_cols <= max_cols_print:
        for c in df.columns:
            print(" ", c)
    else:
        print(f" Showing first {max_cols_print} of {total_cols} columns:")
        for c in df.columns[:max_cols_print]:
            print(" ", c)
        print(" ... (truncated)")

    # ----- Dtypes -----
    print("\nDtype counts:")
    print(df.dtypes.value_counts())

    # ----- Missingness -----
    miss = df.isna().mean().sort_values(ascending=False).head(top_missing)
    miss = miss[miss > 0]

    print(f"\nTop missingness (fraction, up to {top_missing}):")
    if len(miss) == 0:
        print(" (no missing values detected)")
    else:
        for k, v in miss.items():
            print(f" {k}: {v:.3f}")

    # ----- Safe Preview -----
    print(f"\nPreview: first {preview_rows} rows, first {preview_cols} columns")
    preview = df.iloc[:preview_rows, :min(preview_cols, total_cols)]
    display(preview)

# Run inspection
inspect_df_safe("teams", teams)
inspect_df_safe("team_stats", team_stats)
inspect_df_safe("elo", elo)
```

=====

=====

teams

Shape: (368, 3)

Num columns: 3

Columns:

code

name

location

Dtype counts:

object 3

Name: count, dtype: int64

Top missingness (fraction, up to 15):

location: 0.962

Preview: first 2 rows, first 12 columns

	code	name	location
0	AFA	Air Force Falcons	None
1	AKR	Akron Zips	None

```
=====
=====
team_stats
Shape: (6701, 30)
Num columns: 30
```

Columns:

```
team_name
team_code
2fa
2fm
3fa
3fm
apg
ast
blk
bpg
f
fga
fgm
fpg
fta
ftm
g
l
min
oreb
pts
reb
rpg
spg
stl
to
topg
w
winpct
season
```

Dtype counts:

```
float64    26
object      2
int64       1
Int64       1
```

Name: count, dtype: int64

Top missingness (fraction, up to 15):

```
team_name: 0.003
team_code: 0.003
```

Preview: first 2 rows, first 12 columns

	team_name	team_code	2fa	2fm	3fa	3fm	apg	ast	blk	bpg	
0	Air Force Falcons	AFA	657.0	320.0	530.0	199.0	11.9286	334.0	77.0	2.75	48
1	Akron Zips	AKR	1149.0	556.0	782.0	308.0	14.3235	487.0	102.0	3.00	70

```
=====
=====
```

```
elo
Shape: (6690, 5)
Num columns: 5
```

```
Columns:
  team
  code
  elo
  elorank
  season
```

```
Dtype counts:
object      4
Int64       1
Name: count, dtype: int64
```

```
Top missingness (fraction, up to 15):
(no missing values detected)
```

```
Preview: first 2 rows, first 12 columns
```

	team	code	elo	elorank	season
0	Duke	DUKE	2128	1	2008
1	Florida	FLA	2088	2	2008

```
In [8]: # =====
# Step 7 – Clean ELO table (SEASON-AWARE)
# =====

print("Original elo shape:", elo.shape)

elo_keep = ["code", "elo", "elorank"]
if "season" in elo.columns:
    elo_keep.append("season")

elo_clean = elo[elo_keep].copy()

elo_clean["elo"] = pd.to_numeric(elo_clean["elo"], errors="coerce")
elo_clean["elorank"] = pd.to_numeric(elo_clean["elorank"], errors="coerce")
if "season" in elo_clean.columns:
    elo_clean["season"] = pd.to_numeric(elo_clean["season"], errors="coerce")

print("Cleaned elo shape:", elo_clean.shape)
display(elo_clean.head())
```

Original elo shape: (6690, 5)

Cleaned elo shape: (6690, 4)

	code	elo	elrank	season
0	DUKE	2128	1	2008
1	FLA	2088	2	2008
2	MICH	2064	3	2008
3	HOU	2058	4	2008
4	ARIZ	2056	5	2008

```
In [9]: # =====
# Step 8 – Clean team_stats numerics
# =====

team_stats_clean = team_stats.copy()

# Keep useful columns and season key for multi-year joins
keep_cols = [
    "team_code",
    "season",
    "o_ppp", "d_ppp",
    "pace",
    "pts", "g", "min",
    "reb", "reb_d",
    "oreb", "dreb",
    "to", "to_d",
    "fgm", "fga",
    "ftm", "fta",
    "winpct"
]

keep_cols = [c for c in keep_cols if c in team_stats_clean.columns]
team_stats_clean = team_stats_clean[keep_cols]

for c in team_stats_clean.columns:
    if c != "team_code":
        team_stats_clean[c] = pd.to_numeric(team_stats_clean[c], errors="coerce")

# Possession estimate from box-score totals
if {"fga", "fta", "to", "oreb"}.issubset(set(team_stats_clean.columns)):
    team_stats_clean["poss_est"] = (
        team_stats_clean["fga"]
        + 0.44 * team_stats_clean["fta"]
        + team_stats_clean["to"]
        - team_stats_clean["oreb"]
    )

# Fallback offensive PPP
if "o_ppp" not in team_stats_clean.columns:
    team_stats_clean["o_ppp"] = np.nan
if {"pts", "poss_est"}.issubset(set(team_stats_clean.columns)):
    o_ppp_est = pd.Series(np.where(team_stats_clean["poss_est"] > 0, team_st
```

```

team_stats_clean["o_ppp"] = team_stats_clean["o_ppp"].fillna(o_ppp_est)

# Fallback pace (possessions per game)
if "pace" not in team_stats_clean.columns:
    team_stats_clean["pace"] = np.nan
if {"poss_est", "g"}.issubset(set(team_stats_clean.columns)):
    pace_est = pd.Series(np.where(team_stats_clean["g"] > 0, team_stats_clean["poss_est"], 0))
    team_stats_clean["pace"] = team_stats_clean["pace"].fillna(pace_est)

print("Cleaned team_stats shape:", team_stats_clean.shape)
for c in ["o_ppp", "pace", "winpct"]:
    if c in team_stats_clean.columns:
        print(f"nonnull {c}:", int(team_stats_clean[c].notna().sum()))
display(team_stats_clean.head())

```

Cleaned team_stats shape: (6701, 16)

nonnull o_ppp: 6682

nonnull pace: 6701

nonnull winpct: 6701

	team_code	season	pts	g	min	reb	oreb	to	fgm	fga	ftm
0	AFA	2008	1612.0	28	5650.0	685.0	122.0	320.0	519.0	1187.0	375.0
1	AKR	2008	2569.0	34	7275.0	1002.0	342.0	449.0	864.0	1931.0	533.0
2	ALA	2008	2564.0	33	6900.0	1194.0	415.0	443.0	942.0	2057.0	440.0
3	ALAM	2008	1669.0	26	5200.0	814.0	273.0	407.0	559.0	1469.0	395.0
4	ALB	2008	2048.0	30	6098.0	1025.0	319.0	425.0	722.0	1650.0	432.0

```

In [10]: # =====
# Step 9 - Build modeling table (SEASON-AWARE JOINS)
# =====

model_df = games.copy()

if "season" not in model_df.columns:
    raise ValueError("games table must include a season column for all-season")

# --- Merge home stats by team + season ---
home_stats = team_stats_clean.rename(columns={"team_code": "home_code"}).copy()
model_df = model_df.merge(home_stats, on=["home_code", "season"])
for col in team_stats_clean.columns:
    if col not in ["team_code", "season"] and col in model_df.columns:
        model_df.rename(columns={col: f"{col}_home"}, inplace=True)

# --- Merge visitor stats by team + season ---
vis_stats = team_stats_clean.rename(columns={"team_code": "visitor_code"}).copy()
model_df = model_df.merge(vis_stats, on=["visitor_code", "season"])
for col in team_stats_clean.columns:
    if col not in ["team_code", "season"] and col in model_df.columns and f"{col}_home" not in model_df.columns:
        model_df.rename(columns={col: f"{col}_vis"}, inplace=True)

# --- Merge home ELO by team + season ---
home_elo = elo_clean.rename(columns={"code": "home_code"}).copy()

```

```

home_elo_cols = [c for c in ["home_code", "elo", "elorange", "season"] if c in model_df.columns]
model_df = model_df.merge(home_elo[home_elo_cols], on=[c for c in ["home_code", "elo", "elorange", "season"] if c in model_df.columns])
if "elo" in model_df.columns:
    model_df.rename(columns={"elo": "elo_home"}, inplace=True)
if "elorange" in model_df.columns:
    model_df.drop(columns=["elorange"], inplace=True)

# --- Merge visitor ELO by team + season ---
vis_elo = elo_clean.rename(columns={"code": "visitor_code"}).copy()
vis_elo_cols = [c for c in ["visitor_code", "elo", "elorange", "season"] if c in vis_elo.columns]
model_df = model_df.merge(vis_elo[vis_elo_cols], on=[c for c in ["visitor_code", "elo", "elorange", "season"] if c in model_df.columns])
if "elo" in model_df.columns:
    model_df.rename(columns={"elo": "elo_vis"}, inplace=True)
if "elorange" in model_df.columns:
    model_df.drop(columns=["elorange"], inplace=True)

print("Final modeling df shape:", model_df.shape)
display(model_df.head())

```

Final modeling df shape: (100003, 48)

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	300941	Memphis Tigers	MEM	68	Kansas Jayhawks	KU	75
1	300939	North Carolina Tar Heels	UNC	66	Kansas Jayhawks	KU	84
2	300940	UCLA Bruins	UCLA	63	Memphis Tigers	MEM	78
3	502394	Bradley Braves	BRAD	64	Tulsa Golden Hurricane	TULS	70
4	502393	Ohio State Buckeyes	OSU	92	Massachusetts Minutemen	MASS	85

5 rows x 48 columns

```

In [11]: # =====
# Step 10 - Deduplicate games properly
# =====

print("Rows before dedupe:", model_df.shape[0])

if "gamestatus" in model_df.columns:
    gs = model_df["gamestatus"].astype(str).str.strip().str.lower()
    model_df = model_df[gs.str.startswith("final")].copy()

if "id" in model_df.columns:
    model_df = model_df.drop_duplicates(subset=["id"])

print("Rows after dedupe:", model_df.shape[0])
display(model_df.head())

```


Rows before dedupe: 100003

Rows after dedupe: 99803

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	300941	Memphis Tigers	MEM	68	Kansas Jayhawks	KU	75
1	300939	North Carolina Tar Heels	UNC	66	Kansas Jayhawks	KU	84
2	300940	UCLA Bruins	UCLA	63	Memphis Tigers	MEM	78
3	502394	Bradley Braves	BRAD	64	Tulsa Golden Hurricane	TULS	70
4	502393	Ohio State Buckeyes	OSU	92	Massachusetts Minutemen	MASS	85

5 rows x 48 columns

```
In [12]: # =====
# Step 11 – Feature engineering
# =====

model_df["home_win"] = (model_df["winner_code"] == model_df["home_code"]).as

for c in ["elo_home", "elo_vis", "o_ppp_home", "o_ppp_vis", "pace_home", "pa
    if c not in model_df.columns:
        model_df[c] = np.nan

model_df["elo_diff"] = model_df["elo_home"] - model_df["elo_vis"]
model_df["ppp_diff"] = model_df["o_ppp_home"] - model_df["o_ppp_vis"]
model_df["pace_diff"] = model_df["pace_home"] - model_df["pace_vis"]
model_df["winpct_diff"] = model_df["winpct_home"] - model_df["winpct_vis"]

# Extra EDA-friendly derived features
model_df["abs_elo_diff"] = model_df["elo_diff"].abs()
model_df["abs_ppp_diff"] = model_df["ppp_diff"].abs()
model_df["abs_pace_diff"] = model_df["pace_diff"].abs()
model_df["abs_winpct_diff"] = model_df["winpct_diff"].abs()
model_df["home_favored"] = (model_df["elo_diff"] > 0).astype("Int64")
model_df["upset_flag"] = (model_df["home_win"] != model_df["home_favored"]).

if {"score_home", "score_vis"}.issubset(set(model_df.columns)):
    sh = pd.to_numeric(model_df["score_home"], errors="coerce")
    sv = pd.to_numeric(model_df["score_vis"], errors="coerce")
    model_df["game_total"] = sh + sv
    model_df["margin_abs"] = (sh - sv).abs()

# Feature catalog for downstream relationship plots
extended_feature_candidates = [
    "elo_diff", "ppp_diff", "pace_diff", "winpct_diff",
    "abs_elo_diff", "abs_ppp_diff", "abs_pace_diff", "abs_winpct_diff",
```

```

    "home_favored", "upset_flag",
    "elo_home", "elo_vis",
    "o_ppp_home", "o_ppp_vis",
    "pace_home", "pace_vis",
    "winpct_home", "winpct_vis",
    "game_total", "margin_abs"
]
extended_feature_cols = []
for c in extended_feature_candidates:
    if c in model_df.columns and pd.api.types.is_numeric_dtype(model_df[c]):
        if model_df[c].notna().sum() > 100:
            extended_feature_cols.append(c)

display(model_df[[
    "home_win",
    "elo_diff",
    "ppp_diff",
    "pace_diff",
    "winpct_diff",
    "abs_elo_diff",
    "upset_flag"
]].head())
print("Extended feature count:", len(extended_feature_cols))
print("Sample features:", extended_feature_cols[:12])

```

	home_win	elo_diff	ppp_diff	pace_diff	winpct_diff	abs_elo_diff	upset_flag
0	1	363.0	0.030004	-1.439811	-0.0276	363.0	0
1	1	60.0	0.008507	-7.420324	-0.0020	60.0	0
2	1	-234.0	0.004516	5.183396	0.0568	234.0	1
3	1	14.0	-0.033063	2.238947	0.1316	14.0	0
4	0	-507.0	0.013111	11.128000	0.0672	507.0	0

Extended feature count: 20

Sample features: ['elo_diff', 'ppp_diff', 'pace_diff', 'winpct_diff', 'abs_elo_diff', 'abs_ppp_diff', 'abs_pace_diff', 'abs_winpct_diff', 'home_favored', 'upset_flag', 'elo_home', 'elo_vis']

```

In [13]: # =====
# Step 12 – Feature Distribution & Separation (Pure EDA)
# =====

if "extended_feature_cols" in locals() and len(extended_feature_cols) > 0:
    candidate_cols = extended_feature_cols
else:
    candidate_cols = ["elo_diff", "ppp_diff", "pace_diff", "winpct_diff"]

usable_cols = []
for col in candidate_cols:
    if col not in model_df.columns:
        continue
    tmp = model_df[[col, "home_win"]].dropna()
    if tmp.empty or tmp["home_win"].nunique() < 2:

```

```

        continue
    usable_cols.append(col)

# Rank by separability (absolute mean gap) and keep a manageable set for plots
rank_rows = []
for col in usable_cols:
    t = model_df[[col, "home_win"]].dropna()
    gap = t.loc[t.home_win==1, col].mean() - t.loc[t.home_win==0, col].mean()
    rank_rows.append((col, abs(gap)))
ranked = [c for c, _ in sorted(rank_rows, key=lambda x: x[1], reverse=True)]
plot_cols = ranked[:12]

for col in plot_cols:
    plt.figure(figsize=(8,4))
    sns.histplot(data=model_df, x=col, hue="home_win", bins=40, element="step")
    plt.title(f"{col} Distribution by Home Win")
    plt.show()

for col in plot_cols[:8]:
    plt.figure(figsize=(8,4))
    sns.kdeplot(data=model_df, x=col, hue="home_win", fill=True, common_norm=False)
    plt.title(f"{col} KDE Separation")
    plt.show()

for col in plot_cols:
    plot_df = model_df[["home_win", col]].dropna()
    if plot_df.empty:
        continue
    plt.figure(figsize=(6,4))
    sns.boxplot(data=plot_df, x="home_win", y=col)
    plt.title(f"{col} vs Outcome")
    plt.show()



print("\n=== Statistical Separation ===")
for col in plot_cols:
    wins = model_df.loc[model_df["home_win"] == 1, col].dropna()
    losses = model_df.loc[model_df["home_win"] == 0, col].dropna()
    if len(wins) < 2 or len(losses) < 2:
        continue
    mean_diff = wins.mean() - losses.mean()
    t_stat, p_val = ttest_ind(wins, losses, equal_var=False)
    pooled_std = np.sqrt((wins.var() + losses.var()) / 2)
    cohens_d = mean_diff / pooled_std if pooled_std and not np.isnan(pooled_std) else 0
    print(f"{col}: mean_diff={mean_diff:.4f}, d={cohens_d:.3f}, p={p_val:.3e}")

corr_cols = plot_cols + ["home_win"]
if len(corr_cols) >= 3:
    plt.figure(figsize=(10,8))
    corr = model_df[corr_cols].corr()
    sns.heatmap(corr, annot=False, cmap="coolwarm", center=0)
    plt.title("Feature Correlation Matrix (Top Features)")
    plt.show()


if plot_cols:
    relationship_feature_cols = plot_cols
else:
    relationship_feature_cols = []

```


```
relationship_feature_cols = []  
print("Relationship feature set size:", len(relationship_feature_cols))
```




elo_vis Distribution by Home Win




abs_elo_diff Distribution by Home Win







margin_abs Distribution by Home Win







pace_diff Distribution by Home Win




































margin_abs vs Outcome





pace_diff vs Outcome










```

=== Statistical Separation ===
elo_diff: mean_diff=199.8514, d=0.660, p=0.000e+00
elo_home: mean_diff=99.7491, d=0.339, p=0.000e+00
elo_vis: mean_diff=-87.6118, d=-0.303, p=0.000e+00
abs_elo_diff: mean_diff=52.0917, d=0.257, p=0.000e+00
margin_abs: mean_diff=5.2165, d=0.520, p=0.000e+00
pace_diff: mean_diff=4.3403, d=0.331, p=0.000e+00
pace_vis: mean_diff=-3.7703, d=-0.312, p=0.000e+00
abs_pace_diff: mean_diff=3.2822, d=0.272, p=0.000e+00
game_total: mean_diff=1.5521, d=0.078, p=6.912e-32
pace_home: mean_diff=0.5699, d=0.109, p=1.898e-52
winpct_diff: mean_diff=0.2890, d=1.341, p=0.000e+00
home_favored: mean_diff=0.2050, d=0.419, p=0.000e+00

```



Relationship feature set size: 12

Visual Feature Separation — Men's (Verbose)

Elo difference visually provides a strong directional signal, but the overlap is still meaningful. Wins are more frequent when Elo_diff is positive, as expected, but losses still appear throughout moderate Elo gaps. The key visual point is that the density curves

intersect substantially near the center and do not cleanly separate. This implies Elo is useful, but not sufficient to fully resolve outcomes in men's basketball.

Offensive efficiency difference (PPP_diff) appears to provide additional separation. The distribution suggests that teams with higher offensive efficiency relative to their opponent tend to win more often, and this separation may be non-trivial even when Elo overlap remains. Visually, PPP_diff looks like a plausible "second axis" that helps explain why similarly-rated teams can still produce distinct outcomes. This aligns with a league where execution and efficiency can meaningfully differentiate within tiers.

Pace difference shows heavy overlap between winners and losers. The visual separability is weak, meaning tempo alone does not appear to predict victory consistently. Pace may matter in matchup context, but visually it does not function as a stable global predictor.

Win percentage difference mirrors Elo_diff in structure, but again overlap remains substantial. This makes sense: win% reflects schedule and tiering, but in a league with higher parity and more variation, win% alone does not fully discriminate outcomes.

The combined visual interpretation is that men's basketball has multiple partially-informative signals (Elo, efficiency, win%), but none creates a clean boundary. Overlap is the defining visual characteristic. This is a "predictive but noisy" environment, consistent with higher upset rates and more variance conditional on team strength.

```
In [14]: # =====
# Step 13 – Feature Interactions & Nonlinearity
# =====

if "relationship_feature_cols" in locals() and relationship_feature_cols:
    rel_cols = relationship_feature_cols[:8]
else:
    rel_cols = ["elo_diff", "winpct_diff", "ppp_diff", "pace_diff"]

# Pairwise comparisons among selected features
pairs = []
for i in range(len(rel_cols)):
    for j in range(i+1, len(rel_cols)):
        pairs.append((rel_cols[i], rel_cols[j]))

usable_pairs = []
for x_col, y_col in pairs:
    if x_col not in model_df.columns or y_col not in model_df.columns:
        continue
    pair_df = model_df[[x_col, y_col, "home_win"]].dropna()
    if pair_df.empty or pair_df["home_win"].nunique() < 2:
        continue
    usable_pairs.append((x_col, y_col))

# Scatter by class (0/1)
for x_col, y_col in usable_pairs[:12]:
```

```

plt.figure(figsize=(7,6))
sns.scatterplot(data=model_df, x=x_col, y=y_col, hue="home_win", alpha=0.5)
plt.title(f"{x_col} vs {y_col} (home_win 0/1)")
plt.show()

# 2D KDE by class where density is stable
for x_col, y_col in usable_pairs[:8]:
    pair_df = model_df[[x_col, y_col, "home_win"]].dropna()
    if len(pair_df) < 200:
        continue
    plt.figure(figsize=(7,6))
    sns.kdeplot(data=pair_df, x=x_col, y=y_col, hue="home_win", fill=True, t
    plt.title(f"Density: {x_col} x {y_col}")
    plt.show()

# Pairplot matrix for strongest 5 features vs outcome
pairplot_cols = rel_cols[:5]
pairplot_df = model_df[pairplot_cols + ["home_win"]].dropna()
if not pairplot_df.empty:
    if len(pairplot_df) > 4000:
        pairplot_df = pairplot_df.sample(4000, random_state=42)
    sns.pairplot(pairplot_df, vars=pairplot_cols, hue="home_win", corner=True)
    plt.show()


# Monotonicity / win-rate by bins for many features
def win_rate_by_bins(col, bins=10):
    if col not in model_df.columns:
        return
    df = model_df[[col, "home_win"]].dropna().copy()
    if df.empty:
        return
    try:
        df["bin"] = pd.qcut(df[col], q=bins, duplicates="drop")
    except ValueError:
        return
    summary = df.groupby("bin")["home_win"].mean()
    if summary.empty:
        return
    plt.figure(figsize=(8,4))
    summary.plot(marker="o")
    plt.xticks(rotation=45)
    plt.title(f"Home Win Rate by {col} Quantile Bins")
    plt.ylabel("Home Win Rate")
    plt.show()

for col in rel_cols:
    win_rate_by_bins(col, bins=10)


print("Interaction pairs evaluated:", len(usable_pairs))


```

elo_diff vs elo_home (home_win 0/1)




elo_diff vs elo_vis (home_win 0/1)







elo_diff vs margin_abs (home_win 0/1)




elo_diff vs pace_diff (home_win 0/1)




elo_diff vs pace_vis (home_win 0/1)




elo_diff vs abs_pace_diff (home_win 0/1)




elo_home vs elo_vis (home_win 0/1)




elo_home vs abs_elo_diff (home_win 0/1)




elo_home vs margin_abs (home_win 0/1)




elo_home vs pace_diff (home_win 0/1)




elo_home vs pace_vis (home_win 0/1)













Density: elo_diff × margin_abs









Density: elo_diff × pace_vis























Interaction pairs evaluated: 28

```
In [15]: # =====
# Step 14 – Continuous Margin Analysis (FULL SAFE VERSION)
# =====

# ---- Ensure numeric score columns ----
model_df["score_home"] = pd.to_numeric(model_df["score_home"], errors="coerce")
model_df["score_vis"] = pd.to_numeric(model_df["score_vis"], errors="coerce")
```

```

# ---- Create margin if missing ----
if "margin" not in model_df.columns:
    model_df["margin"] = model_df["score_home"] - model_df["score_vis"]

print("Margin summary:")
print(model_df["margin"].describe())

# -----
# 1 Margin Distribution
# -----
plt.figure(figsize=(8,5))
sns.histplot(model_df["margin"].dropna(), bins=60, kde=True)
plt.title("Margin Distribution")
plt.show()

# -----
# 2 Margin vs Diff Features
# -----

diff_features = ["elo_diff", "ppp_diff", "winpct_diff", "pace_diff"]

for col in diff_features:
    if col in model_df.columns:
        plt.figure(figsize=(7,5))
        sns.regplot(
            data=model_df,
            x=col,
            y="margin",
            scatter_kws={"alpha":0.3},
            line_kws={"color":"red"}
        )
        plt.title(f"Margin vs {col}")
        plt.show()

# -----
# 3 Margin by Outcome
# -----

plt.figure(figsize=(6,5))
sns.boxplot(data=model_df, x="home_win", y="margin")
plt.title("Margin by Outcome")
plt.show()

# -----
# 4 Margin Skew & Tail Diagnostics
# -----

skew = model_df["margin"].skew()
kurt = model_df["margin"].kurtosis()

print(f"Margin Skew: {skew:.3f}")
print(f"Margin Kurtosis: {kurt:.3f}")

# Extreme margins
extreme_games = model_df[np.abs(model_df["margin"]) > 30]

```




```
print("Extreme margin games (>30 pts):", len(extreme_games))



plt.figure(figsize=(8,5))
sns.histplot(extreme_games["margin"], bins=40, kde=True)
plt.title("Extreme Margin Distribution (>30)")
plt.show()
```

Margin summary:


count	99803.000000
mean	6.251566
std	15.852792
min	-61.000000
25%	-5.000000
50%	6.000000
75%	15.000000
max	108.000000

Name: margin, dtype: float64







Margin vs winpct_diff



Margin vs pace_diff






Margin Skew: 0.563

Margin Kurtosis: 1.353

Extreme margin games (>30 pts): 7312



```

In [16]: # =====
# Step 15 – PCA Structure Analysis (Full Safe)
# =====

pca_features = ["elo_diff", "ppp_diff", "pace_diff", "winpct_diff"]
pca_df = model_df[[c for c in pca_features if c in model_df.columns] + ["home_win"]]
print("PCA sample size:", len(pca_df))

if len(pca_df) > 1 and len([c for c in pca_features if c in pca_df.columns]):
    use_cols = [c for c in pca_features if c in pca_df.columns]
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(pca_df[use_cols])
    pca = PCA(n_components=2)
    X_pca = pca.fit_transform(X_scaled)
    print("Explained variance ratio:", pca.explained_variance_ratio_)


    plt.figure(figsize=(8,6))
    sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=pca_df["home_win"], alpha=0.5)
    plt.title("PCA Projection (Colored by Home Win)")
    plt.show()

    loadings = pd.DataFrame(pca.components_.T, index=use_cols, columns=["PC1", "PC2"])
    print("\nPCA Loadings:")
    display(loadings)
else:
    print("Skipping PCA: insufficient complete rows/features after dropna.")

```

PCA sample size: 94235

Explained variance ratio: [0.5218556 0.25155676]



PCA Loadings:

	PC1	PC2
elo_diff	0.458693	-0.139064
ppp_diff	0.627347	0.050902
pace_diff	0.024357	0.988897
winpct_diff	0.628843	0.012353

```
In [17]: # =====
# Step 16 – Unsupervised Regime Detection
# =====

if "X_scaled" in locals() and "X_pca" in locals() and "pca_df" in locals():
    kmeans = KMeans(n_clusters=3, random_state=42)
    clusters = kmeans.fit_predict(X_scaled)
    print("Silhouette Score:", silhouette_score(X_scaled, clusters))

    plt.figure(figsize=(8,6))
    sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=clusters, palette="Set2")
    plt.title("KMeans Clusters in PCA Space")
    plt.show()


    cluster_df = pca_df.copy()
    cluster_df["cluster"] = clusters
```

```

cluster_summary = cluster_df.groupby("cluster")["home_win"].mean()
print("\nWin rate by cluster:")
print(cluster_summary)
else:
    print("Skipping Step 16: PCA artifacts unavailable (insufficient data in

```

Silhouette Score: 0.20903924570233695



```

Win rate by cluster:
cluster
0    0.660881
1    0.878137
2    0.366628
Name: home_win, dtype: float64

```

```

In [18]: # =====
# Step 17 - 2D Conditional Win Probability Surface
# =====

def probability_surface(x_col, y_col, bins=12):
    df = model_df[[x_col, y_col, "home_win"]].dropna().copy()

    df["x_bin"] = pd.qcut(df[x_col], q=bins, duplicates="drop")
    df["y_bin"] = pd.qcut(df[y_col], q=bins, duplicates="drop")


    pivot = df.groupby(["x_bin", "y_bin"])["home_win"].mean().unstack()


    plt.figure(figsize=(9,7))
    sns.heatmap(pivot, cmap="coolwarm", center=0.5)

```

```
plt.title(f"Win Probability Surface: {x_col} × {y_col}")
plt.show()

probability_surface("elo_diff", "winpct_diff")
probability_surface("elo_diff", "ppp_diff")
```






```
In [19]: # =====
# Step 18 – Residual Home Advantage After Adjustment
# =====

# Adjust margin for ELO expectation (approx scaling)
model_df["elo_adjusted_margin"] = model_df["margin"] - model_df["elo_diff"] /

plt.figure(figsize=(8,5))
sns.histplot(model_df["elo_adjusted_margin"].dropna(), bins=60, kde=True)
plt.title("ELO-Adjusted Margin Distribution")
plt.show()

print("Mean adjusted margin:", model_df["elo_adjusted_margin"].mean())
```



Mean adjusted margin: 2.122851169947472

Additional EDA Required for Tournament Modeling — MBB

```
In [20]: # Team-season aggregation (MBB)
_tg = games.copy()
for c in ["score_home", "score_vis", "season"]:
    if c in _tg.columns:
        _tg[c] = pd.to_numeric(_tg[c], errors="coerce")
if "gameday" in _tg.columns:
    _tg["gameday"] = pd.to_datetime(_tg["gameday"], errors="coerce")

# NCAA tournament proxy for exclusion from selection features:
# no explicit flag in source, so use null-league games in Mar 15-Apr 15 window
_tg["_mm_proxy"] = (
    _tg.get("league").isna()
    & _tg.get("gameday").notna()
    & (_tg["gameday"].dt.month == 3)
    & (_tg["gameday"].dt.day >= 15)
) | (
    _tg.get("league").isna()
    & _tg.get("gameday").notna()
    & (_tg["gameday"].dt.month == 4)
    & (_tg["gameday"].dt.day <= 15)
)
_tg_selection = _tg.loc[~_tg["_mm_proxy"]].copy()
print("MBB selection pool (non-March-Madness games):", _tg_selection.shape)

home = pd.DataFrame({
```

```

    "season": _tg_selection.get("season"),
    "gameday": _tg_selection.get("gameday"),
    "team_code": _tg_selection.get("home_code"),
    "opp_code": _tg_selection.get("visitor_code"),
    "conference": _tg_selection.get("league") if "league" in _tg_selection.c
    "win": (_tg_selection.get("winner_code") == _tg_selection.get("home_code
    "margin": _tg_selection.get("score_home") - _tg_selection.get("score_vis
})
vis = pd.DataFrame({
    "season": _tg_selection.get("season"),
    "gameday": _tg_selection.get("gameday"),
    "team_code": _tg_selection.get("visitor_code"),
    "opp_code": _tg_selection.get("home_code"),
    "conference": _tg_selection.get("league") if "league" in _tg_selection.c
    "win": (_tg_selection.get("winner_code") == _tg_selection.get("visitor_c
    "margin": _tg_selection.get("score_vis") - _tg_selection.get("score_home
})
team_games = pd.concat([home, vis], ignore_index=True).dropna(subset=["seaso
team_games = team_games.sort_values(["season", "team_code", "gameday"])
team_games["rolling_win_10"] = team_games.groupby(["season", "team_code"])["

selection_df = team_games.groupby(["season", "team_code"], as_index=False).a
    games_played=("win", "size"),
    overall_winpct=("win", "mean"),
    avg_margin=("margin", "mean"),
    conf_winpct=("win", "mean"),
    nonconf_winpct=("win", "mean"),
    last10_winpct=("rolling_win_10", "last"),
)

if "team_stats_clean" in globals() and {"season", "team_code"}.issubset(team
    _ts = team_stats_clean.copy()
    for c in ["o_ppp", "d_ppp"]:
        if c not in _ts.columns:
            _ts[c] = np.nan
    _ts["net_rating"] = _ts["o_ppp"] - _ts["d_ppp"]
    keep = [c for c in ["season", "team_code", "net_rating", "o_ppp", "d_ppp
    selection_df = selection_df.merge(_ts[keep], on=["season", "team_code"],

if "elo_clean" in globals() and {"season", "code", "elo"}.issubset(elo_clean
    e = elo_clean[["season", "code", "elo"]].copy().rename(columns={"code":
    e["end_elo"] = pd.to_numeric(e["end_elo"], errors="coerce")
    selection_df = selection_df.merge(e, on=["season", "team_code"], how="le

selection_df["elo_rank"] = selection_df.groupby("season")["end_elo"].rank(me
selection_df["selected_proxy"] = np.where(selection_df["elo_rank"].notna() &

# SOS proxy
opp_elo = selection_df[["season", "team_code", "end_elo"]].rename(columns={"
_sos = team_games.merge(opp_elo, on=["season", "opp_code"], how="left")
_sos = _sos.groupby(["season", "team_code"], as_index=False)["opp_end_elo"].
selection_df = selection_df.merge(_sos, on=["season", "team_code"], how="lef

# conference proxy: dominant league tag in season
conf = team_games.groupby(["season", "team_code", "conference"], as_index=Fa
conf = conf.drop_duplicates(["season", "team_code"]).drop(columns=['size']).

```

```
selection_df = selection_df.merge(conf, on=["season", "team_code"], how="left")
print("selection_df:", selection_df.shape)
display(selection_df.head())
```

MBB selection pool (non-March-Madness games): (97889, 19)
 selection_df: (6682, 17)

	season	team_code	games_played	overall_winpct	avg_margin	conf_winpct	nonconf_winpct
0	2008	AFA	24	0.541667	1.791667	0.541667	0.541667
1	2008	AKR	28	0.714286	8.250000	0.714286	0.714286
2	2008	ALA	31	0.516129	2.451613	0.516129	0.516129
3	2008	ALAM	23	0.391304	-3.869565	0.391304	0.391304
4	2008	ALB	30	0.500000	1.866667	0.500000	0.500000

Structural Shift — From Game-Level to Team-Season Modeling

At this stage, the modeling problem changes fundamentally.

Game-level analysis focuses on matchup variance and conditional win probability.
 Tournament selection, however, is a team-season classification problem.

Each row now represents a team's full body of work for a given season. The relevant question is not "did this team win a specific game," but rather:

Does this season profile meet the structural threshold for tournament inclusion?

This aggregation reframes the predictive task. Selection modeling depends on:

- Sustained performance,
- Schedule-adjusted strength,
- Conference context,
- And potentially recency effects.

The remaining EDA focuses on identifying whether selected teams form a statistically separable cluster relative to non-selected teams.

```
In [21]: # Selected vs not-selected distribution plots (MBB)
metrics = ["end_elo", "overall_winpct", "net_rating", "avg_margin", "sos_proxy"]
metrics = [m for m in metrics if m in selection_df.columns]
for m in metrics:
    d = selection_df[[m, "selected_proxy"]].dropna()
    if d.empty or d["selected_proxy"].nunique() < 2:
        continue
    plt.figure(figsize=(8,4))
```


```



sns.kdeplot(data=d, x=m, hue="selected_proxy", fill=True, common_norm=False)
plt.title(f"MBB Selected vs Non-Selected (proxy): {m}")
plt.show()



if "end_elo" in selection_df.columns:
    s=selection_df[selection_df["selected_proxy"]==1]["end_elo"].dropna()
    n=selection_df[selection_df["selected_proxy"]==0]["end_elo"].dropna()
    if len(s) and len(n):
        print("min selected elo:", float(s.min()), "max non-selected elo:",

if {"overall_winpct", "sos_proxy", "selected_proxy"}.issubset(selection_df.c
d=selection_df[["overall_winpct", "sos_proxy", "selected_proxy"]].dropna()
if not d.empty and d["selected_proxy"].nunique()>=2:
    plt.figure(figsize=(8,6))
    sns.scatterplot(data=d, x="overall_winpct", y="sos_proxy", hue="sele
    plt.title(f"MBB Win% vs SOS (selection proxy)")
    plt.show()

```








min selected elo: 1566.0 max non-selected elo: 1566.0

MBB Win% vs SOS (selection proxy)



Visual Separation — Selected vs Not Selected

The selected vs non-selected comparisons reveal how deterministic the committee threshold appears.

If the distributions show:

- Minimal overlap → selection behaves like a clear statistical cutoff.
- Substantial overlap → committee behavior introduces subjectivity or secondary criteria.



Elo distribution is particularly important. A clean lower bound among selected teams suggests a de facto rating threshold. Overlap in mid-range Elo values indicates a bubble zone where additional factors likely influence inclusion.


Win percentage alone may appear separated, but without schedule context it risks overstating weaker conference teams. If schedule-adjusted strength reduces overlap more than raw win%, that suggests the committee implicitly values strength-of-schedule.

The key visual question is not whether selected teams are stronger on average (they should be), but whether there exists a clear structural boundary separating them.

```
In [22]: # Bubble zone visuals (MBB)
if "elo_rank" in selection_df.columns:
    bubble = selection_df[(selection_df["elo_rank"]>=30) & (selection_df["elo_rank"]<80)]
    print("bubble rows:", len(bubble))
    b_metrics=[m for m in ["avg_margin", "last10_winpct", "overall_winpct", "net_rating"]]
    for m in b_metrics:
        plot_df=bubble[["selected_proxy", m]].dropna()
        if plot_df.empty or plot_df["selected_proxy"].nunique()<2:
            continue
        plt.figure(figsize=(7,4))
        sns.boxplot(data=plot_df, x="selected_proxy", y=m)
        plt.title(f"MBB Bubble Zone (Ranks 30-80): {m}")
        plt.show()
```

bubble rows: 969





Bubble Zone Analysis — Where Selection Is Decided

The bubble region provides the most informative selection signal.

Top teams are automatically included. Bottom teams are clearly excluded.
The structural complexity lies in the mid-tier.

Within this Elo and performance band:

- Do selected teams have consistently stronger net rating?
- Do they show stronger recent form?
- Is conference strength visibly different?
- Does margin of victory separate them?

If separation inside the bubble band is weak, then selection likely incorporates contextual or non-quantified factors.

If separation is clear even inside this constrained range, then selection behaves closer to a statistical threshold model.

This region should heavily influence feature design for the selection classifier.

```
In [23]: # Conference-level analysis (MBB)
if {"season", "conference_proxy", "selected_proxy"}.issubset(selection_df.columns):
    conf_bids = selection_df.groupby(["season", "conference_proxy"], as_index=False)
    conf_strength = selection_df.groupby("conference_proxy", as_index=False)
    display(conf_bids.head())
```

```


display(conf_strength.sort_values("avg_end_elo", ascending=False).head(2)

top = conf_bids.groupby("conference_proxy", as_index=False)["bids_proxy"]
plt.figure(figsize=(10,4))
sns.barplot(data=top, x="conference_proxy", y="bids_proxy")
plt.title(f"MBB Conference Bid Concentration (proxy)")
plt.xticks(rotation=45)
plt.show()

```

	season	conference_proxy	bids_proxy
0	2008	A-10	3
1	2008	A-EAST	0
2	2008	A-SUN	1
3	2008	ACC	7
4	2008	AMER	2

	conference_proxy	avg_end_elo
25	SEC	1738.292308
5	BIG10	1728.438462
6	BIG12	1709.830189
7	BIGEAST	1630.745614
3	ACC	1606.010989
23	PAC-12	1563.107143
4	AMER	1545.083333
20	MWC	1483.617347
32	WCC	1442.279330
0	A-10	1406.934363
19	MVC	1353.142132
11	C-USA	1302.284444
15	IVY	1285.000000
31	WAC	1280.085714
12	CAA	1275.957944
10	BIGWEST	1242.379121
17	MAC	1231.995633
13	HL	1217.835052
29	SUNBELT	1207.960526
27	SOCON	1196.492386



Conference Effects — Structural Influence on Selection

Conference distribution is not incidental; it is structurally relevant.

If certain conferences consistently receive multiple bids despite overlapping performance distributions, then conference membership is a predictive feature.

Key observations to evaluate visually:

- Concentration of bids within high-Elo conferences.
- Whether mid-major teams require statistically stronger profiles to be selected.
- Differences between MBB and WBB conference clustering.



If women's basketball shows stronger concentration among fewer dominant conferences, that aligns with earlier hierarchy findings.


If men's basketball shows broader conference representation and greater overlap, that reinforces variance and structural diversity.

Conference membership may function as an implicit strength-of-schedule proxy and should likely be encoded in the selection model.

```
In [24]: # Recency / trend visuals (MBB)
trend_cols=[c for c in ["last10_winpct","overall_winpct","avg_margin"] if c
for c in trend_cols:
    d=selection_df[[c, "selected_proxy"]].dropna()
    if d.empty or d["selected_proxy"].nunique()<2:
        continue
    plt.figure(figsize=(8,4))
    sns.kdeplot(data=d, x=c, hue="selected_proxy", fill=True, common_norm=Fa
```

```
plt.title(f"MBB Trend Distribution: {c}")  
plt.show()
```





Recency and Trend Effects

Committees frequently reference how teams are playing late in the season.

Visually, trend features should be evaluated for separation power:

- Late-season Elo change.
- Last 10 game win rate.
- Margin trend in final month.

If selected teams show stronger positive trend distributions relative to non-selected teams — particularly in the bubble region — recency may function as a secondary selection discriminator.

If trend distributions overlap heavily, then season-long structural strength dominates recency.

This distinction matters for model construction: trend features may be nonlinear modifiers rather than primary drivers.

```
In [25]: # Bracket add-ons: upset curve, conditional feature signal, calibration, margin
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import roc_auc_score, log_loss
from sklearn.calibration import calibration_curve

bdf=model_df.copy()
for c in ["elo_diff","ppp_diff","pace_diff","winpct_diff","score_home","score_away"]:
```

```

    if c in bdf.columns:
        bdf[c]=pd.to_numeric(bdf[c], errors='coerce')
if "gameday" in bdf.columns:
    bdf["gameday"]=pd.to_datetime(bdf["gameday"], errors='coerce')

# NCAA tournament proxy: hold out for historical backtesting only.
bdf["_mm_proxy"] = (
    bdf.get("league").isna()
    & bdf.get("gameday").notna()
    & (bdf["gameday"].dt.month == 3)
    & (bdf["gameday"].dt.day >= 15)
) | (
    bdf.get("league").isna()
    & bdf.get("gameday").notna()
    & (bdf["gameday"].dt.month == 4)
    & (bdf["gameday"].dt.day <= 15)
)
train_bdf = bdf.loc[~bdf["_mm_proxy"]].copy()
backtest_bdf = bdf.loc[bdf["_mm_proxy"]].copy()
print("MBB bracket train rows (non-March-Madness):", len(train_bdf))
print("MBB March Madness holdout rows (backtest only):", len(backtest_bdf))

if {"elo_diff", "home_win"}.issubset(train_bdf.columns):
    bins=[0,25,50,75,100,10_000]
    labels=["0-25", "25-50", "50-75", "75-100", "100+"]
    train_bdf["elo_gap_bin"]=pd.cut(train_bdf["elo_diff"].abs(), bins=bins,
    train_bdf["home_favored"]=(train_bdf["elo_diff"]>0).astype(int)
    train_bdf["fav_won"]=np.where(train_bdf["home_favored"]==1, train_bdf["h
    train_bdf["upset"]=1-train_bdf["fav_won"]
    curve=train_bdf.groupby("elo_gap_bin", observed=False).agg(win_prob=("fa
    display(curve)
    plt.figure(figsize=(8,4))
    sns.lineplot(data=curve, x="elo_gap_bin", y="upset_rate", marker='o')
    plt.title(f"MBB Upset Rate by Elo Gap Bin (Regular Season / Conf Tournam
    plt.ylim(0,1)
    plt.show()

auc_rows=[]
for label in ["0-25", "25-50", "50-75", "75-100", "100+"]:
    sub=train_bdf[train_bdf.get("elo_gap_bin")==label]
    if sub is None or len(sub)<200:
        continue
    for feat in ["ppp_diff", "pace_diff", "winpct_diff"]:
        if feat not in sub.columns:
            continue
        d=sub[[feat, "home_win"]].dropna()
        if len(d)<120 or d["home_win"].nunique()<2:
            continue
        try:
            auc=max(roc_auc_score(d["home_win"], d[feat]), 1-roc_auc_score(c
        except Exception:
            auc=np.nan
        auc_rows.append({"elo_gap_bin":label, "feature":feat, "auc_abs":auc,
if auc_rows:
    display(pd.DataFrame(auc_rows))

```



```

cal_feats=[f for f in ["elo_diff","ppp_diff","pace_diff","winpct_diff"] if f
work=train_bdf[cal_feats+["home_win"]].dropna() if cal_feats and "home_win"
if len(work)>1000 and work["home_win"].nunique()==2:
    test=work.sample(frac=0.2, random_state=42)
    train=work.drop(test.index)
    elo_only=["elo_diff"] if "elo_diff" in cal_feats else [cal_feats[0]]

p1=Pipeline([("imp",SimpleImputer(strategy="median")),("sc",StandardScaler)])
p1.fit(train[elo_only], train["home_win"])
pr1=p1.predict_proba(test[elo_only])[:,1]

p2=Pipeline([("imp",SimpleImputer(strategy="median")),("sc",StandardScaler)])
p2.fit(train[cal_feats], train["home_win"])
pr2=p2.predict_proba(test[cal_feats])[:,1]

print("log_loss elo_only:", float(log_loss(test["home_win"], pr1)))
print("log_loss full:", float(log_loss(test["home_win"], pr2)))

pt1,pp1=calibration_curve(test["home_win"], pr1, n_bins=10)
pt2,pp2=calibration_curve(test["home_win"], pr2, n_bins=10)
plt.figure(figsize=(6,6))
plt.plot([0,1],[0,1], '--', color='gray')
plt.plot(pp1,pt1,marker='o',label='elo_only')
plt.plot(pp2,pt2,marker='o',label='full')
plt.title(f"MBB Calibration Curves (Train = Non-March-Madness)")
plt.legend(); plt.show()

if {"score_home","score_vis","elo_gap_bin"}.issubset(train_bdf.columns):
    train_bdf["margin"]=train_bdf["score_home"]-train_bdf["score_vis"]
    mv=train_bdf.groupby("elo_gap_bin", observed=False)["margin"].var().reset_index()
    display(mv)
    plt.figure(figsize=(8,4))
    sns.barplot(data=mv, x="elo_gap_bin", y="margin_variance")
    plt.title(f"MBB Margin Variance by Elo Gap Bin (Train Set)")
    plt.show()


# Persist league-specific datasets for cross-league model comparison
mbb_selection_df = selection_df.copy()
mbb_bracket_train_df = train_bdf.copy()
mbb_bracket_backtest_df = backtest_bdf.copy()

```

MBB bracket train rows (non-March-Madness): 97689

MBB March Madness holdout rows (backtest only): 2114

	elo_gap_bin	win_prob	upset_rate	n
0	0-25	0.494781	0.505219	6132
1	25-50	0.524255	0.475745	6205
2	50-75	0.526562	0.473438	5779
3	75-100	0.562808	0.437192	5684
4	100+	0.668754	0.331246	68339




	elo_gap_bin	feature	auc_abs	n
0	0-25	ppp_diff	0.723793	6132
1	0-25	pace_diff	0.509186	6132
2	0-25	winpct_diff	0.785623	6132
3	25-50	ppp_diff	0.711498	6205
4	25-50	pace_diff	0.506697	6205
5	25-50	winpct_diff	0.784638	6205
6	50-75	ppp_diff	0.704837	5779
7	50-75	pace_diff	0.508370	5779
8	50-75	winpct_diff	0.783184	5779
9	75-100	ppp_diff	0.721950	5684
10	75-100	pace_diff	0.502431	5684
11	75-100	winpct_diff	0.793640	5684
12	100+	ppp_diff	0.776689	68339
13	100+	pace_diff	0.500922	68339
14	100+	winpct_diff	0.830484	68339


log_loss elo_only: 0.6059131685390816

log_loss full: 0.48056721601940483

MBB Calibration Curves (Train = Non-March-Madness)



	elo_gap_bin	margin_variance
0	0-25	154.689905
1	25-50	157.992023
2	50-75	161.038826
3	75-100	164.260691
4	100+	221.314957



Elo Gap vs Upset Probability

Elo gap binning directly tests hierarchy strength.

If win probability increases smoothly and steeply with Elo difference, the league is structurally predictable.

If upset frequency remains high even at moderate Elo gaps, variance remains a meaningful force.

Comparing MBB and WBB curves:

- A steeper WBB curve supports the earlier visual conclusion of stronger hierarchy.
- A flatter MBB curve supports the conclusion of higher conditional variance.

This analysis directly informs bracket simulation assumptions.

WBB Duplicate Section

This section duplicates the same EDA workflow for NCAA Women's data using `wbb_*.parquet`.

NCAA WBB — Exhaustive Visual-First EDA (Local Parquet)

Working locally (API down). Files are in the same folder as this notebook:

- wbb_games.parquet
- wbb_teams.parquet
- wbb_team_stats.parquet
- wbb_elo.parquet

Notebook goals:

1. Load + standardize schemas (dirty-data tolerant)
2. Build core game table with targets (home_win, margin)
3. Merge in ELO + team/team_stats when possible
4. Exhaustive, visual-first EDA (distributions, comparisons, drift, correlation, PCA, clustering)

```
In [26]: # =====
# ALL IMPORTS (single cell)
# =====

# Standard library
import os
import re
import json
import math
import time
import random
import warnings
from pathlib import Path
from datetime import datetime

# Core DS
import numpy as np
import pandas as pd

# Viz
import matplotlib.pyplot as plt
import seaborn as sns

# Stats
from scipy import stats
from scipy.stats import ttest_ind, ks_2samp

# Sklearn (EDA structure)
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.decomposition import PCA, TruncatedSVD
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans, DBSCAN, AgglomerativeClustering
from sklearn.metrics import silhouette_score

# Quick separability sanity checks (EDA signal)
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score

warnings.filterwarnings("ignore")
sns.set_theme()
```

```
plt.style.use("seaborn-v0_8")

# Parquet engine check (required for pd.read_parquet)
ENGINE = None
try:
    import pyarrow # noqa: F401
    ENGINE = "pyarrow"
except Exception:
    try:
        import fastparquet # noqa: F401
        ENGINE = "fastparquet"
    except Exception as e:
        raise ImportError(
            "No parquet engine found. Install one:\n"
            "  pip install pyarrow    (recommended)\n"
            "or\n"
            "  pip install fastparquet\n"
        ) from e

print("Imports loaded. Parquet engine:", ENGINE)
```

Imports loaded. Parquet engine: pyarrow

0) Load Parquets (Local Folder) + Standardize Column Names

We assume the parquet files are in the same folder as the notebook (`Path.cwd()`).
Then we normalize all column names so downstream logic is stable.

```
In [27]: DATA_PATH = Path.cwd()

required_files = ["wbb_games.parquet", "wbb_teams.parquet", "wbb_team_stats.parquet"]
print("CWD:", DATA_PATH)
print("Parquet files in CWD:", [p.name for p in DATA_PATH.glob("*.parquet")])

missing = [f for f in required_files if not (DATA_PATH / f).exists()]
if missing:
    raise FileNotFoundError(f"Missing files in {DATA_PATH}: {missing}")

games = pd.read_parquet(DATA_PATH / "wbb_games.parquet", engine=ENGINE)
teams = pd.read_parquet(DATA_PATH / "wbb_teams.parquet", engine=ENGINE)
team_stats = pd.read_parquet(DATA_PATH / "wbb_team_stats.parquet", engine=ENGINE)
elo = pd.read_parquet(DATA_PATH / "wbb_elo.parquet", engine=ENGINE)

def clean_columns(df: pd.DataFrame) -> pd.DataFrame:
    df = df.copy()
    df.columns = (
        df.columns.astype(str)
        .str.strip()
        .str.lower()
        .str.replace(" ", "_")
        .str.replace("-", "_")
        .str.replace(r"[^a-z0-9_]", "", regex=True)
    )
```

```

    return df

games = clean_columns(games)
teams = clean_columns(teams)
team_stats = clean_columns(team_stats)
elo = clean_columns(elo)

# Optional season filter. Leave as None for all seasons.
TARGET_SEASONS = None # Example: [2024] or list(range(2018, 2025))

for df_name, df in (("games", games), ("team_stats", team_stats), ("elo", elo)):
    if "season" in df.columns:
        df["season"] = pd.to_numeric(df["season"], errors="coerce").astype("Int64")

if TARGET_SEASONS is not None:
    target = set(TARGET_SEASONS)
    if "season" in games.columns:
        games = games[games["season"].isin(target)].copy()
    if "season" in team_stats.columns:
        team_stats = team_stats[team_stats["season"].isin(target)].copy()
    if "season" in elo.columns:
        elo = elo[elo["season"].isin(target)].copy()

print("games:", games.shape)
print("teams:", teams.shape)
print("team_stats:", team_stats.shape)
print("elo:", elo.shape)

```

CWD: /Users/ceverson/Development/ml1_final_code

Parquet files in CWD: ['games.parquet', 'team_stats.parquet', 'wbb_team_stats.parquet', 'elo.parquet', 'wbb_elos.parquet', 'teams.parquet', 'wbb_games.parquet', 'wbb_teams.parquet']

games: (75550, 18)

teams: (366, 6)

team_stats: (5297, 30)

elo: (4937, 5)

Step 1 — Inspect `games` Table Structure

Before engineering targets or merging anything, we must:

- Identify home team column
- Identify away team column
- Identify home score column
- Identify away score column
- Check for season column
- Check for date column

Because the data is not clean, we will not assume column names. We will inspect schema and sample rows first.

```
In [28]: # =====
# Inspect games table
# =====

print("Shape:", games.shape)

print("\nColumns:")
for col in games.columns:
    print(" -", col)

print("\nDtypes:")
print(games.dtypes)

print("\nHead:")
display(games.head(5))

print("\nNumeric summary (coerce only likely-numeric columns):")

likely_numeric = ["score_vis", "score_home", "gameno", "venue_code", "id"]

games_numeric = games.copy()
for col in likely_numeric:
    if col in games_numeric.columns:
        games_numeric[col] = pd.to_numeric(games_numeric[col], errors="coerce")

# describe only the numeric columns we just coerced
numeric_only_df = games_numeric[[c for c in likely_numeric if c in games_numeric.columns]]

display(numeric_only_df.describe())
```


Shape: (75550, 18)

Columns:

- id
- visitor
- visitor_code
- score_vis
- home
- home_code
- score_home
- gamestatus
- overtime
- winner_code
- loser_code
- gameday
- gameno
- venue
- venue_code
- neutral
- league
- season

Dtypes:

id	object
visitor	object
visitor_code	object
score_vis	object
home	object
home_code	object
score_home	object
gamestatus	object
overtime	object
winner_code	object
loser_code	object
gameday	object
gameno	object
venue	object
venue_code	object
neutral	object
league	object
season	Int64

dtype: object

Head:

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	631025	Notre Dame Fightin` Irish	ND	61	Baylor Lady Bears	BAY	80
1	631024	Stanford Cardinal	STAN	47	Baylor Lady Bears	BAY	59
2	631023	Connecticut Huskies	CONN	75	Notre Dame Fightin` Irish	ND	83
3	631022	James Madison Dukes	JMU	68	Oklahoma State Cowboys	OKSU	75
4	631020	Kentucky Wildcats	UK	65	Connecticut Huskies	CONN	80

Numeric summary (coerce only likely-numeric columns):

	score_vis	score_home	gameno	venue_code	id
count	75434.000000	75434.000000	65379.0	75550.000000	7.555000e+04
mean	61.833682	67.907575	1.0	1505.619272	1.010376e+06
std	13.225703	14.039793	0.0	14213.297354	3.055905e+05
min	13.000000	13.000000	1.0	0.000000	5.979670e+05
25%	53.000000	58.000000	1.0	353.000000	6.396992e+05
50%	61.000000	67.000000	1.0	487.000000	1.052720e+06
75%	71.000000	77.000000	1.0	850.000000	1.281713e+06
max	133.000000	159.000000	1.0	518157.000000	1.524504e+06

Step 3 — Clean Types + Create Targets (home_win, margin, total_points)

We will:

- Convert `score_home`, `score_vis` to numeric (already validated)
- Convert `gameday` to datetime
- Normalize `neutral` to Y/N
- Filter to completed games (`gamestatus == "Final"` if present)
- Create:
 - `home_win` (target)
 - `margin`
 - `total_points`

Then we immediately visualize:

- Home vs Visitor score distributions
- Margin distribution (hist + box + violin)
- Home win rate overall and by neutral site

```
In [29]: # =====
# Clean types + create targets
# =====

g = games.copy()

# Scores -> numeric
g["score_home"] = pd.to_numeric(g["score_home"], errors="coerce")
g["score_vis"] = pd.to_numeric(g["score_vis"], errors="coerce")

# Date -> datetime
g["gameday"] = pd.to_datetime(g["gameday"], errors="coerce")

# Keep only rows with scores
g = g.dropna(subset=["score_home", "score_vis"])

# Optional: keep completed games only
if "gamestatus" in g.columns:
    # common values: Final, FINAL, etc.
    g["gamestatus"] = g["gamestatus"].astype(str).str.strip().str.lower()
    g = g[g["gamestatus"].isin(["final", "final/ot", "finalot", "final (ot)"])

# Normalize neutral flag to Y/N
def to_yn(x):
    if pd.isna(x):
        return np.nan
    s = str(x).strip().lower()
    if s in ["y", "yes", "true", "1", "t"]:
        return "Y"
    if s in ["n", "no", "false", "0", "f"]:
        return "N"
    return s.upper()

g["neutral"] = g["neutral"].map(to_yn)

# Targets
g["home_win"] = (g["score_home"] > g["score_vis"]).astype(int)
g["margin"] = g["score_home"] - g["score_vis"]
g["total_points"] = g["score_home"] + g["score_vis"]

print("Cleaned games shape:", g.shape)
print("Home win rate:", round(float(g["home_win"].mean()), 4))

print("\nNeutral distribution (counts):")
print(g["neutral"].value_counts(dropna=False))

print("\nMargin summary:")
print(g["margin"].describe())

# =====
```

```

# Visuals (first core block)
# =====

# Score distributions
plt.figure(figsize=(8, 4))
sns.kdeplot(g["score_home"], label="Home score", fill=True)
sns.kdeplot(g["score_vis"], label="Visitor score", fill=True)
plt.title("Score distributions: Home vs Visitor")
plt.legend()
plt.tight_layout()
plt.show()

# Margin distribution
plt.figure(figsize=(8, 4))
sns.histplot(g["margin"], bins=60, kde=True)
plt.axvline(0, color="red")
plt.title("Margin distribution (home - visitor)")
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 2.2))
sns.boxplot(x=g["margin"])
plt.title("Margin (boxplot)")
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 2.2))
sns.violinplot(x=g["margin"])
plt.title("Margin (violin)")
plt.tight_layout()
plt.show()

# Home win rate overall + by neutral site
plt.figure(figsize=(5, 4))
sns.barplot(x=["away_win", "home_win"], y=[1 - g["home_win"].mean(), g["home_win"].mean()])
plt.title("Home win rate (overall)")
plt.tight_layout()
plt.show()



neutral_wr = g.groupby("neutral")["home_win"].mean().dropna()
plt.figure(figsize=(5, 4))
neutral_wr.plot(kind="bar")
plt.title("Home win rate by neutral flag")
plt.ylabel("home win rate")
plt.tight_layout()
plt.show()




```


Cleaned games shape: (75425, 21)
Home win rate: 0.6208

Neutral distribution (counts):
neutral
NaN 68342
Y 7083
Name: count, dtype: int64

Margin summary:
count 75425.000000
mean 6.074644
std 19.023045
min -86.000000
25% -7.000000
50% 6.000000
75% 17.000000
max 141.000000
Name: margin, dtype: float64







Step 4 — Structural Effects

We now examine structural drivers before merging ELO or stats:

1. Margin distribution by neutral site
2. Home win rate by neutral
3. Overtime impact on margin
4. Top teams by win rate (dominance patterns)
5. Monthly trend of margin (if date is usable)

This helps us understand:

- Tournament dynamics
- Home court strength
- Whether dominance is stable or concentrated

```
In [30]: # =====
# Step 4 – Structural effects (visual-first)
# =====

# 1) Neutral: treat NaN as "N" for analysis (most likely non-neutral games)
g["neutral2"] = g["neutral"].fillna("N")

print("Neutral2 counts:")
print(g["neutral2"].value_counts())

# --- Home win rate by neutral2 ---
wr_by_neutral = g.groupby("neutral2")["home_win"].mean().sort_index()
```

```

plt.figure(figsize=(5,4))
wr_by_neutral.plot(kind="bar")
plt.title("Home win rate: Neutral vs Non-neutral")
plt.ylabel("Home win rate")
plt.tight_layout()
plt.show()

# --- Margin distribution by neutral2 ---
plt.figure(figsize=(8,4))
sns.kdeplot(g[g["neutral2"]=="N"]["margin"], label="Non-neutral (N)", fill=True)
sns.kdeplot(g[g["neutral2"]=="Y"]["margin"], label="Neutral (Y)", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution by neutral site")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure(figsize=(8,4))
sns.boxplot(x="neutral2", y="margin", data=g)
plt.title("Margin by neutral site (boxplot)")
plt.tight_layout()
plt.show()

# 2) Overtime impact (normalize overtime flag)
def to_ot(x):
    if pd.isna(x):
        return "N"
    s = str(x).strip().lower()
    if s in ["y", "yes", "true", "1", "ot", "t"]:
        return "Y"
    return "N"

g["ot_flag"] = g["overtime"].map(to_ot)

print("\nOT counts:")
print(g["ot_flag"].value_counts())

# --- Margin by OT ---
plt.figure(figsize=(8,4))
sns.kdeplot(g[g["ot_flag"]=="N"]["margin"], label="No OT", fill=True)
sns.kdeplot(g[g["ot_flag"]=="Y"]["margin"], label="OT", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: OT vs No OT")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure(figsize=(8,4))
sns.boxplot(x="ot_flag", y="margin", data=g)
plt.title("Margin: OT vs No OT (boxplot)")
plt.tight_layout()
plt.show()

# 3) Top teams by win rate (home team perspective + overall)
# Home-only win rate for teams listed as home

```



```

home_team_wr = g.groupby("home")["home_win"].mean().sort_values(ascending=False)

plt.figure(figsize=(8,5))
home_team_wr.sort_values().plot(kind="barh")
plt.title("Top 15 home teams by home-win rate")
plt.xlabel("Home win rate")
plt.tight_layout()
plt.show()

# Overall win rate using winner_code/loser_code (these exist, use for description)
# Build overall W/L counts per team code
wins = g["winner_code"].value_counts()
losses = g["loser_code"].value_counts()
wl = pd.DataFrame({"wins": wins, "losses": losses}).fillna(0)
wl["games"] = wl["wins"] + wl["losses"]
wl["win_rate"] = wl["wins"] / wl["games"]
wl = wl[wl["games"] >= 10].sort_values("win_rate", ascending=False)

plt.figure(figsize=(8,5))
wl.head(15)["win_rate"].sort_values().plot(kind="barh")
plt.title("Top 15 team codes by overall win rate (min 10 games)")
plt.xlabel("Win rate")
plt.tight_layout()
plt.show()

# 4) Monthly trend of margin (if gameday parsed)
if pd.api.types.is_datetime64_any_dtype(g["gameday"]):
    tmp = g.dropna(subset=["gameday"]).copy()
    tmp["month"] = tmp["gameday"].dt.to_period("M").dt.to_timestamp()

    monthly_margin = tmp.groupby("month")["margin"].mean()
    monthly_home_wr = tmp.groupby("month")["home_win"].mean()

    plt.figure(figsize=(10,4))
    monthly_margin.plot()
    plt.title("Monthly average margin")
    plt.ylabel("Avg margin")
    plt.tight_layout()
    plt.show()

    plt.figure(figsize=(10,4))
    monthly_home_wr.plot()
    plt.title("Monthly home win rate")
    plt.ylabel("Home win rate")
    plt.tight_layout()
    plt.show()
else:
    print("\n'gameday' is not datetime; monthly plots skipped.")

```



Neutral2 counts:


neutral2

N 68342


Y 7083



Name: count, dtype: int64







OT counts:
ot_flag
N 75054
Y 371
Name: count, dtype: int64







Visual Interpretation — Women's Structural Effects
(Verbose)

Visually, women's basketball shows stronger structural hierarchy than men's.

The margin distribution displays a heavier tail of decisive outcomes. Blowouts appear more common and more extreme relative to the league center. The density is not only near competitive games; there is visibly more probability mass in larger margins. This is a strong visual indicator of stratification: when top programs play mid or lower programs, the outcome distribution shifts more decisively.

Neutral-site comparisons still show compression, but the overall structure remains more pronounced than in men's. In other words, removing the home/venue effect does not remove the main separation. This suggests that the primary driver is not location—it is team strength, and the strength gaps appear larger and more consistent.

Overtime patterns again look like variance signals, but the key difference is how the league behaves outside those cases. Women's data visually suggests fewer "random" outcomes in the main mass of games relative to men's; the distribution is more consistent with a league in which large strength gaps translate into consistent results.

Dominance patterns (top teams by win rate) appear steeper and more polarized. The upper tier separates more sharply from the middle. That visual steepness implies that top programs do not just outperform slightly—they appear to sustain higher win rates with less overlap, which is what a top-heavy hierarchy looks like.

Overall, based strictly on visuals, women's basketball appears more hierarchy-driven: strength differences show up more directly in margins and win concentration, with less distributional blending between tiers.

Step 5 — Quantify Structural Effects

We now formally test:

1. Home win rate difference: neutral vs non-neutral
2. Margin difference: neutral vs non-neutral
3. Margin difference: OT vs non-OT

We compute:

- Mean differences
- Cohen's d
- Two-sample t-test

```
In [31]: # =====  
# Step 5 - Quantify structural effects  
# =====  
  
def cohens_d(a, b):
```

```

a = np.asarray(a)
b = np.asarray(b)
a = a[~np.isnan(a)]
b = b[~np.isnan(b)]
if len(a) < 2 or len(b) < 2:
    return np.nan
s = np.sqrt(((a.std(ddof=1) ** 2) + (b.std(ddof=1) ** 2)) / 2)
if s == 0:
    return np.nan
return (a.mean() - b.mean()) / s

def ttest_report(a, b):
    a = np.asarray(a)
    b = np.asarray(b)
    a = a[~np.isnan(a)]
    b = b[~np.isnan(b)]
    if len(a) < 2 or len(b) < 2:
        return (np.nan, np.nan)
    t, p = ttest_ind(a, b, equal_var=False)
    return (t, p)

# --- 1) Home win rate: Neutral vs Non-neutral ---
wr_N = g[g["neutral2"] == "N"]["home_win"].mean()
wr_Y = g[g["neutral2"] == "Y"]["home_win"].mean()

print("Home win rate (Non-neutral N):", round(float(wr_N), 4))
print("Home win rate (Neutral Y):", round(float(wr_Y), 4))
print("Difference (N - Y):", round(float(wr_N - wr_Y), 4))

plt.figure(figsize=(5,4))
sns.barplot(x=["Non-neutral (N)", "Neutral (Y)"], y=[wr_N, wr_Y])
plt.title("Home win rate: Neutral vs Non-neutral")
plt.ylim(0, 1)
plt.tight_layout()
plt.show()

# --- 2) Margin: Neutral vs Non-neutral ---
m_N = g[g["neutral2"] == "N"]["margin"].values
m_Y = g[g["neutral2"] == "Y"]["margin"].values

d_margin = cohens_d(m_N, m_Y)
t_stat, p_val = ttest_report(m_N, m_Y)

print("\nMargin (mean) Non-neutral:", round(float(np.nanmean(m_N)), 3))
print("Margin (mean) Neutral:", round(float(np.nanmean(m_Y)), 3))
print("Mean diff (N - Y):", round(float(np.nanmean(m_N) - np.nanmean(m_Y)), 3))
print("Cohen's d (N - Y):", round(float(d_margin), 3))
print("Welch t-test p-value:", round(float(p_val), 3))

plt.figure(figsize=(8,4))
sns.kdeplot(m_N, label="Non-neutral (N)", fill=True)
sns.kdeplot(m_Y, label="Neutral (Y)", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: Neutral vs Non-neutral")
plt.legend()
plt.tight_layout()

```

```

plt.show()

# --- 3) Margin: OT vs No OT ---
m_no = g[g["ot_flag"] == "N"]["margin"].values
m_ot = g[g["ot_flag"] == "Y"]["margin"].values


d_ot = cohens_d(m_no, m_ot)
t_stat2, p_val2 = ttest_report(m_no, m_ot)

print("\nMargin (mean) No OT:", round(float(np.nanmean(m_no)), 3))
print("Margin (mean) OT: ", round(float(np.nanmean(m_ot)), 3))
print("Mean diff (NoOT - OT):", round(float(np.nanmean(m_no) - np.nanmean(m_
print("Cohen's d (NoOT - OT):", round(float(d_ot), 3))
print("Welch t-test p-value: ", p_val2)


plt.figure(figsize=(8,4))
sns.kdeplot(m_no, label="No OT", fill=True)
sns.kdeplot(m_ot, label="OT", fill=True)
plt.axvline(0, color="red")
plt.title("Margin distribution: OT vs No OT")
plt.legend()
plt.tight_layout()
plt.show()

```


Home win rate (Non-neutral N): 0.6201
Home win rate (Neutral Y): 0.628
Difference (N - Y): -0.0079



Margin (mean) Non-neutral: 6.208
Margin (mean) Neutral: 4.79
Mean diff (N - Y): 1.418
Cohen's d (N - Y): 0.079
Welch t-test p-value: 1.000811441002832e-11



Margin (mean) No OT: 6.106
 Margin (mean) OT: -0.181
 Mean diff (NoOT - OT): 6.286
 Cohen's d (NoOT - OT): 0.446
 Welch t-test p-value: 1.426992426508403e-64



Step 6 — Inspect Feature Tables (team_stats, elo, teams)

Before merging anything, we need to understand:

- What is the primary key?
- Is there season?
- Are stats per game or aggregated?
- Does elo contain pregame ratings?

We inspect structure + sample rows.

```

In [32]: # =====
# Step 6 – Inspect feature tables (TRULY IOPub-safe)
# =====

def inspect_df_safe(name, df, top_missing=15, preview_rows=2, preview_cols=10):
    print("\n" + "="*90)
    print(f"{name}")
    print("Shape:", df.shape)
    print("Num columns:", df.shape[1])

    # ----- Column printing (capped) -----
    print("\nColumns:")
    total_cols = df.shape[1]

    if total_cols <= max_cols_print:
        for c in df.columns:
            print(" ", c)
    else:
        print(f" Showing first {max_cols_print} of {total_cols} columns:")
        for c in df.columns[:max_cols_print]:
            print(" ", c)
        print(" ... (truncated)")

    # ----- Dtypes -----
    print("\nDtype counts:")
    print(df.dtypes.value_counts())

    # ----- Missingness -----
    miss = df.isna().mean().sort_values(ascending=False).head(top_missing)
    miss = miss[miss > 0]

    print(f"\nTop missingness (fraction, up to {top_missing}):")
    if len(miss) == 0:
        print(" (no missing values detected)")
    else:
        for k, v in miss.items():
            print(f" {k}: {v:.3f}")

    # ----- Safe Preview -----
    print(f"\nPreview: first {preview_rows} rows, first {preview_cols} columns")
    preview = df.iloc[:preview_rows, :min(preview_cols, total_cols)]
    display(preview)

# Run inspection
inspect_df_safe("teams", teams)
inspect_df_safe("team_stats", team_stats)
inspect_df_safe("elo", elo)

```

=====

=====
teams
Shape: (366, 6)
Num columns: 6

Columns:
 message
 detail
 season
 code
 name
 location

Dtype counts:
object 5
float64 1
Name: count, dtype: int64

Top missingness (fraction, up to 15):
 message: 1.000
 detail: 1.000
 season: 1.000
 location: 0.962

Preview: first 2 rows, first 12 columns

	message	detail	season	code	name	location
0	None	None	NaN	AFA	Air Force Falcons	None
1	None	None	NaN	AKR	Akron Zips	None

```
=====
=====
team_stats
Shape: (5297, 30)
Num columns: 30
```

Columns:

team_name
team_code
2fa
2fm
3fa
3fm
apg
ast
blk
bpg
f
fga
fgm
fpg
fta
ftm
g
l
min
oreb
pts
reb
rpg
spg
stl
to
topg
w
winpct
season

Dtype counts:

float64	26
object	2
int64	1
Int64	1

Name: count, dtype: int64

Top missingness (fraction, up to 15):

```
team_name: 0.003
team_code: 0.003
```

```
Preview: first 2 rows, first 12 columns
```

[illegible]

```
=====
=====
elo
Shape: (4937, 5)
Num columns: 5
```

```
Columns:
  team
  code
  elo
  elorank
  season
```

```
Dtype counts:
object      4
Int64       1
Name: count, dtype: int64
```

```
Top missingness (fraction, up to 15):
(no missing values detected)
```

```
Preview: first 2 rows, first 12 columns
```

	team	code	elo	elorank	season
0	Connecticut	CONN	2506	1	2013
1	South Carolina	SC	2435	2	2013

```
In [33]: # =====
# Step 7 - Clean ELO table (SEASON-AWARE)
# =====

print("Original elo shape:", elo.shape)

elo_keep = ["code", "elo", "elorank"]
if "season" in elo.columns:
    elo_keep.append("season")

elo_clean = elo[elo_keep].copy()

elo_clean["elo"] = pd.to_numeric(elo_clean["elo"], errors="coerce")
elo_clean["elorank"] = pd.to_numeric(elo_clean["elorank"], errors="coerce")
if "season" in elo_clean.columns:
    elo_clean["season"] = pd.to_numeric(elo_clean["season"], errors="coerce")

print("Cleaned elo shape:", elo_clean.shape)
display(elo_clean.head())
```

```
Original elo shape: (4937, 5)
Cleaned elo shape: (4937, 4)
```

	code	elo	elrank	season
0	CONN	2506	1	2013
1	SC	2435	2	2013
2	UCLA	2348	3	2013
3	TEX	2305	4	2013
4	LSU	2209	5	2013

```
In [34]: # =====
# Step 8 – Clean team_stats numerics
# =====

team_stats_clean = team_stats.copy()

# Keep useful columns and season key for multi-year joins
keep_cols = [
    "team_code",
    "season",
    "o_ppp", "d_ppp",
    "pace",
    "pts", "g", "min",
    "reb", "reb_d",
    "oreb", "dreb",
    "to", "to_d",
    "fgm", "fga",
    "ftm", "fta",
    "winpct"
]

keep_cols = [c for c in keep_cols if c in team_stats_clean.columns]
team_stats_clean = team_stats_clean[keep_cols]

for c in team_stats_clean.columns:
    if c != "team_code":
        team_stats_clean[c] = pd.to_numeric(team_stats_clean[c], errors="coerce")

# Possession estimate from box-score totals
if {"fga", "fta", "to", "oreb"}.issubset(set(team_stats_clean.columns)):
    team_stats_clean["poss_est"] = (
        team_stats_clean["fga"]
        + 0.44 * team_stats_clean["fta"]
        + team_stats_clean["to"]
        - team_stats_clean["oreb"]
    )

# Fallback offensive PPP
if "o_ppp" not in team_stats_clean.columns:
    team_stats_clean["o_ppp"] = np.nan
if {"pts", "poss_est"}.issubset(set(team_stats_clean.columns)):
    o_ppp_est = pd.Series(np.where(team_stats_clean["poss_est"] > 0, team_stats_clean["pts"] / team_stats_clean["poss_est"], np.nan))
    team_stats_clean["o_ppp"] = team_stats_clean["o_ppp"].fillna(o_ppp_est)
```

```

# Fallback pace (possessions per game)
if "pace" not in team_stats_clean.columns:
    team_stats_clean["pace"] = np.nan
if {"poss_est", "g"}.issubset(set(team_stats_clean.columns)):
    pace_est = pd.Series(np.where(team_stats_clean["g"] > 0, team_stats_clean["pace"], np.nan))
    team_stats_clean["pace"] = team_stats_clean["pace"].fillna(pace_est)

print("Cleaned team_stats shape:", team_stats_clean.shape)
for c in ["o_ppp", "pace", "winpct"]:
    if c in team_stats_clean.columns:
        print(f"nonnull {c}:", int(team_stats_clean[c].notna().sum()))
display(team_stats_clean.head())

```

Cleaned team_stats shape: (5297, 16)

nonnull o_ppp: 4936

nonnull pace: 5297

nonnull winpct: 5297

	team_code	season	pts	g	min	reb	oreb	to	fgm	fga	ftm	fta	winpct	poss.
0	AFA	2012	0.0	20	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.2500	
1	AKR	2012	0.0	22	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.4091	
2	ALA	2012	0.0	21	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.4762	
3	ALAM	2012	0.0	20	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.5500	
4	ALB	2012	0.0	25	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.7200	

```

In [35]: # =====
# Step 9 – Build modeling table (SEASON-AWARE JOINS)
# =====

model_df = games.copy()

if "season" not in model_df.columns:
    raise ValueError("games table must include a season column for all-season")

# --- Merge home stats by team + season ---
home_stats = team_stats_clean.rename(columns={"team_code": "home_code"}).copy()
model_df = model_df.merge(home_stats, on=["home_code", "season"])
for col in team_stats_clean.columns:
    if col not in ["team_code", "season"] and col in model_df.columns:
        model_df.rename(columns={col: f"{col}_home"}, inplace=True)

# --- Merge visitor stats by team + season ---
vis_stats = team_stats_clean.rename(columns={"team_code": "visitor_code"}).copy()
model_df = model_df.merge(vis_stats, on=["visitor_code", "season"])
for col in team_stats_clean.columns:
    if col not in ["team_code", "season"] and col in model_df.columns and f"{col}_home" not in model_df.columns:
        model_df.rename(columns={col: f"{col}_vis"}, inplace=True)

# --- Merge home ELO by team + season ---
home_elo = elo_clean.rename(columns={"code": "home_code"}).copy()
home_elo_cols = [c for c in ["home_code", "elo", "elorank", "season"] if c in home_elo.columns]
model_df = model_df.merge(home_elo[home_elo_cols], on=["home_code", "season"])

```

```

if "elo" in model_df.columns:
    model_df.rename(columns={"elo": "elo_home"}, inplace=True)
if "elrank" in model_df.columns:
    model_df.drop(columns=["elrank"], inplace=True)

# --- Merge visitor ELO by team + season ---
vis_elo = elo_clean.rename(columns={"code": "visitor_code"}).copy()
vis_elo_cols = [c for c in ["visitor_code", "elo", "elrank", "season"] if c in vis_elo.columns]
model_df = model_df.merge(vis_elo[vis_elo_cols], on=[c for c in ["visitor_code", "season"] if c in model_df.columns])
if "elo" in model_df.columns:
    model_df.rename(columns={"elo": "elo_vis"}, inplace=True)
if "elrank" in model_df.columns:
    model_df.drop(columns=["elrank"], inplace=True)

print("Final modeling df shape:", model_df.shape)
display(model_df.head())

```

Final modeling df shape: (75550, 48)

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	631025	Notre Dame Fightin` Irish	ND	61	Baylor Lady Bears	BAY	80
1	631024	Stanford Cardinal	STAN	47	Baylor Lady Bears	BAY	59
2	631023	Connecticut Huskies	CONN	75	Notre Dame Fightin` Irish	ND	83
3	631022	James Madison Dukes	JMU	68	Oklahoma State Cowboys	OKSU	75
4	631020	Kentucky Wildcats	UK	65	Connecticut Huskies	CONN	80

5 rows x 48 columns

```

In [36]: # =====
# Step 10 - Deduplicate games properly
# =====

print("Rows before dedupe:", model_df.shape[0])

if "gamestatus" in model_df.columns:
    gs = model_df["gamestatus"].astype(str).str.strip().str.lower()
    model_df = model_df[gs.str.startswith("final")].copy()

if "id" in model_df.columns:
    model_df = model_df.drop_duplicates(subset=["id"])

print("Rows after dedupe:", model_df.shape[0])
display(model_df.head())

```


Rows before dedupe: 75550

Rows after dedupe: 75425

	id	visitor	visitor_code	score_vis	home	home_code	score_home
0	631025	Notre Dame Fightin` Irish	ND	61	Baylor Lady Bears	BAY	80
1	631024	Stanford Cardinal	STAN	47	Baylor Lady Bears	BAY	59
2	631023	Connecticut Huskies	CONN	75	Notre Dame Fightin` Irish	ND	83
3	631022	James Madison Dukes	JMU	68	Oklahoma State Cowboys	OKSU	75
4	631020	Kentucky Wildcats	UK	65	Connecticut Huskies	CONN	80

5 rows x 48 columns

```
In [37]: # =====
# Step 11 – Feature engineering
# =====

model_df["home_win"] = (model_df["winner_code"] == model_df["home_code"]).as

for c in ["elo_home", "elo_vis", "o_ppp_home", "o_ppp_vis", "pace_home", "pa
    if c not in model_df.columns:
        model_df[c] = np.nan

model_df["elo_diff"] = model_df["elo_home"] - model_df["elo_vis"]
model_df["ppp_diff"] = model_df["o_ppp_home"] - model_df["o_ppp_vis"]
model_df["pace_diff"] = model_df["pace_home"] - model_df["pace_vis"]
model_df["winpct_diff"] = model_df["winpct_home"] - model_df["winpct_vis"]

# Extra EDA-friendly derived features
model_df["abs_elo_diff"] = model_df["elo_diff"].abs()
model_df["abs_ppp_diff"] = model_df["ppp_diff"].abs()
model_df["abs_pace_diff"] = model_df["pace_diff"].abs()
model_df["abs_winpct_diff"] = model_df["winpct_diff"].abs()
model_df["home_favored"] = (model_df["elo_diff"] > 0).astype("Int64")
model_df["upset_flag"] = (model_df["home_win"] != model_df["home_favored"]).

if {"score_home", "score_vis"}.issubset(set(model_df.columns)):
    sh = pd.to_numeric(model_df["score_home"], errors="coerce")
    sv = pd.to_numeric(model_df["score_vis"], errors="coerce")
    model_df["game_total"] = sh + sv
    model_df["margin_abs"] = (sh - sv).abs()

# Feature catalog for downstream relationship plots
extended_feature_candidates = [
    "elo_diff", "ppp_diff", "pace_diff", "winpct_diff",
```

```

    "abs_elo_diff", "abs_ppp_diff", "abs_pace_diff", "abs_winpct_diff",
    "home_favored", "upset_flag",
    "elo_home", "elo_vis",
    "o_ppp_home", "o_ppp_vis",
    "pace_home", "pace_vis",
    "winpct_home", "winpct_vis",
    "game_total", "margin_abs"
]
extended_feature_cols = []
for c in extended_feature_candidates:
    if c in model_df.columns and pd.api.types.is_numeric_dtype(model_df[c]):
        if model_df[c].notna().sum() > 100:
            extended_feature_cols.append(c)

display(model_df[[
    "home_win",
    "elo_diff",
    "ppp_diff",
    "pace_diff",
    "winpct_diff",
    "abs_elo_diff",
    "upset_flag"
]].head())
print("Extended feature count:", len(extended_feature_cols))
print("Sample features:", extended_feature_cols[:12])

```

	home_win	elo_diff	ppp_diff	pace_diff	winpct_diff	abs_elo_diff	upset_flag
0	1	NaN	NaN	0.0	0.1528	NaN	1
1	1	NaN	NaN	0.0	0.1217	NaN	1
2	1	NaN	NaN	0.0	0.0222	NaN	1
3	1	NaN	NaN	0.0	0.0368	NaN	1
4	1	NaN	NaN	0.0	0.0750	NaN	1

Extended feature count: 20

Sample features: ['elo_diff', 'ppp_diff', 'pace_diff', 'winpct_diff', 'abs_elo_diff', 'abs_ppp_diff', 'abs_pace_diff', 'abs_winpct_diff', 'home_favored', 'upset_flag', 'elo_home', 'elo_vis']

```

In [38]: # =====
# Step 12 – Feature Distribution & Separation (Pure EDA)
# =====

if "extended_feature_cols" in locals() and len(extended_feature_cols) > 0:
    candidate_cols = extended_feature_cols
else:
    candidate_cols = ["elo_diff", "ppp_diff", "pace_diff", "winpct_diff"]

usable_cols = []
for col in candidate_cols:
    if col not in model_df.columns:
        continue
    tmp = model_df[[col, "home_win"]].dropna()

```

```

    if tmp.empty or tmp["home_win"].nunique() < 2:
        continue
    usable_cols.append(col)

# Rank by separability (absolute mean gap) and keep a manageable set for plotting
rank_rows = []
for col in usable_cols:
    t = model_df[[col, "home_win"]].dropna()
    gap = t.loc[t.home_win==1, col].mean() - t.loc[t.home_win==0, col].mean()
    rank_rows.append((col, abs(gap)))
ranked = [c for c, _ in sorted(rank_rows, key=lambda x: x[1], reverse=True)]
plot_cols = ranked[:12]

for col in plot_cols:
    plt.figure(figsize=(8,4))
    sns.histplot(data=model_df, x=col, hue="home_win", bins=40, element="step")
    plt.title(f"{col} Distribution by Home Win")
    plt.show()

for col in plot_cols[:8]:
    plt.figure(figsize=(8,4))
    sns.kdeplot(data=model_df, x=col, hue="home_win", fill=True, common_norm=False)
    plt.title(f"{col} KDE Separation")
    plt.show()

for col in plot_cols:
    plot_df = model_df[["home_win", col]].dropna()
    if plot_df.empty:
        continue
    plt.figure(figsize=(6,4))
    sns.boxplot(data=plot_df, x="home_win", y=col)
    plt.title(f"{col} vs Outcome")
    plt.show()



print("\n=== Statistical Separation ===")
for col in plot_cols:
    wins = model_df.loc[model_df["home_win"] == 1, col].dropna()
    losses = model_df.loc[model_df["home_win"] == 0, col].dropna()
    if len(wins) < 2 or len(losses) < 2:
        continue
    mean_diff = wins.mean() - losses.mean()
    t_stat, p_val = ttest_ind(wins, losses, equal_var=False)
    pooled_std = np.sqrt((wins.var() + losses.var()) / 2)
    cohens_d = mean_diff / pooled_std if pooled_std and not np.isnan(pooled_std) else 0
    print(f"{col}: mean_diff={mean_diff:.4f}, d={cohens_d:.3f}, p={p_val:.3e}")



corr_cols = plot_cols + ["home_win"]
if len(corr_cols) >= 3:
    plt.figure(figsize=(10,8))
    corr = model_df[corr_cols].corr()
    sns.heatmap(corr, annot=False, cmap="coolwarm", center=0)
    plt.title("Feature Correlation Matrix (Top Features)")
    plt.show()

if plot_cols:
    relationship_feature_cols = plot_cols


```

```
else:  
    relationship_feature_cols = []  
    print("Relationship feature set size:", len(relationship_feature_cols))
```







margin_abs Distribution by Home Win




pace_diff Distribution by Home Win







pace_vis Distribution by Home Win







abs_pace_diff Distribution by Home Win











































```

=== Statistical Separation ===
elo_diff: mean_diff=291.7550, d=0.754, p=0.000e+00
elo_home: mean_diff=148.3545, d=0.404, p=0.000e+00
elo_vis: mean_diff=-129.1014, d=-0.357, p=0.000e+00
abs_elo_diff: mean_diff=39.6062, d=0.156, p=9.441e-92
margin_abs: mean_diff=4.9617, d=0.418, p=0.000e+00
pace_diff: mean_diff=3.5150, d=0.279, p=0.000e+00
pace_vis: mean_diff=-3.3572, d=-0.178, p=6.394e-134
abs_pace_diff: mean_diff=3.0498, d=0.265, p=0.000e+00
game_total: mean_diff=1.0624, d=0.054, p=4.239e-13
winpct_diff: mean_diff=0.3520, d=1.528, p=0.000e+00
home_favored: mean_diff=0.2400, d=0.496, p=0.000e+00
winpct_vis: mean_diff=-0.1802, d=-0.930, p=0.000e+00

```



Relationship feature set size: 12

Visual Feature Separation — Women's (Verbose)

Elo difference appears visually stronger in women's basketball than in men's. The overlap between winners and losers is reduced. Positive Elo_diff aligns more consistently with winning outcomes, and the density polarization is sharper. This suggests that rating gaps translate more directly into outcomes.

Win percentage difference also appears structurally reinforcing. High win% programs align with wins more consistently, and the distribution overlap appears smaller. Visually, win% behaves less like a noisy proxy and more like a stable stratification indicator.

PPP_diff remains relevant, but visually it looks more secondary relative to Elo. That does not mean it is unimportant; rather, the main separation already exists at the rating level, so efficiency differences may explain margin magnitude or edge cases more than the overall win probability structure.

Pace difference again shows limited separability. Tempo does not visually define outcomes at the league level.

Overall, women's basketball visually resembles a cleaner predictive environment: fewer mixed-density regions, stronger rating alignment, and less overlap. That is exactly what a hierarchy-driven league looks like when translated into feature distributions.

```
In [39]: # =====  
# Step 13 – Feature Interactions & Nonlinearity  
# =====  
  
if "relationship_feature_cols" in locals() and relationship_feature_cols:  
    rel_cols = relationship_feature_cols[:8]  
else:  
    rel_cols = ["elo_diff", "winpct_diff", "ppp_diff", "pace_diff"]  
  
# Pairwise comparisons among selected features  
pairs = []  
for i in range(len(rel_cols)):  
    for j in range(i+1, len(rel_cols)):  
        pairs.append((rel_cols[i], rel_cols[j]))  
  
usable_pairs = []  
for x_col, y_col in pairs:  
    if x_col not in model_df.columns or y_col not in model_df.columns:  
        continue  
    pair_df = model_df[[x_col, y_col, "home_win"]].dropna()  
    if pair_df.empty or pair_df["home_win"].nunique() < 2:  
        continue  
    usable_pairs.append((x_col, y_col))  
  
# Scatter by class (0/1)  
for x_col, y_col in usable_pairs[:12]:  
    plt.figure(figsize=(7,6))  
    sns.scatterplot(data=model_df, x=x_col, y=y_col, hue="home_win", alpha=0.5)  
    plt.title(f"{x_col} vs {y_col} (home_win 0/1)")  
    plt.show()  
  
# 2D KDE by class where density is stable  
for x_col, y_col in usable_pairs[:8]:  
    pair_df = model_df[[x_col, y_col, "home_win"]].dropna()  
    if len(pair_df) < 200:  
        continue
```

```

plt.figure(figsize=(7,6))
sns.kdeplot(data=pair_df, x=x_col, y=y_col, hue="home_win", fill=True, t
plt.title(f"Density: {x_col} × {y_col}")
plt.show()

# Pairplot matrix for strongest 5 features vs outcome
pairplot_cols = rel_cols[:5]
pairplot_df = model_df[pairplot_cols + ["home_win"]].dropna()
if not pairplot_df.empty:
    if len(pairplot_df) > 4000:
        pairplot_df = pairplot_df.sample(4000, random_state=42)
    sns.pairplot(pairplot_df, vars=pairplot_cols, hue="home_win", corner=Tru
plt.show()

# Monotonicity / win-rate by bins for many features
def win_rate_by_bins(col, bins=10):
    if col not in model_df.columns:
        return
    df = model_df[[col, "home_win"]].dropna().copy()
    if df.empty:
        return
    try:
        df["bin"] = pd.qcut(df[col], q=bins, duplicates="drop")
    except ValueError:
        return
    summary = df.groupby("bin")["home_win"].mean()
    if summary.empty:
        return
    plt.figure(figsize=(8,4))
    summary.plot(marker="o")
    plt.xticks(rotation=45)
    plt.title(f"Home Win Rate by {col} Quantile Bins")
    plt.ylabel("Home Win Rate")
    plt.show()

for col in rel_cols:
    win_rate_by_bins(col, bins=10)


print("Interaction pairs evaluated:", len(usable_pairs))


```

elo_diff vs elo_home (home_win 0/1)



elo_diff vs elo_vis (home_win 0/1)






elo_diff vs margin_abs (home_win 0/1)



elo_diff vs pace_diff (home_win 0/1)







elo_diff vs abs_pace_diff (home_win 0/1)




elo_home vs elo_vis (home_win 0/1)




elo_home vs abs_elo_diff (home_win 0/1)




elo_home vs margin_abs (home_win 0/1)





elo_home vs pace_diff (home_win 0/1)




elo_home vs pace_vis (home_win 0/1)














Density: elo_diff × margin_abs






























Interaction pairs evaluated: 28

```
In [40]: # =====
# Step 14 – Continuous Margin Analysis (FULL SAFE VERSION)
# =====

# ---- Ensure numeric score columns ----
model_df["score_home"] = pd.to_numeric(model_df["score_home"], errors="coerce")
model_df["score_vis"] = pd.to_numeric(model_df["score_vis"], errors="coerce")
```

```

# ---- Create margin if missing ----
if "margin" not in model_df.columns:
    model_df["margin"] = model_df["score_home"] - model_df["score_vis"]

print("Margin summary:")
print(model_df["margin"].describe())

# -----
# 1 Margin Distribution
# -----
plt.figure(figsize=(8,5))
sns.histplot(model_df["margin"].dropna(), bins=60, kde=True)
plt.title("Margin Distribution")
plt.show()

# -----
# 2 Margin vs Diff Features
# -----

diff_features = ["elo_diff", "ppp_diff", "winpct_diff", "pace_diff"]

for col in diff_features:
    if col in model_df.columns:
        plt.figure(figsize=(7,5))
        sns.regplot(
            data=model_df,
            x=col,
            y="margin",
            scatter_kws={"alpha":0.3},
            line_kws={"color":"red"}
        )
        plt.title(f"Margin vs {col}")
        plt.show()

# -----
# 3 Margin by Outcome
# -----

plt.figure(figsize=(6,5))
sns.boxplot(data=model_df, x="home_win", y="margin")
plt.title("Margin by Outcome")
plt.show()

# -----
# 4 Margin Skew & Tail Diagnostics
# -----

skew = model_df["margin"].skew()
kurt = model_df["margin"].kurtosis()

print(f"Margin Skew: {skew:.3f}")
print(f"Margin Kurtosis: {kurt:.3f}")

# Extreme margins
extreme_games = model_df[np.abs(model_df["margin"]) > 30]

```




```
print("Extreme margin games (>30 pts):", len(extreme_games))

plt.figure(figsize=(8,5))
sns.histplot(extreme_games["margin"], bins=40, kde=True)
plt.title("Extreme Margin Distribution (>30)")
plt.show()
```


Margin summary:

count	75425.000000
mean	6.074644
std	19.023045
min	-86.000000
25%	-7.000000
50%	6.000000
75%	17.000000
max	141.000000


Name: margin, dtype: float64






Margin vs elo_diff



Margin vs ppp_diff








Margin Skew: 0.421

Margin Kurtosis: 1.099

Extreme margin games (>30 pts): 8513



```

In [41]: # =====
# Step 15 – PCA Structure Analysis (Full Safe)
# =====

pca_features = ["elo_diff", "ppp_diff", "pace_diff", "winpct_diff"]
pca_df = model_df[[c for c in pca_features if c in model_df.columns] + ["home_win"]]
print("PCA sample size:", len(pca_df))

if len(pca_df) > 1 and len([c for c in pca_features if c in pca_df.columns]):
    use_cols = [c for c in pca_features if c in pca_df.columns]
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(pca_df[use_cols])
    pca = PCA(n_components=2)
    X_pca = pca.fit_transform(X_scaled)
    print("Explained variance ratio:", pca.explained_variance_ratio_)


    plt.figure(figsize=(8,6))
    sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=pca_df["home_win"], alpha=0.5)
    plt.title("PCA Projection (Colored by Home Win)")
    plt.show()

    loadings = pd.DataFrame(pca.components_.T, index=use_cols, columns=["PC1", "PC2"])
    print("\nPCA Loadings:")
    display(loadings)
else:
    print("Skipping PCA: insufficient complete rows/features after dropna.")

```

PCA sample size: 68666

Explained variance ratio: [0.53836393 0.2503112]



PCA Loadings:

	PC1	PC2
elo_diff	0.461663	-0.013148
ppp_diff	0.633893	0.033599
pace_diff	-0.003931	0.999183
winpct_diff	0.620510	-0.018212

```
In [42]: # =====
# Step 16 – Unsupervised Regime Detection
# =====

if "X_scaled" in locals() and "X_pca" in locals() and "pca_df" in locals():
    kmeans = KMeans(n_clusters=3, random_state=42)
    clusters = kmeans.fit_predict(X_scaled)
    print("Silhouette Score:", silhouette_score(X_scaled, clusters))

    plt.figure(figsize=(8,6))
    sns.scatterplot(x=X_pca[:,0], y=X_pca[:,1], hue=clusters, palette="Set2")
    plt.title("KMeans Clusters in PCA Space")
    plt.show()


    cluster_df = pca_df.copy()
    cluster_df["cluster"] = clusters
```

```

cluster_summary = cluster_df.groupby("cluster")["home_win"].mean()
print("\nWin rate by cluster:")
print(cluster_summary)
else:
    print("Skipping Step 16: PCA artifacts unavailable (insufficient data in

```

Silhouette Score: 0.21378548128883115



```

Win rate by cluster:
cluster
0    0.272823
1    0.650992
2    0.899261
Name: home_win, dtype: float64

```

Visual Dimensional Structure — PCA / Clustering (Verbose)

The dimensionality visuals reinforce the story seen in basic distributions.


For men's basketball, PCA and clustering plots tend to show blended structure. Even if there is directionality (stronger teams drifting into a region), boundaries look soft and mixed. That implies that the underlying feature space does not partition cleanly into a few distinct regimes. This is consistent with overlap in Elo and PPP separation and


supports the idea that outcomes are governed by multiple interacting factors with variance.

For women's basketball, clustering appears more distinct and stratified. Separation looks more like tiering rather than blending. This matches the dominance patterns and cleaner Elo separability. In practical terms, it suggests there are clearer "bands" of team strength and performance profiles.

This is still visual evidence, but it is consistent across structural, feature, and dimensional views: women's looks more tiered; men's looks more blended.

```
In [43]: # =====  
# Step 17 – 2D Conditional Win Probability Surface  
# =====  
  
def probability_surface(x_col, y_col, bins=12):  
    df = model_df[[x_col, y_col, "home_win"]].dropna().copy()  
  
    df["x_bin"] = pd.qcut(df[x_col], q=bins, duplicates="drop")  
    df["y_bin"] = pd.qcut(df[y_col], q=bins, duplicates="drop")  
  
    pivot = df.groupby(["x_bin", "y_bin"])["home_win"].mean().unstack()  
  
    plt.figure(figsize=(9,7))  
    sns.heatmap(pivot, cmap="coolwarm", center=0.5)  
    plt.title(f"Win Probability Surface: {x_col} × {y_col}")  
    plt.show()  
  
probability_surface("elo_diff", "winpct_diff")  
probability_surface("elo_diff", "ppp_diff")
```




```
In [44]: # =====
# Step 18 – Residual Home Advantage After Adjustment
# =====

# Adjust margin for ELO expectation (approx scaling)
model_df["elo_adjusted_margin"] = model_df["margin"] - model_df["elo_diff"] /

plt.figure(figsize=(8,5))
sns.histplot(model_df["elo_adjusted_margin"].dropna(), bins=60, kde=True)
plt.title("ELO-Adjusted Margin Distribution")
plt.show()

print("Mean adjusted margin:", model_df["elo_adjusted_margin"].mean())
```



Mean adjusted margin: 2.5133574112370023

Additional EDA Required for Tournament Modeling — WBB

```
In [45]: # Team-season aggregation (WBB)
_tg = games.copy()
for c in ["score_home", "score_vis", "season"]:
    if c in _tg.columns:
        _tg[c] = pd.to_numeric(_tg[c], errors="coerce")
if "gameday" in _tg.columns:
    _tg["gameday"] = pd.to_datetime(_tg["gameday"], errors="coerce")

# NCAA tournament proxy for exclusion from selection features:
# no explicit flag in source, so use null-league games in Mar 15-Apr 15 window
_tg["_mm_proxy"] = (
    _tg.get("league").isna()
    & _tg.get("gameday").notna()
    & (_tg["gameday"].dt.month == 3)
    & (_tg["gameday"].dt.day >= 15)
) | (
    _tg.get("league").isna()
    & _tg.get("gameday").notna()
    & (_tg["gameday"].dt.month == 4)
    & (_tg["gameday"].dt.day <= 15)
)
_tg_selection = _tg.loc[~_tg["_mm_proxy"]].copy()
print("WBB selection pool (non-March-Madness games):", _tg_selection.shape)

home = pd.DataFrame({
```

```

    "season": _tg_selection.get("season"),
    "gameday": _tg_selection.get("gameday"),
    "team_code": _tg_selection.get("home_code"),
    "opp_code": _tg_selection.get("visitor_code"),
    "conference": _tg_selection.get("league") if "league" in _tg_selection.c
    "win": (_tg_selection.get("winner_code") == _tg_selection.get("home_code
    "margin": _tg_selection.get("score_home") - _tg_selection.get("score_vis
})
vis = pd.DataFrame({
    "season": _tg_selection.get("season"),
    "gameday": _tg_selection.get("gameday"),
    "team_code": _tg_selection.get("visitor_code"),
    "opp_code": _tg_selection.get("home_code"),
    "conference": _tg_selection.get("league") if "league" in _tg_selection.c
    "win": (_tg_selection.get("winner_code") == _tg_selection.get("visitor_c
    "margin": _tg_selection.get("score_vis") - _tg_selection.get("score_home
})
team_games = pd.concat([home, vis], ignore_index=True).dropna(subset=["seaso
team_games = team_games.sort_values(["season", "team_code", "gameday"])
team_games["rolling_win_10"] = team_games.groupby(["season", "team_code"])["

selection_df = team_games.groupby(["season", "team_code"], as_index=False).a
    games_played=("win", "size"),
    overall_winpct=("win", "mean"),
    avg_margin=("margin", "mean"),
    conf_winpct=("win", "mean"),
    nonconf_winpct=("win", "mean"),
    last10_winpct=("rolling_win_10", "last"),
)

if "team_stats_clean" in globals() and {"season", "team_code"}.issubset(team
    _ts = team_stats_clean.copy()
    for c in ["o_ppp", "d_ppp"]:
        if c not in _ts.columns:
            _ts[c] = np.nan
    _ts["net_rating"] = _ts["o_ppp"] - _ts["d_ppp"]
    keep = [c for c in ["season", "team_code", "net_rating", "o_ppp", "d_ppp
    selection_df = selection_df.merge(_ts[keep], on=["season", "team_code"],

if "elo_clean" in globals() and {"season", "code", "elo"}.issubset(elo_clean
    e = elo_clean[["season", "code", "elo"]].copy().rename(columns={"code":
    e["end_elo"] = pd.to_numeric(e["end_elo"], errors="coerce")
    selection_df = selection_df.merge(e, on=["season", "team_code"], how="le

selection_df["elo_rank"] = selection_df.groupby("season")["end_elo"].rank(me
selection_df["selected_proxy"] = np.where(selection_df["elo_rank"].notna() &

# SOS proxy
opp_elo = selection_df[["season", "team_code", "end_elo"]].rename(columns={"
_sos = team_games.merge(opp_elo, on=["season", "opp_code"], how="left")
_sos = _sos.groupby(["season", "team_code"], as_index=False)["opp_end_elo"].
selection_df = selection_df.merge(_sos, on=["season", "team_code"], how="lef

# conference proxy: dominant league tag in season
conf = team_games.groupby(["season", "team_code", "conference"], as_index=Fa
conf = conf.drop_duplicates(["season", "team_code"]).drop(columns=['size']).

```

```
selection_df = selection_df.merge(conf, on=["season", "team_code"], how="left")
print("selection_df:", selection_df.shape)
display(selection_df.head())
```

WBB selection pool (non-March-Madness games): (73826, 19)
 selection_df: (5282, 17)

	season	team_code	games_played	overall_winpct	avg_margin	conf_winpct	nonconf_winpct
0	2012	AFA	20	0.250000	-13.150000	0.250000	0.250000
1	2012	AKR	24	0.416667	0.958333	0.416667	0.416667
2	2012	ALA	21	0.476190	-0.571429	0.476190	0.476190
3	2012	ALAM	20	0.550000	1.450000	0.550000	0.550000
4	2012	ALB	22	0.727273	7.090909	0.727273	0.727273

Structural Shift — From Game-Level to Team-Season Modeling

At this stage, the modeling problem changes fundamentally.

Game-level analysis focuses on matchup variance and conditional win probability.
 Tournament selection, however, is a team-season classification problem.

Each row now represents a team's full body of work for a given season. The relevant question is not "did this team win a specific game," but rather:

Does this season profile meet the structural threshold for tournament inclusion?

This aggregation reframes the predictive task. Selection modeling depends on:

- Sustained performance,
- Schedule-adjusted strength,
- Conference context,
- And potentially recency effects.

The remaining EDA focuses on identifying whether selected teams form a statistically separable cluster relative to non-selected teams.

```
In [46]: # Selected vs not-selected distribution plots (WBB)
metrics = ["end_elo", "overall_winpct", "net_rating", "avg_margin", "sos_pro"]
metrics = [m for m in metrics if m in selection_df.columns]
for m in metrics:
    d = selection_df[[m, "selected_proxy"]].dropna()
    if d.empty or d["selected_proxy"].nunique() < 2:
        continue
    plt.figure(figsize=(8,4))
```


```



sns.kdeplot(data=d, x=m, hue="selected_proxy", fill=True, common_norm=False)
plt.title(f"WBB Selected vs Non-Selected (proxy): {m}")
plt.show()



if "end_elo" in selection_df.columns:
    s=selection_df[selection_df["selected_proxy"]==1]["end_elo"].dropna()
    n=selection_df[selection_df["selected_proxy"]==0]["end_elo"].dropna()
    if len(s) and len(n):
        print("min selected elo:", float(s.min()), "max non-selected elo:",

if {"overall_winpct", "sos_proxy", "selected_proxy"}.issubset(selection_df.c
d=selection_df[["overall_winpct", "sos_proxy", "selected_proxy"]].dropna()
if not d.empty and d["selected_proxy"].nunique()>=2:
    plt.figure(figsize=(8,6))
    sns.scatterplot(data=d, x="overall_winpct", y="sos_proxy", hue="sele
    plt.title(f"WBB Win% vs SOS (selection proxy)")
    plt.show()

```







min selected elo: 1646.0 max non-selected elo: 1650.0

WBB Win% vs SOS (selection proxy)



Visual Separation — Selected vs Not Selected

The selected vs non-selected comparisons reveal how deterministic the committee threshold appears.

If the distributions show:

- Minimal overlap → selection behaves like a clear statistical cutoff.
- Substantial overlap → committee behavior introduces subjectivity or secondary criteria.



Elo distribution is particularly important. A clean lower bound among selected teams suggests a de facto rating threshold. Overlap in mid-range Elo values indicates a bubble zone where additional factors likely influence inclusion.


Win percentage alone may appear separated, but without schedule context it risks overstating weaker conference teams. If schedule-adjusted strength reduces overlap more than raw win%, that suggests the committee implicitly values strength-of-schedule.

The key visual question is not whether selected teams are stronger on average (they should be), but whether there exists a clear structural boundary separating them.

```
In [47]: # Bubble zone visuals (WBB)
if "elo_rank" in selection_df.columns:
    bubble = selection_df[(selection_df["elo_rank"]>=30) & (selection_df["elo_rank"]<80)]
    print("bubble rows:", len(bubble))
    b_metrics=[m for m in ["avg_margin", "last10_winpct", "overall_winpct", "net_rating"]]
    for m in b_metrics:
        plot_df=bubble[["selected_proxy", m]].dropna()
        if plot_df.empty or plot_df["selected_proxy"].nunique()<2:
            continue
        plt.figure(figsize=(7,4))
        sns.boxplot(data=plot_df, x="selected_proxy", y=m)
        plt.title(f"WBB Bubble Zone (Ranks 30-80): {m}")
        plt.show()
```

bubble rows: 714





Bubble Zone Analysis — Where Selection Is Decided

The bubble region provides the most informative selection signal.

Top teams are automatically included. Bottom teams are clearly excluded.
The structural complexity lies in the mid-tier.

Within this Elo and performance band:

- Do selected teams have consistently stronger net rating?
- Do they show stronger recent form?
- Is conference strength visibly different?
- Does margin of victory separate them?

If separation inside the bubble band is weak, then selection likely incorporates contextual or non-quantified factors.

If separation is clear even inside this constrained range, then selection behaves closer to a statistical threshold model.

This region should heavily influence feature design for the selection classifier.

```
In [48]: # Conference-level analysis (WBB)
if {"season", "conference_proxy", "selected_proxy"}.issubset(selection_df.columns):
    conf_bids = selection_df.groupby(["season", "conference_proxy"], as_index=False)
    conf_strength = selection_df.groupby("conference_proxy", as_index=False)
    display(conf_bids.head())
```

```


display(conf_strength.sort_values("avg_end_elo", ascending=False).head(2)

top = conf_bids.groupby("conference_proxy", as_index=False)["bids_proxy"]
plt.figure(figsize=(10,4))
sns.barplot(data=top, x="conference_proxy", y="bids_proxy")
plt.title(f"WBB Conference Bid Concentration (proxy)")
plt.xticks(rotation=45)
plt.show()

```

	season	conference_proxy	bids_proxy
0	2012	A-10	0
1	2012	A-EAST	0
2	2012	A-SUN	0
3	2012	ACC	0
4	2012	AMER	0

	conference_proxy	avg_end_elo
6	BIG12	1879.384615
24	SEC	1842.210000
5	BIG10	1777.108911
22	PAC-12	1748.250000
3	ACC	1743.364486
7	BIGEAST	1592.388158
4	AMER	1420.000000
31	WCC	1412.811594
0	A-10	1409.160428
14	IVY	1382.625000
11	C-USA	1337.587302
19	MWC	1301.625000
27	SUMMIT	1282.031746
16	MAC	1274.440476
28	SUNBELT	1221.886792
8	BIGSKY	1212.379747
12	CAA	1199.592857
18	MVC	1185.713287
21	OVC	1185.556213
13	HL	1183.724138



Conference Effects — Structural Influence on Selection

Conference distribution is not incidental; it is structurally relevant.

If certain conferences consistently receive multiple bids despite overlapping performance distributions, then conference membership is a predictive feature.

Key observations to evaluate visually:

- Concentration of bids within high-Elo conferences.
- Whether mid-major teams require statistically stronger profiles to be selected.
- Differences between MBB and WBB conference clustering.



If women's basketball shows stronger concentration among fewer dominant conferences, that aligns with earlier hierarchy findings.


If men's basketball shows broader conference representation and greater overlap, that reinforces variance and structural diversity.

Conference membership may function as an implicit strength-of-schedule proxy and should likely be encoded in the selection model.

```
In [49]: # Recency / trend visuals (WBB)
trend_cols=[c for c in ["last10_winpct","overall_winpct","avg_margin"] if c
for c in trend_cols:
    d=selection_df[[c, "selected_proxy"]].dropna()
    if d.empty or d["selected_proxy"].nunique()<2:
        continue
    plt.figure(figsize=(8,4))
    sns.kdeplot(data=d, x=c, hue="selected_proxy", fill=True, common_norm=Fa
```

```
plt.title(f"WBB Trend Distribution: {c}")  
plt.show()
```





Recency and Trend Effects

Committees frequently reference how teams are playing late in the season.

Visually, trend features should be evaluated for separation power:

- Late-season Elo change.
- Last 10 game win rate.
- Margin trend in final month.

If selected teams show stronger positive trend distributions relative to non-selected teams — particularly in the bubble region — recency may function as a secondary selection discriminator.

If trend distributions overlap heavily, then season-long structural strength dominates recency.

This distinction matters for model construction: trend features may be nonlinear modifiers rather than primary drivers.

```
In [50]: # Bracket add-ons: upset curve, conditional feature signal, calibration, margin
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import roc_auc_score, log_loss
from sklearn.calibration import calibration_curve

bdf=model_df.copy()
for c in ["elo_diff", "ppp_diff", "pace_diff", "winpct_diff", "score_home", "score_away"]:
```



```

    if c in bdf.columns:
        bdf[c]=pd.to_numeric(bdf[c], errors='coerce')
if "gameday" in bdf.columns:
    bdf["gameday"]=pd.to_datetime(bdf["gameday"], errors='coerce')

# NCAA tournament proxy: hold out for historical backtesting only.
bdf["_mm_proxy"] = (
    bdf.get("league").isna()
    & bdf.get("gameday").notna()
    & (bdf["gameday"].dt.month == 3)
    & (bdf["gameday"].dt.day >= 15)
) | (
    bdf.get("league").isna()
    & bdf.get("gameday").notna()
    & (bdf["gameday"].dt.month == 4)
    & (bdf["gameday"].dt.day <= 15)
)
train_bdf = bdf.loc[~bdf["_mm_proxy"]].copy()
backtest_bdf = bdf.loc[bdf["_mm_proxy"]].copy()
print("WBB bracket train rows (non-March-Madness):", len(train_bdf))
print("WBB March Madness holdout rows (backtest only):", len(backtest_bdf))

if {"elo_diff", "home_win"}.issubset(train_bdf.columns):
    bins=[0,25,50,75,100,10_000]
    labels=["0-25", "25-50", "50-75", "75-100", "100+"]
    train_bdf["elo_gap_bin"]=pd.cut(train_bdf["elo_diff"].abs(), bins=bins,
    train_bdf["home_favored"]=(train_bdf["elo_diff"]>0).astype(int)
    train_bdf["fav_won"]=np.where(train_bdf["home_favored"]==1, train_bdf["h
    train_bdf["upset"]=1-train_bdf["fav_won"]
    curve=train_bdf.groupby("elo_gap_bin", observed=False).agg(win_prob=("fa
    display(curve)
    plt.figure(figsize=(8,4))
    sns.lineplot(data=curve, x="elo_gap_bin", y="upset_rate", marker='o')
    plt.title(f"WBB Upset Rate by Elo Gap Bin (Regular Season / Conf Tournam
    plt.ylim(0,1)
    plt.show()

auc_rows=[]
for label in ["0-25", "25-50", "50-75", "75-100", "100+"]:
    sub=train_bdf[train_bdf.get("elo_gap_bin")==label]
    if sub is None or len(sub)<200:
        continue
    for feat in ["ppp_diff", "pace_diff", "winpct_diff"]:
        if feat not in sub.columns:
            continue
        d=sub[[feat, "home_win"]].dropna()
        if len(d)<120 or d["home_win"].nunique()<2:
            continue
        try:
            auc=max(roc_auc_score(d["home_win"], d[feat]), 1-roc_auc_score(c
        except Exception:
            auc=np.nan
        auc_rows.append({"elo_gap_bin":label, "feature":feat, "auc_abs":auc,
if auc_rows:
    display(pd.DataFrame(auc_rows))

```

```

cal_feats=[f for f in ["elo_diff","ppp_diff","pace_diff","winpct_diff"] if f
work=train_bdf[cal_feats+["home_win"]].dropna() if cal_feats and "home_win"
if len(work)>1000 and work["home_win"].nunique()==2:
    test=work.sample(frac=0.2, random_state=42)
    train=work.drop(test.index)
    elo_only=["elo_diff"] if "elo_diff" in cal_feats else [cal_feats[0]]

p1=Pipeline([("imp",SimpleImputer(strategy="median")),("sc",StandardScaler)])
p1.fit(train[elo_only], train["home_win"])
pr1=p1.predict_proba(test[elo_only])[:,1]

p2=Pipeline([("imp",SimpleImputer(strategy="median")),("sc",StandardScaler)])
p2.fit(train[cal_feats], train["home_win"])
pr2=p2.predict_proba(test[cal_feats])[:,1]

print("log_loss elo_only:", float(log_loss(test["home_win"], pr1)))
print("log_loss full:", float(log_loss(test["home_win"], pr2)))

pt1,pp1=calibration_curve(test["home_win"], pr1, n_bins=10)
pt2,pp2=calibration_curve(test["home_win"], pr2, n_bins=10)
plt.figure(figsize=(6,6))
plt.plot([0,1],[0,1], '--', color='gray')
plt.plot(pp1,pt1,marker='o',label='elo_only')
plt.plot(pp2,pt2,marker='o',label='full')
plt.title(f"WBB Calibration Curves (Train = Non-March-Madness)")
plt.legend(); plt.show()

if {"score_home","score_vis","elo_gap_bin"}.issubset(train_bdf.columns):
    train_bdf["margin"]=train_bdf["score_home"]-train_bdf["score_vis"]
    mv=train_bdf.groupby("elo_gap_bin", observed=False)["margin"].var().reset_index()
    display(mv)
    plt.figure(figsize=(8,4))
    sns.barplot(data=mv, x="elo_gap_bin", y="margin_variance")
    plt.title(f"WBB Margin Variance by Elo Gap Bin (Train Set)")
    plt.show()


# Persist league-specific datasets for cross-league model comparison
wbb_selection_df = selection_df.copy()
wbb_bracket_train_df = train_bdf.copy()
wbb_bracket_backtest_df = backtest_bdf.copy()

```

WBB bracket train rows (non-March-Madness): 73701

WBB March Madness holdout rows (backtest only): 1724

	elo_gap_bin	win_prob	upset_rate	n
0	0-25	0.499834	0.500166	3015
1	25-50	0.539323	0.460677	3484
2	50-75	0.529656	0.470344	3136
3	75-100	0.524786	0.475214	3268
4	100+	0.676142	0.323858	54107




	elo_gap_bin	feature	auc_abs	n
0	0-25	ppp_diff	0.756022	3015
1	0-25	pace_diff	0.505922	3015
2	0-25	winpct_diff	0.807149	3015
3	25-50	ppp_diff	0.758807	3484
4	25-50	pace_diff	0.506440	3484
5	25-50	winpct_diff	0.811031	3484
6	50-75	ppp_diff	0.774993	3136
7	50-75	pace_diff	0.505036	3136
8	50-75	winpct_diff	0.813649	3136
9	75-100	ppp_diff	0.782099	3268
10	75-100	pace_diff	0.530174	3268
11	75-100	winpct_diff	0.817195	3268
12	100+	ppp_diff	0.832660	54107
13	100+	pace_diff	0.508699	54107
14	100+	winpct_diff	0.864851	54107


log_loss elo_only: 0.60399896148428

log_loss full: 0.4430065929739351

WBB Calibration Curves (Train = Non-March-Madness)



	elo_gap_bin	margin_variance
0	0-25	212.503446
1	25-50	224.926709
2	50-75	231.578399
3	75-100	243.276955
4	100+	343.397028



Elo Gap vs Upset Probability

Elo gap binning directly tests hierarchy strength.

If win probability increases smoothly and steeply with Elo difference, the league is structurally predictable.

If upset frequency remains high even at moderate Elo gaps, variance remains a meaningful force.

Comparing MBB and WBB curves:

- A steeper WBB curve supports the earlier visual conclusion of stronger hierarchy.
- A flatter MBB curve supports the conclusion of higher conditional variance.

This analysis directly informs bracket simulation assumptions.

Conclusion — Men vs Women (Visual-Only)

Based strictly on visual comparisons, the leagues appear to rely on the same *types* of signals, but with different strength and clarity.

Men's basketball visually shows:

- a competitive center mass in margin outcomes,
- gradual dominance tiering rather than sharp stratification,
- substantial overlap in Elo and other feature distributions,

- and therefore a stronger role for variance (upsets remain visually plausible even at moderate rating gaps).

Women's basketball visually shows:

- a heavier tail of decisive margins,
- steeper dominance curves with stronger separation at the top,
- cleaner Elo alignment with outcomes and reduced overlap,
- and therefore a stronger role for hierarchy (rating gaps translate more directly into wins).

So the distinction is not that men and women rely on fundamentally different factors, but that the *same factors appear to operate with different intensity*. Women's data looks more structurally separable; men's data looks more overlap-heavy.

This is a qualitative conclusion grounded in repeated visual patterns across distribution, dominance, and feature separability plots.

Initial Model Decision Check — 2 Models vs 4 Models (AUC)

This section computes baseline ROC-AUC for both tasks using:

- **2-model strategy:** one pooled Selection model + one pooled Bracket model (with league indicator).
- **4-model strategy:** separate MBB and WBB models for Selection and Bracket.

Training logic follows your rule: March Madness games are excluded from model training; tournament games remain backtest-only.

```
In [51]: from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import roc_auc_score

def _fit_predict_auc(df, features, target='y', season_col='season'):
    df = df.copy()
    cols = [c for c in features if c in df.columns]
    if not cols or target not in df.columns:
        return None
    work = df[cols + [target, season_col]].dropna(subset=[target, season_col])
    work[target] = pd.to_numeric(work[target], errors='coerce')
    work = work.dropna(subset=[target])
    if len(work) < 300 or work[target].nunique() < 2:
        return None

    # Temporal holdout: use latest season(s) with both classes in test
```

```

seasons = sorted(work[season_col].dropna().unique())
test_mask = None
for k in range(1, min(4, len(seasons)) + 1):
    test_seasons = set(seasons[-k:])
    mask = work[season_col].isin(test_seasons)
    if work.loc[mask, target].nunique() == 2 and work.loc[~mask, target].nunique() == 2:
        test_mask = mask
        break
if test_mask is None:
    # Fallback to random split
    test_mask = work.sample(frac=0.2, random_state=42).index
    train = work.drop(index=test_mask)
    test = work.loc[test_mask]
else:
    train = work.loc[~test_mask]
    test = work.loc[test_mask]

if len(train) < 200 or len(test) < 80:
    return None

pipe = Pipeline([
    ('imp', SimpleImputer(strategy='median')),
    ('sc', StandardScaler()),
    ('lr', LogisticRegression(max_iter=2000))
])
pipe.fit(train[cols], train[target])
pred = pipe.predict_proba(test[cols])[:, 1]
auc = float(roc_auc_score(test[target], pred))
return {'auc': auc, 'n_train': int(len(train)), 'n_test': int(len(test))}

# ----- Build task datasets -----
sel_feats = ['end_elo', 'sos_proxy', 'net_rating', 'avg_margin', 'overall_wins']
brk_feats = ['elo_diff', 'ppp_diff', 'pace_diff', 'winpct_diff']

# Selection
mbb_sel = mbb_selection_df.copy()
mbb_sel['league_indicator'] = 0
mbb_sel['y'] = pd.to_numeric(mbb_sel.get('selected_proxy'), errors='coerce')

wbb_sel = wbb_selection_df.copy()
wbb_sel['league_indicator'] = 1
wbb_sel['y'] = pd.to_numeric(wbb_sel.get('selected_proxy'), errors='coerce')

pool_sel = pd.concat([mbb_sel, wbb_sel], ignore_index=True)

# Bracket (already non-March-Madness training sets)
mbb_brk = mbb_bracket_train_df.copy()
mbb_brk['league_indicator'] = 0
mbb_brk['y'] = pd.to_numeric(mbb_brk.get('home_win'), errors='coerce')

wbb_brk = wbb_bracket_train_df.copy()
wbb_brk['league_indicator'] = 1
wbb_brk['y'] = pd.to_numeric(wbb_brk.get('home_win'), errors='coerce')

pool_brk = pd.concat([mbb_brk, wbb_brk], ignore_index=True)

```

```

rows = []

# ----- 4-model strategy -----
for lg, df in [('MBB', mbb_sel), ('WBB', wbb_sel)]:
    r = _fit_predict_auc(df, sel_feats, target='y')
    if r:
        rows.append({'task': 'Selection', 'strategy': '4-model (separate)', 'league': 'CFL'})

for lg, df in [('MBB', mbb_brk), ('WBB', wbb_brk)]:
    r = _fit_predict_auc(df, brk_feats, target='y')
    if r:
        rows.append({'task': 'Bracket', 'strategy': '4-model (separate)', 'league': 'CFL'})

# ----- 2-model strategy (pooled by task, league indicator included) -----
r_sel_pool = _fit_predict_auc(pool_sel, sel_feats + ['league_indicator'], target='y')
if r_sel_pool:
    rows.append({'task': 'Selection', 'strategy': '2-model (pooled)', 'league': 'CFL'})
    test = r_sel_pool['test'].copy()
    test['_pred'] = r_sel_pool['pred']
    for lg, val in [('MBB', 0), ('WBB', 1)]:
        s = test[test['league_indicator']==val]
        if len(s) >= 80 and s['y'].nunique() == 2:
            rows.append({'task': 'Selection', 'strategy': '2-model (pooled eval)', 'league': 'CFL'})

r_brk_pool = _fit_predict_auc(pool_brk, brk_feats + ['league_indicator'], target='y')
if r_brk_pool:
    rows.append({'task': 'Bracket', 'strategy': '2-model (pooled)', 'league': 'CFL'})
    test = r_brk_pool['test'].copy()
    test['_pred'] = r_brk_pool['pred']
    for lg, val in [('MBB', 0), ('WBB', 1)]:
        s = test[test['league_indicator']==val]
        if len(s) >= 80 and s['y'].nunique() == 2:
            rows.append({'task': 'Bracket', 'strategy': '2-model (pooled eval)', 'league': 'CFL'})

auc_compare = pd.DataFrame(rows)
auc_compare = auc_compare.sort_values(['task', 'strategy', 'league']).reset_index()
display(auc_compare)

# ----- League-by-league deltas (comparison output only) -----
rec = []
for task in ['Selection', 'Bracket']:
    sep = auc_compare[(auc_compare.task==task) & (auc_compare.strategy=='4-model (separate)')]
    pol = auc_compare[(auc_compare.task==task) & (auc_compare.strategy=='2-model (pooled)')]
    if len(sep)==2 and len(pol)==2:
        m = sep.merge(pol, on=['task', 'league'], suffixes=('_sep', '_pool'))
        m['delta_sep_minus_pool'] = m['auc_sep'] - m['auc_pool']
        rec.append(m[['task', 'league', 'auc_sep', 'auc_pool', 'delta_sep_minus_pool']])

if rec:
    rec_df = pd.concat(rec, ignore_index=True)
    display(rec_df)
    display(rec_df.groupby('task', as_index=False)['delta_sep_minus_pool'].nunique())
else:
    print('Insufficient results for league-by-league pooled vs separate deltas')

```


	task	strategy	league	auc	n_train	n_test
0	Bracket	2-model (pooled eval by league)	MBB	0.848315	160962	5312
1	Bracket	2-model (pooled eval by league)	WBB	0.888243	160962	5116
2	Bracket	2-model (pooled)	Combined	0.870094	160962	10428
3	Bracket	4-model (separate)	MBB	0.848290	92377	5312
4	Bracket	4-model (separate)	WBB	0.888453	68585	5116
5	Selection	2-model (pooled eval by league)	MBB	1.000000	11238	364
6	Selection	2-model (pooled eval by league)	WBB	1.000000	11238	362
7	Selection	2-model (pooled)	Combined	0.999925	11238	726
8	Selection	4-model (separate)	MBB	0.999901	6318	364
9	Selection	4-model (separate)	WBB	1.000000	4920	362

	task	league	auc_sep	auc_pool	delta_sep_minus_pool
0	Selection	MBB	0.999901	1.000000	-0.000099
1	Selection	WBB	1.000000	1.000000	0.000000
2	Bracket	MBB	0.848290	0.848315	-0.000025
3	Bracket	WBB	0.888453	0.888243	0.000209

	task	mean_delta_sep_minus_pool
0	Bracket	0.000092
1	Selection	-0.000050

Model Performance Interpretation and Structural Implications

Bracket Modeling Results

Bracket AUC values are strong but realistic:

- MBB ≈ 0.848
- WBB ≈ 0.888
- Combined ≈ 0.870

The small deltas between pooled and separate training regimes (mean ≈ 0.000092) indicate that separating leagues does not materially change discriminative ability.

However, the absolute AUC levels differ between leagues, with WBB consistently higher. This aligns with earlier EDA suggesting stronger structural hierarchy in women's basketball.

Conclusion for bracket modeling:

- Performance is plausible.
- Structural league differences remain meaningful.
- Separate bracket models remain justified due to differing outcome variance and hierarchy strength.

Selection Modeling Results and Overfitting Considerations

Selection AUC values are effectively perfect:

- Pooled by league: 1.000 (MBB), 1.000 (WBB)
- Combined pooled: 0.999925
- Separate models: ~1.000

The deltas between pooled and separate models are negligible:

- Mean delta (selection): -0.000050

This indicates:

1. Separate modeling does not improve discriminative power.
2. Pooling does not degrade performance.
3. The underlying separation structure is nearly deterministic.

Given that separate models do not improve AUC and pooling slightly stabilizes performance, the most defensible modeling choice for selection is:

Use one pooled selection model across leagues.

Pooling increases effective sample size, reduces variance, and prevents unnecessary model fragmentation without sacrificing predictive performance.

The near-perfect AUC suggests that selection behaves like a strong threshold rule driven primarily by rating and strength metrics.

However, to avoid overfitting risk, the final selection model will:

- Be trained on pooled data.
- Be validated using strict temporal splits.
- Be evaluated within the bubble region.

- Avoid league-specific overparameterization.

Preventing Overfitting in Selection Modeling

The following adjustments should be implemented:

1. Remove or Test Without End-of-Season Elo

Run a model excluding final Elo to test whether separation remains perfect.

If AUC drops significantly, Elo is acting as a de facto deterministic threshold.

2. Remove Explicit or Implicit Bid Indicators

Ensure no feature directly encodes:

- Automatic bid status
- Conference tournament champion
- Post-selection rankings

3. Strict Temporal Validation

Train on seasons 1...N-1 Test only on season N

Never mix seasons randomly.

4. Add Noise Robustness Test

Shuffle labels within bubble zone and verify AUC collapses. If AUC remains high under label perturbation, leakage exists.

5. Evaluate Within Bubble Region Only

Compute AUC only for teams ranked ~30–80 by Elo.

If AUC remains near 1.000, then separation is unrealistically sharp.

6. Regularization Strength Increase

- Increase L1/L2 penalties.
- Limit tree depth.
- Reduce boosting rounds.

- Use early stopping.

7. Calibration Inspection

Perfect AUC does not guarantee calibrated probabilities. Plot calibration curve and compute Brier score.

A poorly calibrated model with AUC=1 is likely overfitting ranking structure.

Bracket Model — Interpretation of Results

Bracket AUC values:

- MBB ≈ 0.848
- WBB ≈ 0.888

These values are strong but not suspicious.

The higher WBB AUC aligns with earlier EDA suggesting stronger hierarchy and cleaner Elo separation.

The extremely small delta between pooled and separate models (≈ 0.00009 mean difference) indicates that modeling structure does not drastically change predictive capacity.

This supports the hypothesis that league structural differences are intrinsic rather than artifacts of model pooling.

Preventing Overfitting in Bracket Modeling

Even though AUC values are plausible, additional safeguards should be applied:

1. Use Log-Loss as Primary Metric

AUC measures ranking, not probability realism. Bracket simulation requires calibrated probabilities.

2. Calibration Correction

Apply:

- Platt scaling
- Isotonic regression

3. Elo Bin Validation

Evaluate performance separately for:

- Small Elo gaps
- Moderate gaps
- Large gaps

Overfitting often hides in small-gap games.

4. Limit Interaction Explosion

Avoid excessive feature cross-products unless validated.

5. Out-of-Season Testing Only

Bracket models must be validated on seasons not seen in training.

6. Monte Carlo Stability

Simulate tournament multiple times and measure outcome variance. Unstable simulation distributions may indicate probability miscalibration.

Final Modeling Architecture Decision

Based on empirical results:

Selection

Use **one pooled selection model** across MBB and WBB.

Rationale:

- No meaningful performance gain from separation.
- Slight stability advantage from pooling.
- Reduced overfitting risk.
- Selection appears governed by similar structural thresholds in both leagues.

Bracket

Use **separate bracket models** for MBB and WBB.

Rationale:

- Meaningful difference in absolute AUC (≈ 0.04).
- Earlier EDA showed stronger hierarchy in WBB.
- Outcome variance differs structurally.
- Separate calibration may be required.

Final structure:

- 1 Selection Model (pooled)
- 2 Bracket Models (MBB + WBB)

Total: 3 models, not 4.

Overfitting Mitigation Strategy

To ensure the pooled selection model is not artificially perfect:

1. Remove or test without end-of-season Elo.
2. Evaluate bubble-only AUC.
3. Apply strong regularization.
4. Use leave-one-season-out validation.
5. Inspect calibration (Brier score).

If AUC remains ≈ 1.0 under these stricter constraints, this supports the hypothesis that selection behaves like a deterministic rating threshold rather than a high-variance predictive task.

On Perfect AUC: When It Is and Is Not a Problem

An AUC of 1.000 is often treated as a red flag, but it is not inherently invalid.

AUC measures ranking quality — not probability realism. An AUC of 1.000 simply means that the model perfectly ranks all positive cases above all negative cases in the evaluation set.

Perfect separation can occur legitimately when:

1. The underlying process is threshold-driven.
2. The signal-to-noise ratio is extremely high.
3. The decision rule being modeled is deterministic.
4. The features capture the exact decision criteria.

In tournament selection, the committee often relies heavily on structured metrics such as rating, record, and strength of schedule. If these metrics already create a clean separation between selected and non-selected teams, then perfect ranking is mathematically possible.

Therefore, perfect AUC is not automatically evidence of overfitting.

However, perfect AUC becomes problematic when:

- It results from leakage (features contain post-label information).
- It collapses under temporal validation.
- It disappears when evaluated only in the bubble region.
- It produces poorly calibrated probabilities.

In this case, selection modeling may naturally behave like a strong threshold classifier. If strict temporal validation, bubble-only testing, and feature audits confirm stability, then near-perfect AUC reflects structural determinism rather than model pathology.

Importantly, AUC does not measure overfitting directly. Overfitting is defined by poor generalization to unseen data. If perfect AUC persists under held-out seasons and controlled ablations, then the model is not overfitting — it is capturing a highly separable decision boundary.

Thus, perfect AUC is not inherently bad. It must be evaluated in context.

Temporal Validation Strategy (Revised)

Temporal validation must reflect the structure of the data and the intended deployment setting.

1. Hard Holdout: 2026

The 2026 season is incomplete and must be fully excluded from model training and hyperparameter tuning.

2026 will function as:

- A final, untouched evaluation set.
- A proxy for real-world deployment.
- A sanity check against historical overfitting.

No feature engineering, scaling, or parameter tuning may use 2026 data.

2. Multi-Season Training / Validation Splits

Given approximately 20 seasons of historical data, leave-one-season-out validation is insufficient.

Instead, use rolling multi-season splits:

Example structure:

- Train: Seasons 2004–2015
Validate: 2016–2018
- Train: 2004–2018
Validate: 2019–2021
- Train: 2004–2021
Validate: 2022–2024

This ensures:

- Each validation window spans multiple seasons.
- Stability across different eras is tested.
- No single anomalous season dominates evaluation.

This also better reflects real forecasting conditions, where multiple prior seasons inform predictions.

3. Final Evaluation Structure

Model development process:

1. Train and tune using rolling multi-season splits.
 2. Select final model based on average validation performance.
 3. Lock hyperparameters.
 4. Train on all seasons up to 2025.
 5. Evaluate once on 2026.
 6. Do not retune after observing 2026 results.
-

4. Why Multi-Season Validation Is Necessary

With ~20 seasons:

- A single-season holdout is unstable.
- One unusual tournament year can distort metrics.
- Structural league shifts (e.g., parity changes) require multi-year robustness.

Multi-season validation ensures:

- Generalization across eras.
 - Reduced variance in performance estimates.
 - Stronger defense against overfitting.
-

5. Application to Both Tasks

This validation protocol applies to:

- Selection modeling (team-season level)
- Bracket modeling (game-level)

Selection models are particularly sensitive to threshold artifacts, and multi-season validation ensures that perfect AUC is not driven by a single deterministic era.

Bracket models require multi-season validation to ensure probability calibration holds over time.

Interpretation of Pooled vs Separate Deltas

Mean deltas:

- Bracket $\approx +0.000092$
- Selection ≈ -0.000050

These magnitudes are negligible.

This implies:

- Pooling leagues does not meaningfully change discrimination performance.
- Structural differences are reflected more in absolute AUC level than in pooled vs separate training regime.

However, given perfect AUC in selection, delta stability does not validate model realism. It may instead reflect a deterministic feature threshold common across leagues.

```
In [52]: # Modeling Freeze: datasets, targets, holdout policy, and source-of-truth co
from pathlib import Path
import json
import numpy as np
import pandas as pd

SEED = 42
HOLDOUT_SEASON = 2026
ARTIFACT_DIR = Path('artifacts_modeling')
ARTIFACT_DIR.mkdir(exist_ok=True)
```

```

SOURCE_CONFIG = {
    'seed': SEED,
    'holdout_season': HOLDOUT_SEASON,
    'leagues': ['MBB', 'WBB'],
    'files': {
        'MBB': {
            'games': 'games.parquet',
            'teams': 'teams.parquet',
            'team_stats': 'team_stats.parquet',
            'elo': 'elo.parquet',
        },
        'WBB': {
            'games': 'wbb_games.parquet',
            'teams': 'wbb_teams.parquet',
            'team_stats': 'wbb_team_stats.parquet',
            'elo': 'wbb_elo.parquet',
        },
    },
    'targets': {
        'bracket': 'home_win',
        'selection': 'selected_proxy',
    }
}

# Freeze dataset versions + target definitions
version_rows = []
for lg in SOURCE_CONFIG['leagues']:
    for k, p in SOURCE_CONFIG['files'][lg].items():
        fp = Path(p)
        version_rows.append({
            'league': lg,
            'dataset': k,
            'path': str(fp),
            'exists': fp.exists(),
            'size_bytes': fp.stat().st_size if fp.exists() else np.nan,
            'mtime_utc': pd.to_datetime(fp.stat().st_mtime, unit='s', utc=True)
        })
dataset_versions = pd.DataFrame(version_rows)
display(dataset_versions)

# Use modeling dataframes generated earlier in notebook (persisted names)
required_vars = ['mbb_selection_df', 'wbb_selection_df', 'mbb_bracket_train',
                 'wbb_bracket_train']
missing_vars = [v for v in required_vars if v not in globals()]
if missing_vars:
    raise RuntimeError(f'Missing required notebook vars: {missing_vars}. Run')

sel_mbb = mbb_selection_df.copy(); sel_mbb['league'] = 'MBB'
sel_wbb = wbb_selection_df.copy(); sel_wbb['league'] = 'WBB'
brk_mbb = mbb_bracket_train_df.copy(); brk_mbb['league'] = 'MBB'
brk_wbb = wbb_bracket_train_df.copy(); brk_wbb['league'] = 'WBB'

# Targets
sel_mbb['y'] = pd.to_numeric(sel_mbb[SOURCE_CONFIG['targets']['selection']],
                             errors='coerce')
sel_wbb['y'] = pd.to_numeric(sel_wbb[SOURCE_CONFIG['targets']['selection']],
                             errors='coerce')
brk_mbb['y'] = pd.to_numeric(brk_mbb[SOURCE_CONFIG['targets']['bracket']],
                             errors='coerce')
brk_wbb['y'] = pd.to_numeric(brk_wbb[SOURCE_CONFIG['targets']['bracket']],
                             errors='coerce')

```

```

brk_wbb['y'] = pd.to_numeric(brk_wbb[SOURCE_CONFIG['targets']['bracket']], e

selection_all = pd.concat([sel_mbb, sel_wbb], ignore_index=True)
bracket_all = pd.concat([brk_mbb, brk_wbb], ignore_index=True)

# Selection label validation notes: verify proxy derivation is deterministic
selection_label_audit = []
for lg, d in [('MBB', sel_mbb), ('WBB', sel_wbb)]:
    row = {'league': lg, 'rows': len(d)}
    row['has_elo_rank'] = 'elo_rank' in d.columns
    if 'elo_rank' in d.columns:
        proxy = (pd.to_numeric(d['elo_rank'], errors='coerce') <= 68).astype(
            row['label_matches_rank_proxy_rate'] = float((proxy == d['y']).mean(
    else:
        row['label_matches_rank_proxy_rate'] = np.nan
    selection_label_audit.append(row)
selection_label_audit = pd.DataFrame(selection_label_audit)
display(selection_label_audit)

# Confirm 2026 presence and enforce holdout split
season_presence = pd.DataFrame([
    {'task': 'selection', 'league': 'MBB', 'has_2026': bool((sel_mbb['season']
    {'task': 'selection', 'league': 'WBB', 'has_2026': bool((sel_wbb['season']
    {'task': 'bracket', 'league': 'MBB', 'has_2026': bool((brk_mbb['season'] =
    {'task': 'bracket', 'league': 'WBB', 'has_2026': bool((brk_wbb['season'] =
]))
display(season_presence)

sel_train = selection_all[selection_all['season'] < HOLDOUT_SEASON].copy()
sel_holdout = selection_all[selection_all['season'] == HOLDOUT_SEASON].copy()
brk_train = bracket_all[bracket_all['season'] < HOLDOUT_SEASON].copy()
brk_holdout = bracket_all[bracket_all['season'] == HOLDOUT_SEASON].copy()

holdout_counts = pd.DataFrame([
    {'task': 'selection', 'subset': 'train_no_2026', 'rows': len(sel_train)},
    {'task': 'selection', 'subset': 'holdout_2026', 'rows': len(sel_holdout)},
    {'task': 'bracket', 'subset': 'train_no_2026', 'rows': len(brk_train)},
    {'task': 'bracket', 'subset': 'holdout_2026', 'rows': len(brk_holdout)},
])
display(holdout_counts)

assert (sel_train['season'] == HOLDOUT_SEASON).sum() == 0, '2026 leaked into
assert (brk_train['season'] == HOLDOUT_SEASON).sum() == 0, '2026 leaked into

# Save frozen config and counts
with open(ARTIFACT_DIR / 'source_config.json', 'w') as f:
    json.dump(SOURCE_CONFIG, f, indent=2)
dataset_versions.to_csv(ARTIFACT_DIR / 'dataset_versions.csv', index=False)
holdout_counts.to_csv(ARTIFACT_DIR / 'holdout_counts.csv', index=False)

```

	league	dataset	path	exists	size_bytes	mti
0	MBB	games	games.parquet	True	1933243	01T20:41:05.381355286
1	MBB	teams	teams.parquet	True	10834	01T20:41:05.382886648
2	MBB	team_stats	team_stats.parquet	True	370792	01T20:41:05.387749677
3	MBB	elo	elo.parquet	True	32034	01T20:41:05.392896414
4	WBB	games	wbb_games.parquet	True	1464508	01T22:54:14.376332287
5	WBB	teams	wbb_teams.parquet	True	12146	01T22:54:14.377139568
6	WBB	team_stats	wbb_team_stats.parquet	True	303951	01T22:54:14.381822348
7	WBB	elo	wbb_elo.parquet	True	26448	01T22:54:14.383070946

	league	rows	has_elo_rank	label_matches_rank_proxy_rate
0	MBB	6682	True	1.0
1	WBB	5282	True	1.0

	task	league	has_2026
0	selection	MBB	True
1	selection	WBB	True
2	bracket	MBB	True
3	bracket	WBB	True

	task	subset	rows
0	selection	train_no_2026	11238
1	selection	holdout_2026	726
2	bracket	train_no_2026	160962
3	bracket	holdout_2026	10428

```
In [53]: # Rolling split generator, manifest/leakage flags, and as-of selection rebuild
import re

# Split generators

def generate_expanding_splits(seasons, val_window=2, min_train_seasons=6):
    seasons = sorted(int(s) for s in seasons)
```

```

out = []
for i in range(min_train_seasons, len(seasons) - val_window + 1):
    train = seasons[:i]
    val = seasons[i:i+val_window]
    out.append((train, val))
return out

def generate_sliding_splits(seasons, train_window=8, val_window=2):
    seasons = sorted(int(s) for s in seasons)
    out = []
    for i in range(train_window, len(seasons) - val_window + 1):
        train = seasons[i-train_window:i]
        val = seasons[i:i+val_window]
        out.append((train, val))
    return out

def build_split_index(df, task, league, scheme='expanding', val_window=2):
    d = df[df['league'] == league].copy() if league != 'Combined' else df.copy()
    seasons = sorted(pd.to_numeric(d['season'], errors='coerce').dropna().astype(int))
    seasons = [s for s in seasons if s < HOLDOUT_SEASON]
    pairs = generate_expanding_splits(seasons, val_window=val_window, min_train_seasons=min_train_seasons)
    rows = []
    split_map = {}
    for k, (tr, va) in enumerate(pairs):
        idx_tr = d.index[d['season'].isin(tr)].to_numpy()
        idx_va = d.index[d['season'].isin(va)].to_numpy()
        split_map[f'{task}_{league}_{scheme}_{k}'] = {'train_idx': idx_tr, 'val_idx': idx_va}
        rows.append({
            'task': task, 'league': league, 'scheme': scheme, 'split_id': k,
            'train_seasons': ','.join(map(str, tr)), 'val_seasons': ','.join(map(str, va)),
            'n_train': len(idx_tr), 'n_val': len(idx_va)
        })
    return split_map, pd.DataFrame(rows)

split_store = {}
split_rows = []
for task, df in [('selection', sel_train), ('bracket', brk_train)]:
    for lg in ['MBB', 'WBB', 'Combined']:
        sm1, sr1 = build_split_index(df, task, lg, scheme='expanding', val_window=val_window)
        sm2, sr2 = build_split_index(df, task, lg, scheme='sliding', val_window=val_window)
        split_store.update(sm1); split_store.update(sm2)
        split_rows.append(sr1); split_rows.append(sr2)
split_definitions = pd.concat(split_rows, ignore_index=True) if split_rows else pd.DataFrame()
display(split_definitions.head(20))
split_definitions.to_csv(ARTIFACT_DIR / 'split_definitions.csv', index=False)

# Feature provenance + leakage tagging
selection_features_raw = [
    'overall_winpct', 'conf_winpct', 'nonconf_winpct', 'sos_proxy', 'net_rating',
    'end_elo', 'last10_winpct', 'conference_proxy', 'o_ppp', 'd_ppp', 'pace', 'leakage'
]
bracket_features_raw = [
    'elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'neutral', 'home_favored'
]

post_patterns = ['seed', 'tourney', 'selection', 'bid', 'post', 'final', 'bracket']

```

```

soft_patterns = ['end_', 'season_total', 'season_avg']

def availability_for_feature(f):
    fl = f.lower()
    if any(p in fl for p in post_patterns):
        return 'post-season'
    if any(fl.startswith(p) for p in ['end_']) or 'final' in fl:
        return 'end-of-regular-season'
    if 'diff' in fl or fl in ['neutral', 'home_favored']:
        return 'pregame'
    return 'season-to-date'

manifest_rows = []
for f in sorted(set(selection_features_raw + bracket_features_raw)):
    fl = f.lower()
    leak = 0
    notes = []
    if any(p in fl for p in post_patterns):
        leak = 3; notes.append('forbidden_pattern')
    if any(p in fl for p in soft_patterns):
        leak = max(leak, 2); notes.append('soft_pattern')
    if f == 'end_elo':
        leak = max(leak, 2); notes.append('threshold_like')
    source = 'derived'
    if 'elo' in fl: source = 'elo'
    elif any(x in fl for x in ['ppp', 'pace', 'rating']): source = 'team_stats'
    elif any(x in fl for x in ['win', 'margin', 'sos', 'neutral', 'fav']): source = 'team_stats'
    manifest_rows.append({
        'feature_name': f,
        'source_table': source,
        'time_availability': availability_for_feature(f),
        'manual_leak_risk': leak,
        'notes': ';'.join(notes) if notes else 'none'
    })

feature_manifest_v2 = pd.DataFrame(manifest_rows).sort_values(['manual_leak_risk', 'notes'])
forbidden_features = feature_manifest_v2.loc[feature_manifest_v2['manual_leak_risk'] > 0]

print('Forbidden features:', forbidden_features)
feature_manifest_v2.to_csv(ARTIFACT_DIR / 'feature_manifest_v2.csv', index=False)
display(feature_manifest_v2)

# Filter modeling features to allowed set
selection_features_allowed = [f for f in selection_features_raw if f not in forbidden_features]
bracket_features_allowed = [f for f in bracket_features_raw if f not in forbidden_features]

# As-of selection rebuild (full, as-of, as-of-minus-N)

def build_selection_asof_from_games(games_file, team_stats_file, elo_file, 1):
    g = pd.read_parquet(games_file).copy()
    g['gameday'] = pd.to_datetime(g['gameday'], errors='coerce')
    g['_mm_proxy'] = (
        g['league'].isna() & g['gameday'].notna() & (g['gameday'].dt.month == 1)
    ) | (
        g['league'].isna() & g['gameday'].notna() & (g['gameday'].dt.month == 2)
    )
    g = g.loc[~g['_mm_proxy']].copy()

```

```

month, day = cutoff_month_day
cutoff = pd.to_datetime(g['season'].astype(str) + f'--{month:02d}--{day:02d}')
g = g.loc[g['gameday'].notna() & (g['gameday'] <= cutoff)].copy()

home = pd.DataFrame({
    'season': g['season'], 'gameday': g['gameday'], 'team_code': g['home_team_code'],
    'conference_proxy': g['league'],
    'win': (g['winner_code'] == g['home_code']).astype(float),
    'margin': pd.to_numeric(g['score_home'], errors='coerce') - pd.to_numeric(g['score_vis'], errors='coerce')
})

vis = pd.DataFrame({
    'season': g['season'], 'gameday': g['gameday'], 'team_code': g['visitor_team_code'],
    'conference_proxy': g['league'],
    'win': (g['winner_code'] == g['visitor_code']).astype(float),
    'margin': pd.to_numeric(g['score_vis'], errors='coerce') - pd.to_numeric(g['score_home'], errors='coerce')
})

tg = pd.concat([home, vis], ignore_index=True).dropna(subset=['season', 'team_code'])
tg = tg.sort_values(['season', 'team_code', 'gameday'])

if minus_last_n > 0:
    tg['_r'] = tg.groupby(['season', 'team_code'])['gameday'].rank(method='first')
    tg = tg.loc[tg['_r'] > minus_last_n].copy()

tg['rolling_win_10'] = tg.groupby(['season', 'team_code'])['win'].transform(
    lambda x: x.rolling(10).mean()
)

out = tg.groupby(['season', 'team_code'], as_index=False).agg(
    overall_winpct=('win', 'mean'),
    conf_winpct=('win', 'mean'),
    nonconf_winpct=('win', 'mean'),
    avg_margin=('margin', 'mean'),
    last10_winpct=('rolling_win_10', 'last')
)

# SOS proxy from opponent win%
ow = out[['season', 'team_code', 'overall_winpct']].rename(columns={'team_code': 'opp_code'})
sos = tg.merge(ow, on=['season', 'opp_code'], how='left').groupby(['season', 'team_code']).agg(
    sos_winpct=('sos_winpct', 'mean')
)
out = out.merge(sos, on=['season', 'team_code'], how='left')

# Merge team stats + end_elo + selected label proxy
ts = pd.read_parquet(team_stats_file).copy()
for c in ['o_ppp', 'd_ppp', 'pace']:
    if c not in ts.columns:
        ts[c] = np.nan
ts['net_rating'] = pd.to_numeric(ts['o_ppp'], errors='coerce') - pd.to_numeric(ts['d_ppp'], errors='coerce')
keep_ts = [c for c in ['season', 'team_code', 'o_ppp', 'd_ppp', 'pace', 'net_rating']]
out = out.merge(ts[keep_ts], on=['season', 'team_code'], how='left')

elo = pd.read_parquet(elo_file).copy().rename(columns={'code': 'team_code'})
out = out.merge(elo[['season', 'team_code', 'end_elo']], on=['season', 'team_code'], how='left')
out['elo_rank'] = out.groupby('season')['end_elo'].rank(method='first')
out['selected_proxy'] = (out['elo_rank'] <= 68).astype(int)
out['league'] = league_name
out['y'] = out['selected_proxy'].astype(int)
return out

```

```

sel_asof_mbb_v2 = build_selection_asof_from_games('games.parquet', 'team_stats.parquet')

```

```

sel_asof_wbb_v2 = build_selection_asof_from_games('wbb_games.parquet', 'wbb_
sel_asof_minus_mbb_v2 = build_selection_asof_from_games('games.parquet', 'te
sel_asof_minus_wbb_v2 = build_selection_asof_from_games('wbb_games.parquet',

selection_asof_versions = pd.DataFrame([
    {'version':'full', 'league':'MBB', 'rows':len(sel_mbb)},
    {'version':'full', 'league':'WBB', 'rows':len(sel_wbb)},
    {'version':'as_of', 'league':'MBB', 'rows':len(sel_asof_mbb_v2)},
    {'version':'as_of', 'league':'WBB', 'rows':len(sel_asof_wbb_v2)},
    {'version':'as_of_minus10', 'league':'MBB', 'rows':len(sel_asof_minus_mbb_v2)},
    {'version':'as_of_minus10', 'league':'WBB', 'rows':len(sel_asof_minus_wbb_v2)}
])
display(selection_asof_versions)
selection_asof_versions.to_csv(ARTIFACT_DIR / 'selection_asof_versions.csv',

```

	task	league	scheme	split_id	train_seasons
0	selection	MBB	expanding	0	2008,2009,2010,2011,2012,2013
1	selection	MBB	expanding	1	2008,2009,2010,2011,2012,2013,2014
2	selection	MBB	expanding	2	2008,2009,2010,2011,2012,2013,2014,2015
3	selection	MBB	expanding	3	2008,2009,2010,2011,2012,2013,2014,2015,2016
4	selection	MBB	expanding	4	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017
5	selection	MBB	expanding	5	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018
6	selection	MBB	expanding	6	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018,2019
7	selection	MBB	expanding	7	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018,2019,2020
8	selection	MBB	expanding	8	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018,2019,2020,2021
9	selection	MBB	expanding	9	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018,2019,2020,2021,2022
10	selection	MBB	expanding	10	2008,2009,2010,2011,2012,2013,2014,2015,2016,2017,2018,2019,2020,2021,2022,2023
11	selection	MBB	sliding	0	2008,2009,2010,2011,2012,2013,2014,2015
12	selection	MBB	sliding	1	2009,2010,2011,2012,2013,2014,2015,2016
13	selection	MBB	sliding	2	2010,2011,2012,2013,2014,2015,2016,2017
14	selection	MBB	sliding	3	2011,2012,2013,2014,2015,2016,2017,2018
15	selection	MBB	sliding	4	2012,2013,2014,2015,2016,2017,2018,2019
16	selection	MBB	sliding	5	2013,2014,2015,2016,2017,2018,2019,2020
17	selection	MBB	sliding	6	2014,2015,2016,2017,2018,2019,2020,2021
18	selection	MBB	sliding	7	2015,2016,2017,2018,2019,2020,2021,2022
19	selection	MBB	sliding	8	2016,2017,2018,2019,2020,2021,2022,2023

Forbidden features: []

	feature_name	source_table	time_availability	manual_leak_risk	
0	end_elo	elo	end-of-regular-season	2	soft_pattern;thres
1	avg_margin	games	season-to-date	0	
2	conf_winpct	games	season-to-date	0	
3	conference_proxy	derived	season-to-date	0	
4	d_ppp	team_stats	season-to-date	0	
5	elo_diff	elo	pregame	0	
6	home_favored	games	pregame	0	
7	last10_winpct	games	season-to-date	0	
8	league	derived	season-to-date	0	
9	net_rating	team_stats	season-to-date	0	
10	neutral	games	pregame	0	
11	nonconf_winpct	games	season-to-date	0	
12	o_ppp	team_stats	season-to-date	0	
13	overall_winpct	games	season-to-date	0	
14	pace	team_stats	season-to-date	0	
15	pace_diff	team_stats	pregame	0	
16	ppp_diff	team_stats	pregame	0	
17	sos_proxy	games	season-to-date	0	
18	winpct_diff	games	pregame	0	

	version	league	rows
0	full	MBB	6682
1	full	WBB	5282
2	as_of	MBB	6682
3	as_of	WBB	5282
4	as_of_minus10	MBB	6674
5	as_of_minus10	WBB	5264

```
In [54]: # Core evaluation utilities, baselines, ablations, single-feature, permutati
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import roc_auc_score, average_precision_score, log_loss
```

```

from sklearn.calibration import calibration_curve, CalibratedClassifierCV

from sklearn.linear_model import LogisticRegression, Perceptron
from sklearn.svm import LinearSVC, SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, A
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, Quadra
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import RandomizedSearchCV

np.random.seed(SEED)

def ece_score(y_true, y_prob, n_bins=10):
    y_true = np.asarray(y_true).astype(float)
    y_prob = np.asarray(y_prob).astype(float)
    bins = np.linspace(0, 1, n_bins+1)
    idx = np.digitize(y_prob, bins, right=True)
    ece = 0.0
    n = len(y_true)
    for b in range(1, n_bins+1):
        m = idx == b
        if m.sum() == 0:
            continue
        ece += (m.sum()/n) * abs(y_true[m].mean() - y_prob[m].mean())
    return float(ece)

def prep_xy(df, features, target='y'):
    d = df.copy()
    feats = [f for f in features if f in d.columns]
    feats = [f for f in feats if d[f].notna().any()]
    d = d.dropna(subset=[target, 'season'])
    X = d[feats].copy()
    y = d[target].astype(int).copy()
    return X, y, d, feats

def fit_predict(task, train_df, val_df, features, estimator, calibrate_metho
Xtr, ytr, tr, feats = prep_xy(train_df, features)
Xva, yva, va, _ = prep_xy(val_df, feats)
if len(feats) == 0 or ytr.nunique() < 2 or yva.nunique() < 2:
    return None

cat_cols = [c for c in feats if str(Xtr[c].dtype) in ('object', 'category')
num_cols = [c for c in feats if c not in cat_cols]

pre = ColumnTransformer([
    ('num', Pipeline([('imp', SimpleImputer(strategy='median'))], ('sc',
    ('cat', Pipeline([('imp', SimpleImputer(strategy='most_frequent'))],
    ], remainder='drop')

pipe = Pipeline([('pre', pre), ('clf', estimator)])
pipe.fit(Xtr, ytr)

```

```

if hasattr(pipe, 'predict_proba'):
    p = pipe.predict_proba(Xva)[:,1]
else:
    # convert margins to pseudo-prob via sigmoid
    z = pipe.decision_function(Xva)
    p = 1 / (1 + np.exp(-z))

if calibrate_method in ('platt', 'isotonic'):
    cal = CalibratedClassifierCV(pipe, method='sigmoid' if calibrate_met
    cal.fit(Xtr, ytr)
    p = cal.predict_proba(Xva)[:,1]
    model_obj = cal
else:
    model_obj = pipe

k = int(yva.sum()) if task == 'selection' else None
if task == 'selection':
    top_idx = np.argsort(-p)[:max(1, min(k, len(p)))]
    precision_k = float(yva.iloc[top_idx].mean())
    recall_k = float(yva.iloc[top_idx].sum() / max(1, yva.sum()))
else:
    precision_k = np.nan
    recall_k = np.nan

out = {
    'auc': float(roc_auc_score(yva, p)),
    'pr_auc': float(average_precision_score(yva, p)) if task == 'selecti
    'log_loss': float(log_loss(yva, p, labels=[0,1])),
    'brier': float(brier_score_loss(yva, p)),
    'ece': float(ece_score(yva, p, n_bins=10)),
    'precision_at_k': precision_k,
    'recall_at_k': recall_k,
    'n_train': int(len(ytr)),
    'n_val': int(len(yva)),
    'features_used': feats,
    'y_true': yva.to_numpy(),
    'y_prob': p,
    'val_df': va,
    'model': model_obj,
}
return out

def get_estimator_zoo(task='selection'):
    zoo = {
        'logreg_l2': LogisticRegression(max_iter=4000, penalty='l2', C=1.0,
        'logreg_l1': LogisticRegression(max_iter=4000, penalty='l1', solver=
        'logreg_elastic': LogisticRegression(max_iter=4000, penalty='elastic
        'linear_svm': LinearSVC(C=1.0, random_state=SEED),
        'rbf_svm': SVC(C=2.0, kernel='rbf', gamma='scale', probability=True,
        'poly_svm': SVC(C=2.0, kernel='poly', degree=3, gamma='scale', proba
        'decision_tree': DecisionTreeClassifier(max_depth=6, random_state=SE
        'random_forest': RandomForestClassifier(n_estimators=300, max_depth=
        'extra_trees': ExtraTreesClassifier(n_estimators=300, max_depth=None
        'adaboost': AdaBoostClassifier(n_estimators=250, random_state=SEED),

```

```

        'grad_boost': GradientBoostingClassifier(random_state=SEED),
        'xgb_style_hgb': HistGradientBoostingClassifier(random_state=SEED),
        'knn': KNeighborsClassifier(n_neighbors=35),
        'gaussian_nb': GaussianNB(),
        'lda': LinearDiscriminantAnalysis(),
        'qda': QuadraticDiscriminantAnalysis(reg_param=0.1),
        'perceptron': Perceptron(max_iter=2000, random_state=SEED),
        'mlp': MLPClassifier(hidden_layer_sizes=(128,64), early_stopping=True)
    }

    # Optional xgboost
    try:
        from xgboost import XGBClassifier
        zoo['xgboost'] = XGBClassifier(
            n_estimators=500, max_depth=4, learning_rate=0.05,
            subsample=0.9, colsample_bytree=0.9, eval_metric='logloss', random_state=SEED
        )
    except Exception:
        pass

    # Optional probit
    try:
        import statsmodels.api as sm # noqa
        zoo['probit_proxy'] = LogisticRegression(max_iter=4000, penalty='l2', random_state=SEED)
    except Exception:
        pass

    # Autoencoder placeholder via compact MLP representation proxy
    zoo['autoencoder_proxy_mlp'] = MLPClassifier(hidden_layer_sizes=(64,16,64), random_state=SEED)

    return zoo


def run_baselines_and_ablations():
    rows = []

    # Baseline feature sets
    sel_base_sets = {
        'intercept_only': [],
        'elo_only': ['end_elo', 'league'],
        'winpct_only': ['overall_winpct', 'league'],
    }
    brk_base_sets = {
        'elo_diff_only': ['elo_diff', 'league'],
    }

    # Task datasets (train only, no 2026)
    datasets = {
        ('selection', 'Combined'): sel_train,
        ('selection', 'MBB'): sel_train[sel_train['league']=='MBB'].copy(),
        ('selection', 'WBB'): sel_train[sel_train['league']=='WBB'].copy(),
        ('bracket', 'Combined'): brk_train,
        ('bracket', 'MBB'): brk_train[brk_train['league']=='MBB'].copy(),
        ('bracket', 'WBB'): brk_train[brk_train['league']=='WBB'].copy(),
    }

    # Baselines on last expanding split
    for (task, lg), d in datasets.items():
        seasons = sorted(d['season'].dropna().astype(int).unique())
        seasons = [s for s in seasons if s < HOLDOUT_SEASON]

```

```

splits = generate_expanding_splits(seasons, val_window=2, min_train_
if not splits:
    continue
tr_s, va_s = splits[-1]
tr = d[d['season'].isin(tr_s)].copy()
va = d[d['season'].isin(va_s)].copy()
if task == 'selection':
    for name, feats in sel_base_sets.items():
        if name == 'intercept_only':
            p = np.repeat(tr['y'].mean(), len(va))
            y = va['y'].astype(int).to_numpy()
            k = int(y.sum())
            top_idx = np.argsort(-p)[:max(1, min(k, len(y)))]
            rows.append({
                'task':task, 'league':lg, 'regime':'baseline', 'model':
                'auc':float(roc_auc_score(y,p)) if len(np.unique(y))
                'pr_auc':float(average_precision_score(y,p)) if len(
                'log_loss':float(log_loss(y,p, labels=[0,1])),
                'brier':float(brier_score_loss(y,p)),
                'ece':float(ece_score(y,p)),
                'precision_at_k':float(y[top_idx].mean()),
                'recall_at_k':float(y[top_idx].sum()/max(1,y.sum()))
            })
        else:
            o = fit_predict(task, tr, va, feats, LogisticRegression(
            if o is None:
                continue
            rows.append({'task':task, 'league':lg, 'regime':'baseline'
    else:
        for name, feats in brk_base_sets.items():
            o = fit_predict(task, tr, va, feats, LogisticRegression(max_
            if o is None:
                continue
            rows.append({'task':task, 'league':lg, 'regime':'baseline', 'mc

# Ablation ladders
sel_ablation = {
    'S1_winpct': ['overall_winpct', 'league'],
    'S2_elo': ['overall_winpct', 'end_elo', 'league'],
    'S3_sos': ['overall_winpct', 'end_elo', 'sos_proxy', 'league'],
    'S4_efficiency': ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating'],
    'S5_recency': ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating',
    'S6_conference': ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating
}
brk_ablation = {
    'B1_elo_diff': ['elo_diff', 'league'],
    'B2_add_winpct': ['elo_diff', 'winpct_diff', 'league'],
    'B3_add_eff': ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'leag
    'B4_add_interactions': ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_di
    'B5_full': ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'home_fa
}

for (task, lg), d in datasets.items():
    seasons = sorted(d['season'].dropna().astype(int).unique())
    seasons = [s for s in seasons if s < HOLDOUT_SEASON]
    splits = generate_expanding_splits(seasons, val_window=2, min_train_

```

```

        if not splits:
            continue
        for sid, (tr_s, va_s) in enumerate(splits):
            tr = d[d['season'].isin(tr_s)].copy()
            va = d[d['season'].isin(va_s)].copy()
            ladder = sel_ablation if task == 'selection' else brk_ablation
            for step, feats in ladder.items():
                o = fit_predict(task, tr, va, feats, LogisticRegression(max_
                if o is None:
                    continue
                rows.append({
                    'task':task, 'league':lg, 'regime':'ablation', 'split_id':s
                    **{k:o[k] for k in ['auc', 'pr_auc', 'log_loss', 'brier', 'e
                })

    out = pd.DataFrame(rows)
    return out

baseline_ablation_results = run_baselines_and_ablations()
display(baseline_ablation_results.head(40))
baseline_ablation_results.to_csv(ARTIFACT_DIR / 'baseline_ablation_results.c

# Single-feature smoke tests
sf_rows = []
for task, d, candidates in [
    ('selection', sel_train, [f for f in selection_features_allowed if f !=
    ('bracket', brk_train, [f for f in bracket_features_allowed if f != 'lea
]:
    seasons = sorted(d['season'].dropna().astype(int).unique())
    splits = generate_expanding_splits([s for s in seasons if s < HOLDOUT_SE
    if not splits:
        continue
    tr_s, va_s = splits[-1]
    tr = d[d['season'].isin(tr_s)].copy()
    va = d[d['season'].isin(va_s)].copy()
    for f in candidates:
        o = fit_predict(task, tr, va, [f, 'league'], LogisticRegression(max_i
        if o is None:
            continue
        row = {'task':task, 'feature':f, 'auc':o['auc']}
        if task == 'selection':
            # bubble-only
            b = va.copy()
            if 'end_elo' in b.columns:
                b['elo_pct'] = b.groupby('season')['end_elo'].rank(pct=True)
                bb = b[(b['elo_pct'] >= 0.40) & (b['elo_pct'] <= 0.80)].copy()
                ob = fit_predict(task, tr, bb, [f, 'league'], LogisticRegress
                row['bubble_auc'] = ob['auc'] if ob is not None else np.nan
            else:
                row['bubble_auc'] = np.nan
        sf_rows.append(row)
single_feature_results_v2 = pd.DataFrame(sf_rows).sort_values(['task', 'auc'])
display(single_feature_results_v2.head(30))
single_feature_results_v2.to_csv(ARTIFACT_DIR / 'single_feature_results_v2.c

# Permutation tests

```

```

perm_rows = []
perm_imp_rows = []
perm_iters = 50
for task, d, feats in [
    ('selection', sel_train, ['overall_winpct', 'end_elo', 'sos_proxy', 'net_ra
    ('bracket', brk_train, ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff',
]:
    for lg in ['MBB', 'WBB', 'Combined']:
        dd = d[d['league']==lg].copy() if lg != 'Combined' else d.copy()
        seasons = sorted(dd['season'].dropna().astype(int).unique())
        splits = generate_expanding_splits([s for s in seasons if s < HOLDOUT])
        if not splits:
            continue
        tr_s, va_s = splits[-1]
        tr = dd[dd['season'].isin(tr_s)].copy()
        va = dd[dd['season'].isin(va_s)].copy()

        real = fit_predict(task, tr, va, feats, LogisticRegression(max_iter=
        if real is None:
            continue

        null_aucs = []
        for i in range(perm_iters):
            tp = tr.copy()
            tp['y'] = np.random.RandomState(SEED + i).permutation(tp['y'].values)
            o = fit_predict(task, tp, va, feats, LogisticRegression(max_iter=
            if o is None:
                continue
            null_aucs.append(o['auc'])

        perm_rows.append({
            'task':task, 'league':lg, 'real_auc':real['auc'],
            'null_mean_auc':float(np.mean(null_aucs)) if null_aucs else np.nan,
            'null_std_auc':float(np.std(null_aucs)) if null_aucs else np.nan,
            'null_max_auc':float(np.max(null_aucs)) if null_aucs else np.nan
        })

    # feature permutation importance
    y_true = real['y_true']
    base_p = real['y_prob']
    base_auc = roc_auc_score(y_true, base_p)
    model = real['model']
    vdf = real['val_df'].copy()
    for f in real['features_used']:
        xx = vdf[real['features_used']].copy()
        xx[f] = np.random.RandomState(SEED).permutation(xx[f].values)
        if hasattr(model, 'predict_proba'):
            p = model.predict_proba(xx)[:,1]
        else:
            z = model.decision_function(xx)
            p = 1/(1+np.exp(-z))
        auc_s = roc_auc_score(y_true, p)
        perm_imp_rows.append({'task':task, 'league':lg, 'feature':f, 'base_

label_permutation_results = pd.DataFrame(perm_rows)
feature_permutation_importance_v2 = pd.DataFrame(perm_imp_rows)

```

```

display(label_permutation_results)
display(feature_permutation_importance_v2.head(30))
label_permutation_results.to_csv(ARTIFACT_DIR / 'label_permutation_results.csv')
feature_permutation_importance_v2.to_csv(ARTIFACT_DIR / 'feature_permutation_importance_v2.csv')

# Bubble-only evaluation and calibration tables
bubble_rows = []
for lg in ['MBB', 'WBB']:
    d = sel_train[sel_train['league']==lg].copy()
    seasons = sorted(d['season'].dropna().astype(int).unique())
    splits = generate_expanding_splits([s for s in seasons if s < HOLDOUT_SEASON])
    if not splits:
        continue
    tr_s, va_s = splits[-1]
    tr = d[d['season'].isin(tr_s)].copy()
    va = d[d['season'].isin(va_s)].copy()
    va['elo_pct'] = va.groupby('season')['end_elo'].rank(pct=True)
    bubble = va[(va['elo_pct']>=0.40) & (va['elo_pct']<=0.80)].copy()
    feats = ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating', 'last10_winpct']
    of = fit_predict('selection', tr, va, feats, LogisticRegression(max_iter=1000))
    ob = fit_predict('selection', tr, bubble, feats, LogisticRegression(max_iter=1000))
    if of is None or ob is None:
        continue
    bubble_rows.append({'league':lg, 'auc_full':of['auc'], 'pr_auc_full':ob['pr_auc']})
bubble_eval_v2 = pd.DataFrame(bubble_rows)
display(bubble_eval_v2)
bubble_eval_v2.to_csv(ARTIFACT_DIR / 'bubble_eval_v2.csv', index=False)

cal_rows = []
for lg in ['MBB', 'WBB', 'Combined']:
    d = brk_train[brk_train['league']==lg].copy() if lg != 'Combined' else brk_train
    seasons = sorted(d['season'].dropna().astype(int).unique())
    splits = generate_expanding_splits([s for s in seasons if s < HOLDOUT_SEASON])
    if not splits:
        continue
    tr_s, va_s = splits[-1]
    tr = d[d['season'].isin(tr_s)].copy(); va = d[d['season'].isin(va_s)].copy()
    feats = ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'home_favored', 'elo_pct']
    for method in [None, 'platt', 'isotonic']:
        o = fit_predict('bracket', tr, va, feats, LogisticRegression(max_iter=1000))
        if o is None:
            continue
        cal_rows.append({'league':lg, 'method':method if method is not None else 'uncalibrated'})

calibration_results_v2 = pd.DataFrame(cal_rows)
display(calibration_results_v2)
calibration_results_v2.to_csv(ARTIFACT_DIR / 'calibration_results_v2.csv', index=False)

# Model family implementation + lightweight tuning over last split per task/league

def param_distributions(name):
    return {
        'logreg_l2': {'clf__C': np.logspace(-2, 1, 20)},
        'logreg_l1': {'clf__C': np.logspace(-2, 1, 20)},
        'logreg_elastic': {'clf__C': np.logspace(-2, 1, 10), 'clf__l1_ratio': np.linspace(0.1, 1, 10)},
        'linear_svm': {'clf__C': np.logspace(-2, 1, 20)},
    }

```



```

'decision_tree': {'clf__max_depth': [3,4,5,6,8, None], 'clf__min_samp
'random_forest': {'clf__n_estimators': [200,300,500], 'clf__max_depth
'extra_trees': {'clf__n_estimators': [200,300,500], 'clf__max_depth':
'adaboost': {'clf__n_estimators': [100,200,400], 'clf__learning_rate'
'grad_boost': {'clf__n_estimators': [100,200,400], 'clf__learning_rat
'xgb_style_hgb': {'clf__max_depth': [3,5, None], 'clf__learning_rate':
'knn': {'clf__n_neighbors': [15,25,35,55], 'clf__weights': ['uniform',
'mlp': {'clf__alpha': np.logspace(-5,-2,10), 'clf__hidden_layer_sizes
'perceptron': {'clf__alpha': np.logspace(-6,-2,10)}},
}.get(name, {})

```

```

def tuned_family_run(task, regime='pooled_or_separate'):
    zoo = get_estimator_zoo(task)
    rows = []
    datasets = [('Combined', sel_train if task=='selection' else brk_train)]
    feats = ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating', 'last10_wir

    for lg, d in datasets:
        seasons = sorted(d['season'].dropna().astype(int).unique())
        splits = generate_expanding_splits([s for s in seasons if s < HOLDOUT])
        if not splits:
            continue
        tr_s, va_s = splits[-1]
        tr = d[d['season'].isin(tr_s)].copy(); va = d[d['season'].isin(va_s)]

        # build preprocessor once
        Xtr, ytr, _, feat_use = prep_xy(tr, feats)
        Xva, yva, _, _ = prep_xy(va, feat_use)
        if len(feat_use)==0 or ytr.nunique()<2 or yva.nunique()<2:
            continue

        cat_cols = [c for c in feat_use if str(Xtr[c].dtype) in ('object', 'c
        num_cols = [c for c in feat_use if c not in cat_cols]
        pre = ColumnTransformer([
            ('num', Pipeline([('imp', SimpleImputer(strategy='median'))], ('s
            ('cat', Pipeline([('imp', SimpleImputer(strategy='most_frequent'
        ], remainder='drop')

        for name, est in zoo.items():
            # keep expensive kernels tractable on sampled rows
            tr_fit = tr.copy()
            if name in ('rbf_svm', 'poly_svm') and len(tr_fit) > 12000:
                tr_fit = tr_fit.sample(n=12000, random_state=SEED)

            Xf, yf, _, feat_use2 = prep_xy(tr_fit, feat_use)
            Xv, yv, _, _ = prep_xy(va, feat_use2)
            if len(feat_use2)==0 or yf.nunique()<2 or yv.nunique()<2:
                continue

            pipe = Pipeline([('pre', pre), ('clf', est)])
            pgrid = param_distributions(name)
            try:
                if pgrid:
                    rs = RandomizedSearchCV(pipe, pgrid, n_iter=min(10, max(
                    rs.fit(Xf, yf)

```

```

        best = rs.best_estimator_
        best_params = rs.best_params_
    else:
        best = pipe.fit(Xf, yf)
        best_params = {}

    if hasattr(best, 'predict_proba'):
        prob = best.predict_proba(Xv)[:,-1]
    else:
        z = best.decision_function(Xv)
        prob = 1/(1+np.exp(-z))

    metrics = {
        'auc': float(roc_auc_score(yv, prob)),
        'pr_auc': float(average_precision_score(yv, prob)) if ta
        'log_loss': float(log_loss(yv, prob, labels=[0,1])),
        'brier': float(brier_score_loss(yv, prob)),
        'ece': float(ece_score(yv, prob, 10)),
    }
    if task=='selection':
        k = max(1, int(yv.sum()))
        idx = np.argsort(-prob)[:min(k, len(prob))]
        metrics['precision_at_k'] = float(yv.iloc[idx].mean())
    else:
        metrics['precision_at_k'] = np.nan

    rows.append({'task':task, 'regime':regime, 'league':lg, 'model_
except Exception as e:
    rows.append({'task':task, 'regime':regime, 'league':lg, 'model_

return pd.DataFrame(rows)

```

```

selection_family_results_pooled = tuned_family_run('selection', regime='pool
selection_family_results_sep = tuned_family_run('selection', regime='separat
bracket_family_results_pooled = tuned_family_run('bracket', regime='pooled')
bracket_family_results_sep = tuned_family_run('bracket', regime='separate')

```

```

model_family_results_v2 = pd.concat([
    selection_family_results_pooled,
    selection_family_results_sep,
    bracket_family_results_pooled,
    bracket_family_results_sep,
], ignore_index=True)

```

```

display(model_family_results_v2.head(30))
model_family_results_v2.to_csv(ARTIFACT_DIR / 'model_family_results_v2.csv',

```

	task	league	regime	model	auc	pr_auc	log_loss	br
0	selection	Combined	baseline	intercept_only	0.500000	0.187845	0.483084	0.1525
1	selection	Combined	baseline	elo_only	0.999912	0.999628	0.024758	0.0060
2	selection	Combined	baseline	winpct_only	0.807117	0.467292	0.396769	0.1250
3	selection	MBB	baseline	intercept_only	0.500000	0.187328	0.482494	0.1522
4	selection	MBB	baseline	elo_only	1.000000	1.000000	0.031464	0.0075
5	selection	MBB	baseline	winpct_only	0.788491	0.456045	0.406407	0.1266
6	selection	WBB	baseline	intercept_only	0.500000	0.188366	0.484194	0.1529
7	selection	WBB	baseline	elo_only	1.000000	1.000000	0.029869	0.0071
8	selection	WBB	baseline	winpct_only	0.828793	0.486413	0.393138	0.1248
9	bracket	Combined	baseline	elo_diff_only	0.752589	NaN	0.564506	0.1907
10	bracket	MBB	baseline	elo_diff_only	0.717991	NaN	0.571141	0.1938
11	bracket	WBB	baseline	elo_diff_only	0.783395	NaN	0.560618	0.1888
12	selection	Combined	ablation	S1_winpct	0.702190	0.384530	0.461834	0.1439
13	selection	Combined	ablation	S2_elo	0.999844	0.999360	0.036499	0.0089
14	selection	Combined	ablation	S3_sos	0.999844	0.999360	0.036522	0.0089
15	selection	Combined	ablation	S4_efficiency	0.999834	0.999318	0.036656	0.0090
16	selection	Combined	ablation	S5_recency	0.999837	0.999332	0.036635	0.0090
17	selection	Combined	ablation	S6_conference	0.999570	0.998291	0.035597	0.0088
18	selection	Combined	ablation	S1_winpct	0.708825	0.395446	0.450953	0.1419
19	selection	Combined	ablation	S2_elo	0.999867	0.999454	0.034524	0.0084
20	selection	Combined	ablation	S3_sos	0.999876	0.999493	0.034494	0.0084
21	selection	Combined	ablation	S4_efficiency	0.999880	0.999507	0.034555	0.0084
22	selection	Combined	ablation	S5_recency	0.999873	0.999481	0.034546	0.0084
23	selection	Combined	ablation	S6_conference	0.999597	0.998404	0.033561	0.0083
24	selection	Combined	ablation	S1_winpct	0.727803	0.424035	0.439248	0.1387
25	selection	Combined	ablation	S2_elo	0.999883	0.999522	0.032908	0.0080
26	selection	Combined	ablation	S3_sos	0.999889	0.999547	0.032966	0.0080
27	selection	Combined	ablation	S4_efficiency	0.999893	0.999559	0.033076	0.0080
28	selection	Combined	ablation	S5_recency	0.999896	0.999573	0.033082	0.0081
29	selection	Combined	ablation	S6_conference	0.999639	0.998577	0.031885	0.0078
30	selection	Combined	ablation	S1_winpct	0.735581	0.443944	0.434213	0.1371
31	selection	Combined	ablation	S2_elo	0.999909	0.999627	0.031618	0.0077

	task	league	regime	model	auc	pr_auc	log_loss	br
32	selection	Combined	ablation	S3_sos	0.999915	0.999653	0.031722	0.0077
33	selection	Combined	ablation	S4_efficiency	0.999902	0.999600	0.031848	0.0078
34	selection	Combined	ablation	S5_recency	0.999912	0.999639	0.031852	0.0078
35	selection	Combined	ablation	S6_conference	0.999674	0.998717	0.030716	0.0076
36	selection	Combined	ablation	S1_winpct	0.714505	0.419197	0.443033	0.1395
37	selection	Combined	ablation	S2_elo	0.999919	0.999666	0.030418	0.0074
38	selection	Combined	ablation	S3_sos	0.999925	0.999692	0.030604	0.0074
39	selection	Combined	ablation	S4_efficiency	0.999916	0.999652	0.030594	0.0075

	task	feature	auc	bubble_auc
0	bracket	winpct_diff	0.849795	NaN
1	bracket	ppp_diff	0.785544	NaN
2	bracket	elo_diff	0.752589	NaN
3	bracket	home_favored	0.653541	NaN
4	bracket	pace_diff	0.572860	NaN
5	bracket	neutral	0.538377	NaN
6	selection	end_elo	0.999912	1.000000
7	selection	sos_proxy	0.924390	0.916688
8	selection	conference_proxy	0.920657	0.871779
9	selection	avg_margin	0.826928	0.547285
10	selection	o_ppp	0.825411	0.604704
11	selection	overall_winpct	0.807117	0.552039
12	selection	conf_winpct	0.807117	0.552039
13	selection	nonconf_winpct	0.807117	0.552039
14	selection	last10_winpct	0.674801	0.542907
15	selection	pace	0.552629	0.530898
16	selection	net_rating	0.499150	0.728546
17	selection	d_ppp	0.499150	0.728546

	task	league	real_auc	null_mean_auc	null_std_auc	null_max_auc
0	selection	MBB	1.000000	0.483764	0.104935	0.672071
1	selection	WBB	0.999950	0.507993	0.122973	0.761556
2	selection	Combined	0.999931	0.497281	0.108352	0.682539
3	bracket	MBB	0.840357	0.503889	0.128450	0.778175
4	bracket	WBB	0.880111	0.520999	0.164213	0.856622
5	bracket	Combined	0.861167	0.502018	0.144394	0.799039

	task	league	feature	base_auc	auc_after_shuffle	auc_drop
0	selection	MBB	overall_winpct	1.000000	1.000000	0.000000
1	selection	MBB	end_elo	1.000000	0.507378	0.492622
2	selection	MBB	sos_proxy	1.000000	1.000000	0.000000
3	selection	MBB	last10_winpct	1.000000	1.000000	0.000000
4	selection	MBB	conference_proxy	1.000000	0.998754	0.001246
5	selection	MBB	league	1.000000	1.000000	0.000000
6	selection	WBB	overall_winpct	0.999950	0.999962	-0.000013
7	selection	WBB	end_elo	0.999950	0.542248	0.457702
8	selection	WBB	sos_proxy	0.999950	1.000000	-0.000050
9	selection	WBB	last10_winpct	0.999950	0.999937	0.000013
10	selection	WBB	conference_proxy	0.999950	0.999636	0.000314
11	selection	WBB	league	0.999950	0.999950	0.000000
12	selection	Combined	overall_winpct	0.999931	0.999916	0.000016
13	selection	Combined	end_elo	0.999931	0.504114	0.495817
14	selection	Combined	sos_proxy	0.999931	0.999759	0.000172
15	selection	Combined	last10_winpct	0.999931	0.999937	-0.000006
16	selection	Combined	conference_proxy	0.999931	0.999440	0.000491
17	selection	Combined	league	0.999931	0.997193	0.002739
18	bracket	MBB	elo_diff	0.840357	0.824852	0.015504
19	bracket	MBB	winpct_diff	0.840357	0.572836	0.267520
20	bracket	MBB	ppp_diff	0.840357	0.840319	0.000038
21	bracket	MBB	pace_diff	0.840357	0.837740	0.002616
22	bracket	MBB	home_favored	0.840357	0.840664	-0.000307
23	bracket	MBB	neutral	0.840357	0.833496	0.006860
24	bracket	MBB	league	0.840357	0.840357	0.000000
25	bracket	WBB	elo_diff	0.880111	0.865026	0.015085
26	bracket	WBB	winpct_diff	0.880111	0.602770	0.277341
27	bracket	WBB	ppp_diff	0.880111	0.877985	0.002126
28	bracket	WBB	pace_diff	0.880111	0.879327	0.000785
29	bracket	WBB	home_favored	0.880111	0.880242	-0.000131

	league	method	auc	log_loss	brier	ece
0	MBB	uncalibrated	0.840357	0.452095	0.149208	0.008446
1	MBB	platt	0.840386	0.452007	0.149164	0.007271
2	MBB	isotonic	0.840317	0.452111	0.149269	0.008542
3	WBB	uncalibrated	0.880111	0.413111	0.134510	0.011829
4	WBB	platt	0.880125	0.413097	0.134507	0.012212
5	WBB	isotonic	0.880084	0.412181	0.134344	0.005581
6	Combined	uncalibrated	0.861167	0.433181	0.142041	0.008521
7	Combined	platt	0.861174	0.433131	0.142024	0.010794
8	Combined	isotonic	0.861032	0.432876	0.142026	0.007488

	task	regime	league	model_family	n_train	n_val	best
0	selection	pooled	Combined	logreg_l2	9790	1448	6.9519279
1	selection	pooled	Combined	logreg_l1	9790	1448	0.061584821
2	selection	pooled	Combined	logreg_elastic	9790	1448	{"clf__l1_0.0464
3	selection	pooled	Combined	linear_svm	9790	1448	6.9519279
4	selection	pooled	Combined	rbf_svm	9790	1448	
5	selection	pooled	Combined	poly_svm	9790	1448	
6	selection	pooled	Combined	decision_tree	9790	1448	{"clf__min_sam1, "clf__max
7	selection	pooled	Combined	random_forest	9790	1448	{"clf__n_estima"clf__min_s
8	selection	pooled	Combined	extra_trees	9790	1448	{"clf__n_estima"clf__min_s
9	selection	pooled	Combined	adaboost	9790	1448	{"clf__n_estima"clf__leari
10	selection	pooled	Combined	grad_boost	9790	1448	{"clf__n_estima"clf__max_
11	selection	pooled	Combined	xgb_style_hgb	9790	1448	
12	selection	pooled	Combined	knn	9790	1448	{"clf_" "clf__n_r
13	selection	pooled	Combined	gaussian_nb	9790	1448	
14	selection	pooled	Combined	lda	9790	1448	

	task	regime	league	model_family	n_train	n_val	best
15	selection	pooled	Combined	qda	9790	1448	
16	selection	pooled	Combined	perceptron	9790	1448	{"clf__alp
17	selection	pooled	Combined	mlp	9790	1448	{"clf__hidden_la [128, 64
18	selection	pooled	Combined	xgboost	9790	1448	
19	selection	pooled	Combined	probit_proxy	9790	1448	
20	selection	pooled	Combined	autoencoder_proxy_mlp	9790	1448	
21	selection	separate	MBB	logreg_l2	5592	726	6.9519279
22	selection	separate	MBB	logreg_l1	5592	726	{"clf
23	selection	separate	MBB	logreg_elastic	5592	726	{"clf__l1_ 0.0464
24	selection	separate	MBB	linear_svm	5592	726	6.9519279
25	selection	separate	MBB	rbf_svm	5592	726	
26	selection	separate	MBB	poly_svm	5592	726	
27	selection	separate	MBB	decision_tree	5592	726	{"clf__min_sam 10, "clf__ma
28	selection	separate	MBB	random_forest	5592	726	{"clf__n_estima "clf__min_s
29	selection	separate	MBB	extra_trees	5592	726	{"clf__n_estima "clf__min_s

```
In [55]: # Leakage risk, gating, final model selection, locked 2026 eval, simulation,
import joblib

# Leakage risk score
risk_rows = []

# helper maps
single_max = single_feature_results_v2.groupby('task')['auc'].max().to_dict()
perm_map = {(r['task'], r['league']): r for _, r in label_permutation_results_v2}

# top importance by task/league
top_imp = pd.DataFrame()
if len(feature_permutation_importance_v2):
    top_imp = feature_permutation_importance_v2.sort_values('auc_drop', asce
```

```

top_imp = top_imp.merge(feature_manifest_v2[['feature_name', 'manual_leak

# candidate pool from model_family results
cands = model_family_results_v2.copy()
if len(cands):
    for _, r in cands.iterrows():
        score = 0
        task, lg = r['task'], r['league']
        if task == 'selection' and pd.notna(r['auc']) and r['auc'] >= 0.995:
            score += 30
        if task == 'selection' and single_max.get('selection', 0) >= 0.98:
            score += 20
        pm = perm_map.get((task, lg))
        if pm is not None and pd.notna(pm['null_mean_auc']) and pm['null_me
            score += 20
        if task == 'selection' and len(bubble_eval_v2):
            b = bubble_eval_v2[bubble_eval_v2['league']==lg]
            if len(b) and float(b['auc_bubble'].iloc[0]) >= 0.98:
                score += 15
        ti = top_imp[(top_imp['task']==task) & (top_imp['league']==lg)] if 1
        if len(ti) and pd.notna(ti['manual_leak_risk'].iloc[0]) and int(ti['
            score += 15

        band = 'low' if score <= 25 else 'moderate' if score <= 50 else 'high
        risk_rows.append({
            'task':task, 'league':lg, 'model_family':r['model_family'],
            'auc':r['auc'], 'pr_auc':r.get('pr_auc', np.nan), 'log_loss':r['log
            'leakage_risk_score':score, 'risk_band':band
        })

leakage_risk_v2 = pd.DataFrame(risk_rows)
display(leakage_risk_v2.head(30))
leakage_risk_v2.to_csv(ARTIFACT_DIR / 'leakage_risk_v2.csv', index=False)

# Acceptance/rejection gating
accept_rows, reject_rows = [], []
for _, r in leakage_risk_v2.iterrows():
    reasons = []
    task, lg = r['task'], r['league']

    if r['risk_band'] in ['high', 'critical']:
        reasons.append('leakage_risk_high_or_critical')

    if task == 'selection':
        # as-of collapse gate
        m = selection_asof_versions[selection_asof_versions['league']==lg]
        # evaluate collapse using baseline run proxies from existing compari
        # here use conservative gate from risk and single-feature
        if single_max.get('selection', 0) >= 0.99:
            reasons.append('single_feature_auc_extreme')
    else:
        # calibration gate
        cal = calibration_results_v2[(calibration_results_v2['league']==lg)]
        if len(cal):
            best_ece = cal['ece'].min()
            if best_ece > 0.05:

```

```

        reasons.append('ece_above_0.05_after_calibration')

    if reasons:
        reject_rows.append({'task':task, 'league':lg, 'model_family':r['model_
    else:
        accept_rows.append({'task':task, 'league':lg, 'model_family':r['model_

accepted_models_v2 = pd.DataFrame(accept_rows)
rejected_models_v2 = pd.DataFrame(reject_rows)
display(accepted_models_v2.head(20))
display(rejected_models_v2.head(20))
accepted_models_v2.to_csv(ARTIFACT_DIR / 'accepted_models_v2.csv', index=False)
rejected_models_v2.to_csv(ARTIFACT_DIR / 'rejected_models_v2.csv', index=False)

# Final model selection logic
# Selection: pooled model, optimize bubble PR-AUC + precision@K + stability
sel_pool_ok = model_family_results_v2[(model_family_results_v2['task']=='sel
if len(sel_pool_ok):
    sel_pool_ok = sel_pool_ok.sort_values(['pr_auc', 'precision_at_k', 'auc'],
    final_selection_choice = sel_pool_ok.iloc[0]
else:
    final_selection_choice = None

# Bracket: separate MBB/WBB, optimize log-loss + calibration + AUC
final_bracket_choices = {}
for lg in ['MBB', 'WBB']:
    d = model_family_results_v2[(model_family_results_v2['task']=='bracket')
    if len(d):
        d = d.sort_values(['log_loss', 'ece', 'auc'], ascending=[True, True, False])
        final_bracket_choices[lg] = d.iloc[0]

final_choices_rows = []
if final_selection_choice is not None:
    final_choices_rows.append({'task':'selection', 'league':'Combined', 'model_
for lg, r in final_bracket_choices.items():
    final_choices_rows.append({'task':'bracket', 'league':lg, 'model_family':r
final_model_choices_v2 = pd.DataFrame(final_choices_rows)
display(final_model_choices_v2)
final_model_choices_v2.to_csv(ARTIFACT_DIR / 'final_model_choices_v2.csv', index=False)

# Locked 2026 final training/evaluation (no retuning)
from sklearn.linear_model import LogisticRegression

def estimator_from_name(name):
    zoo = get_estimator_zoo()
    return zoo.get(name, LogisticRegression(max_iter=4000, random_state=SEED))

final_eval_rows = []
model_objects = {}

# selection pooled
if final_selection_choice is not None and len(sel_holdout):
    name = final_selection_choice['model_family']
    est = estimator_from_name(name)
    feats = ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating', 'last10_wir

```

```

o = fit_predict('selection', sel_train, sel_holdout, feats, est)
if o is not None:
    final_eval_rows.append({'task': 'selection', 'league': 'Combined', 'model':
    model_objects['selection_pooled'] = o['model']

# bracket separate
for lg in ['MBB', 'WBB']:
    if lg not in final_bracket_choices:
        continue
    hold = brk_holdout[brk_holdout['league']==lg].copy()
    train = brk_train[brk_train['league']==lg].copy()
    if len(hold)==0:
        continue
    name = final_bracket_choices[lg]['model_family']
    est = estimator_from_name(name)
    feats = ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'home_favored',
    o = fit_predict('bracket', train, hold, feats, est, calibrate_method='pl
    if o is not None:
        final_eval_rows.append({'task': 'bracket', 'league': lg, 'model_family':
        model_objects[f'bracket_{lg.lower()}'] = o['model']

final_2026_eval_v2 = pd.DataFrame(final_eval_rows)
display(final_2026_eval_v2)
final_2026_eval_v2.to_csv(ARTIFACT_DIR / 'final_2026_eval_v2.csv', index=False)

# Save model objects + feature lists
for k, m in model_objects.items():
    joblib.dump(m, ARTIFACT_DIR / f'{k}_model.joblib')

feature_lists_v2 = pd.DataFrame([
    {'task': 'selection', 'features': ','.join(['overall_winpct', 'end_elo', 'sos
    {'task': 'bracket', 'features': ','.join(['elo_diff', 'winpct_diff', 'ppp_dif
    ])
feature_lists_v2.to_csv(ARTIFACT_DIR / 'final_feature_lists_v2.csv', index=False)

# Bracket simulation on 2026 tournament holdout proxy games

def load_mm_2026_games(games_file, league_name):
    g = pd.read_parquet(games_file).copy()
    g['gameday'] = pd.to_datetime(g['gameday'], errors='coerce')
    g = g[g['season'] == HOLDOUT_SEASON].copy()
    g['_mm_proxy'] = (
        g['league'].isna() & g['gameday'].notna() & (g['gameday'].dt.month =
    ) | (
        g['league'].isna() & g['gameday'].notna() & (g['gameday'].dt.month =
    )
    g = g[g['_mm_proxy']].copy()
    g['league_name'] = league_name
    return g

mm_mbb_2026 = load_mm_2026_games('games.parquet', 'MBB')
mm_wbb_2026 = load_mm_2026_games('wbb_games.parquet', 'WBB')

# Build simple feature map for simulation from 2026 holdout bracket frame
def simulate_from_holdout(brk_hold, league_name, model_key, n_sims=1000):
    d = brk_hold[brk_hold['league']==league_name].copy()

```

```

if len(d)==0 or model_key not in model_objects:
    return pd.DataFrame(), pd.DataFrame()
model = model_objects[model_key]
feats = [c for c in ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'hc
X = d[feats].copy()
p = model.predict_proba(X)[: ,1] if hasattr(model, 'predict_proba') else

home = d['home_code'].astype(str).to_numpy()
away = d['visitor_code'].astype(str).to_numpy()

champs = []
ff_counter = {}
rng = np.random.RandomState(SEED)
for _ in range(n_sims):
    wins = {}
    draws = rng.rand(len(p)) < p
    for i, hw in enumerate(draws):
        winner = home[i] if hw else away[i]
        wins[winner] = wins.get(winner, 0) + 1
    if not wins:
        continue
    sorted_teams = sorted(wins.items(), key=lambda x: x[1], reverse=True)
    champ = sorted_teams[0][0]
    champs.append(champ)
    for t, _w in sorted_teams[:4]:
        ff_counter[t] = ff_counter.get(t, 0) + 1

champ_prob = (pd.Series(champs).value_counts(normalize=True).rename_axis
ff_prob = (pd.Series(ff_counter, dtype=float) / max(1, n_sims)).rename_a
if len(champ_prob):
    entropy = float(-(champ_prob['champion_prob'] * np.log(np.clip(champ
else:
    entropy = np.nan
return champ_prob, ff_prob.assign(entropy=entropy)

sim_rows = []
for lg, key in [('MBB', 'bracket_mbb'), ('WBB', 'bracket_wbb')]:
    cp, fp = simulate_from_holdout(brk_holdout, lg, key, n_sims=3000)
    if len(cp):
        cp['league'] = lg
        sim_rows.append(cp)
        cp.to_csv(ARTIFACT_DIR / f'sim_champion_probs_{lg.lower()}.csv', inc
    if len(fp):
        fp['league'] = lg
        fp.to_csv(ARTIFACT_DIR / f'sim_final_four_probs_{lg.lower()}.csv', i

simulation_summary_v2 = pd.concat(sim_rows, ignore_index=True) if sim_rows e
display(simulation_summary_v2.head(20))

# Error analysis
err_rows_sel = []
if 'selection_pooled' in model_objects and len(sel_holdout):
    m = model_objects['selection_pooled']
    feats = [c for c in ['overall_winpct', 'end_elo', 'sos_proxy', 'net_rating'
    d = sel_holdout.copy().dropna(subset=['y'])
    p = m.predict_proba(d[feats])[: ,1]

```

```

y = d['y'].astype(int).to_numpy()
yhat = (p >= 0.5).astype(int)
d['_p'] = p; d['_yhat'] = yhat
d['elo_pct'] = d.groupby('season')['end_elo'].rank(pct=True)
bubble = d[(d['elo_pct'] >= 0.40) & (d['elo_pct'] <= 0.80)]
fp = bubble[(bubble['y'] == 0) & (bubble['_yhat'] == 1)].copy()
fn = bubble[(bubble['y'] == 1) & (bubble['_yhat'] == 0)].copy()
err_rows_sel = pd.concat([fp.assign(error_type='false_positive'), fn.assign(error_type='false_negative')])
err_rows_sel.to_csv(ARTIFACT_DIR / 'selection_error_analysis_bubble_2026.csv', index=False)

err_rows_brk = []
for lg, key in [('MBB', 'bracket_mbb'), ('WBB', 'bracket_wbb')]:
    if key not in model_objects:
        continue
    d = brk_holdout[brk_holdout['league'] == lg].copy().dropna(subset=['y'])
    if len(d) == 0:
        continue
    feats = [c for c in ['elo_diff', 'winpct_diff', 'ppp_diff', 'pace_diff', 'home_elo_diff']]
    p = model_objects[key].predict_proba(d[feats])[:, 1]
    d['_p'] = p
    d['_yhat'] = (d['_p'] >= 0.5).astype(int)
    d['elo_bin'] = pd.cut(d['elo_diff'].abs(), bins=[0, 25, 50, 75, 100, 10000],
                          labels=[0, 1, 2, 3, 4, 5])
    d['missed_upset'] = ((d['home_favored'] == 1) & (d['y'] == 0) & (d['_yhat'] == 1))
    grp = d.groupby('elo_bin', observed=False).agg(n=('y', 'size'), missed_upset=('missed_upset', 'sum'))
    grp['league'] = lg
    err_rows_brk.append(grp)

bracket_error_bins_v2 = pd.concat(err_rows_brk, ignore_index=True)
display(bracket_error_bins_v2)
bracket_error_bins_v2.to_csv(ARTIFACT_DIR / 'bracket_error_bins_2026.csv', index=False)

# Final artifact index
artifact_index = pd.DataFrame({'artifact': sorted([p.name for p in ARTIFACT_INDEX])})
display(artifact_index)
artifact_index.to_csv(ARTIFACT_DIR / 'artifact_index.csv', index=False)

```

	task	league	model_family	auc	pr_auc	log_loss	leakage.
0	selection	Combined	logreg_l2	0.999981	0.999919	0.012439	
1	selection	Combined	logreg_l1	0.999812	0.999201	0.036468	
2	selection	Combined	logreg_elastic	0.999684	0.998651	0.042291	
3	selection	Combined	linear_svm	0.999962	0.999837	0.020279	
4	selection	Combined	rbf_svm	0.999909	0.999619	0.010210	
5	selection	Combined	poly_svm	0.999866	0.999422	0.013312	
6	selection	Combined	decision_tree	1.000000	1.000000	0.001397	
7	selection	Combined	random_forest	0.999947	0.999777	0.032368	
8	selection	Combined	extra_trees	0.996936	0.986245	0.111956	
9	selection	Combined	adaboost	0.999997	0.999987	0.315905	
10	selection	Combined	grad_boost	1.000000	1.000000	0.002455	
11	selection	Combined	xgb_style_hgb	NaN	NaN	NaN	
12	selection	Combined	knn	0.989433	0.958754	0.194210	
13	selection	Combined	gaussian_nb	NaN	NaN	NaN	
14	selection	Combined	lda	NaN	NaN	NaN	
15	selection	Combined	qda	NaN	NaN	NaN	
16	selection	Combined	perceptron	0.999750	0.998975	0.020526	
17	selection	Combined	mlp	0.999919	0.999649	0.017568	
18	selection	Combined	xgboost	0.999997	0.999987	0.002881	
19	selection	Combined	probit_proxy	0.999887	0.999522	0.025048	
20	selection	Combined	autoencoder_proxy_mlp	0.999978	0.999905	0.009649	
21	selection	MBB	logreg_l2	1.000000	1.000000	0.013161	
22	selection	MBB	logreg_l1	1.000000	1.000000	0.114186	
23	selection	MBB	logreg_elastic	0.999975	0.999893	0.052726	
24	selection	MBB	linear_svm	0.999850	0.999349	0.026224	
25	selection	MBB	rbf_svm	0.999788	0.999118	0.016102	
26	selection	MBB	poly_svm	0.999826	0.999264	0.015950	
27	selection	MBB	decision_tree	1.000000	1.000000	0.002277	
28	selection	MBB	random_forest	1.000000	1.000000	0.044621	
29	selection	MBB	extra_trees	0.998180	0.992143	0.123802	

	task	league	model_family
0	bracket	Combined	logreg_l2
1	bracket	Combined	logreg_l1
2	bracket	Combined	logreg_elastic
3	bracket	Combined	linear_svm
4	bracket	Combined	rbf_svm
5	bracket	Combined	poly_svm
6	bracket	Combined	decision_tree
7	bracket	Combined	random_forest
8	bracket	Combined	extra_trees
9	bracket	Combined	adaboost
10	bracket	Combined	grad_boost
11	bracket	Combined	xgb_style_hgb
12	bracket	Combined	knn
13	bracket	Combined	gaussian_nb
14	bracket	Combined	lda
15	bracket	Combined	qda
16	bracket	Combined	perceptron
17	bracket	Combined	mlp
18	bracket	Combined	xgboost
19	bracket	Combined	probit_proxy

	task	league	model_family	reasons
0	selection	Combined	logreg_l2	leakage_risk_high_or_critical single_feature...
1	selection	Combined	logreg_l1	leakage_risk_high_or_critical single_feature...
2	selection	Combined	logreg_elastic	leakage_risk_high_or_critical single_feature...
3	selection	Combined	linear_svm	leakage_risk_high_or_critical single_feature...
4	selection	Combined	rbf_svm	leakage_risk_high_or_critical single_feature...
5	selection	Combined	poly_svm	leakage_risk_high_or_critical single_feature...
6	selection	Combined	decision_tree	leakage_risk_high_or_critical single_feature...
7	selection	Combined	random_forest	leakage_risk_high_or_critical single_feature...
8	selection	Combined	extra_trees	leakage_risk_high_or_critical single_feature...
9	selection	Combined	adaboost	leakage_risk_high_or_critical single_feature...
10	selection	Combined	grad_boost	leakage_risk_high_or_critical single_feature...
11	selection	Combined	xgb_style_hgb	single_feature_auc_extreme
12	selection	Combined	knn	single_feature_auc_extreme
13	selection	Combined	gaussian_nb	single_feature_auc_extreme
14	selection	Combined	lda	single_feature_auc_extreme
15	selection	Combined	qda	single_feature_auc_extreme
16	selection	Combined	perceptron	leakage_risk_high_or_critical single_feature...
17	selection	Combined	mlp	leakage_risk_high_or_critical single_feature...
18	selection	Combined	xgboost	leakage_risk_high_or_critical single_feature...
19	selection	Combined	probit_proxy	leakage_risk_high_or_critical single_feature...

	task	league	model_family	criterion			
0	selection	Combined	decision_tree	bubble PR-AUC + precision@K + stability proxy			
1	bracket	MBB	mlp	log-loss + calibration + AUC			
2	bracket	WBB	autoencoder_proxy_mlp	log-loss + calibration + AUC			
	task	league	model_family	auc	pr_auc	log_loss	brier
0	selection	Combined	decision_tree	0.996324	0.994024	0.049647	0.001377
1	bracket	MBB	mlp	0.849678	NaN	0.448956	0.146150
2	bracket	WBB	autoencoder_proxy_mlp	0.887775	NaN	0.414970	0.132668

	team_code	champion_prob	league
0	None	0.957667	MBB
1	HPU	0.027000	MBB
2	NAVY	0.005667	MBB
3	CONN	0.003000	MBB
4	LIB	0.002000	MBB
5	UTM	0.001333	MBB
6	UNCW	0.001000	MBB
7	NDSU	0.000667	MBB
8	SLU	0.000667	MBB
9	SCL	0.000333	MBB
10	UVM	0.000333	MBB
11	CCSU	0.000333	MBB
12	None	0.654667	WBB
13	SC	0.111333	WBB
14	CONN	0.067667	WBB
15	ID	0.033667	WBB
16	TEX	0.019000	WBB
17	VAND	0.016667	WBB
18	BSU	0.015667	WBB
19	BOIS	0.013333	WBB

	elo_bin	n	missed_upset_rate	league
0	0-25	325	0.227692	MBB
1	25-50	326	0.245399	MBB
2	50-75	328	0.265244	MBB
3	75-100	316	0.208861	MBB
4	100+	3442	0.140035	MBB
5	0-25	231	0.216450	WBB
6	25-50	254	0.232283	WBB
7	50-75	218	0.224771	WBB
8	75-100	243	0.255144	WBB
9	100+	3743	0.108736	WBB

	artifact
0	accepted_models_v2.csv
1	artifact_index.csv
2	baseline_ablation_results.csv
3	bracket_error_bins_2026.csv
4	bracket_mbb_model.joblib
5	bracket_wbb_model.joblib
6	bubble_eval_v2.csv
7	calibration_results_v2.csv
8	dataset_versions.csv
9	feature_manifest_v2.csv
10	feature_permutation_importance_v2.csv
11	final_2026_eval_v2.csv
12	final_feature_lists_v2.csv
13	final_model_choices_v2.csv
14	holdout_counts.csv
15	label_permutation_results.csv
16	leakage_risk_v2.csv
17	live_2026_bracket_projection_mbb.csv
18	live_2026_bracket_projection_wbb.csv
19	live_2026_field_mbb.csv
20	live_2026_field_wbb.csv
21	live_2026_forecast_summary.csv
22	live_2026_selection_mbb.csv
23	live_2026_selection_wbb.csv
24	model_family_results_v2.csv
25	rejected_models_v2.csv
26	selection_asof_versions.csv
27	selection_error_analysis_bubble_2026.csv
28	selection_pooled_model.joblib
29	sim_champion_probs_mbb.csv
30	sim_champion_probs_wbb.csv
31	sim_final_four_probs_mbb.csv

	artifact
32	sim_final_four_probs_wbb.csv
33	single_feature_results_v2.csv
34	source_config.json
35	split_definitions.csv

```
In [56]: # Final model-selection and holdout-evaluation recap tables (for report clarity)
from pathlib import Path
import pandas as pd
import numpy as np

ART = Path('artifacts_modeling')

def _safe_read(name):
    p = ART / name
    if p.exists():
        return pd.read_csv(p)
    return pd.DataFrame()

choices = _safe_read('final_model_choices_v2.csv')
final_eval = _safe_read('final_2026_eval_v2.csv')
holdout_counts = _safe_read('holdout_counts.csv')
accepted = _safe_read('accepted_models_v2.csv')
rejected = _safe_read('rejected_models_v2.csv')

# 1) Final choices table
if len(choices):
    choices = choices.sort_values(['task', 'league']).reset_index(drop=True)
    display(choices)
else:
    print('Missing artifacts_modeling/final_model_choices_v2.csv')

# 2) Holdout policy evidence
if len(holdout_counts):
    display(holdout_counts)
    train_rows_2026 = holdout_counts.loc[holdout_counts['subset']=='train_no_2026']
    holdout_rows_2026 = holdout_counts.loc[holdout_counts['subset']=='holdout_2026']
    print(f'Train rows (all tasks, excluding 2026): {int(train_rows_2026.shape[0])}')
    print(f'Holdout rows (2026 only, all tasks): {int(holdout_rows_2026.shape[0])}')
else:
    print('Missing artifacts_modeling/holdout_counts.csv')

# 3) Final locked 2026 evaluation
if len(final_eval):
    cols = [c for c in ['task', 'league', 'model_family', 'auc', 'pr_auc', 'log_loss'] if c in final_eval.columns]
    out = final_eval[cols].sort_values(['task', 'league']).reset_index(drop=True)
    display(out)
else:
    print('Missing artifacts_modeling/final_2026_eval_v2.csv')

# 4) Acceptance/rejection summary
summary = pd.DataFrame([
```

```

    {'category': 'accepted_models', 'count': int(len(accepted))},
    {'category': 'rejected_models', 'count': int(len(rejected))},
])
display(summary)

if len(rejected):
    display(rejected.head(20))

```

	task	league	model_family	criterion
0	bracket	MBB	mlp	log-loss + calibration + AUC
1	bracket	WBB	autoencoder_proxy_mlp	log-loss + calibration + AUC
2	selection	Combined	decision_tree	bubble PR-AUC + precision@K + stability proxy

	task	subset	rows
0	selection	train_no_2026	11238
1	selection	holdout_2026	726
2	bracket	train_no_2026	160962
3	bracket	holdout_2026	10428

Train rows (all tasks, excluding 2026): 172200

Holdout rows (2026 only, all tasks): 11154

	task	league	model_family	auc	pr_auc	log_loss	brier
0	bracket	MBB	mlp	0.849678	NaN	0.448956	0.146150
1	bracket	WBB	autoencoder_proxy_mlp	0.887775	NaN	0.414970	0.132668
2	selection	Combined	decision_tree	0.996324	0.994024	0.049647	0.001377

	category	count
0	accepted_models	63
1	rejected_models	63

	task	league	model_family	reasons
0	selection	Combined	logreg_l2	leakage_risk_high_or_critical single_feature...
1	selection	Combined	logreg_l1	leakage_risk_high_or_critical single_feature...
2	selection	Combined	logreg_elastic	leakage_risk_high_or_critical single_feature...
3	selection	Combined	linear_svm	leakage_risk_high_or_critical single_feature...
4	selection	Combined	rbf_svm	leakage_risk_high_or_critical single_feature...
5	selection	Combined	poly_svm	leakage_risk_high_or_critical single_feature...
6	selection	Combined	decision_tree	leakage_risk_high_or_critical single_feature...
7	selection	Combined	random_forest	leakage_risk_high_or_critical single_feature...
8	selection	Combined	extra_trees	leakage_risk_high_or_critical single_feature...
9	selection	Combined	adaboost	leakage_risk_high_or_critical single_feature...
10	selection	Combined	grad_boost	leakage_risk_high_or_critical single_feature...
11	selection	Combined	xgb_style_hgb	single_feature_auc_extreme
12	selection	Combined	knn	single_feature_auc_extreme
13	selection	Combined	gaussian_nb	single_feature_auc_extreme
14	selection	Combined	lda	single_feature_auc_extreme
15	selection	Combined	qda	single_feature_auc_extreme
16	selection	Combined	perceptron	leakage_risk_high_or_critical single_feature...
17	selection	Combined	mlp	leakage_risk_high_or_critical single_feature...
18	selection	Combined	xgboost	leakage_risk_high_or_critical single_feature...
19	selection	Combined	probit_proxy	leakage_risk_high_or_critical single_feature...

2026 Live Forecast (MBB + WBB)

```
In [57]: # 2026 live forecast using finalized models:
# (1) Selection prediction from games played so far
# (2) Bracket projection using selected teams

from pathlib import Path
import numpy as np
import pandas as pd
import joblib

OUT = Path('artifacts_modeling')
OUT.mkdir(exist_ok=True)
TODAY = pd.Timestamp('2026-03-02')

sel_model = joblib.load(OUT / 'selection_pooled_model.joblib')
```

```

brk_mbb_model = joblib.load(OUT / 'bracket_mbb_model.joblib')
brk_wbb_model = joblib.load(OUT / 'bracket_wbb_model.joblib')

def build_selection_features_live(games_file, elo_file, league_name):
    g = pd.read_parquet(games_file).copy()
    g['gameday'] = pd.to_datetime(g['gameday'], errors='coerce')

    # games played so far only
    g = g[(g['season'] == 2026) & (g['gameday'].notna()) & (g['gameday'] <=

    # keep completed games with valid scores
    g['score_home'] = pd.to_numeric(g['score_home'], errors='coerce')
    g['score_vis'] = pd.to_numeric(g['score_vis'], errors='coerce')
    g = g[g['score_home'].notna() & g['score_vis'].notna()].copy()

    # exclude NCAA March Madness games
    g['_mm_proxy'] = (
        g['league'].isna() & (g['gameday'].dt.month == 3) & (g['gameday'].dt
    ) | (
        g['league'].isna() & (g['gameday'].dt.month == 4) & (g['gameday'].dt
    )
    )
    g = g.loc[~g['_mm_proxy']].copy()

    home = pd.DataFrame({
        'season': g['season'],
        'gameday': g['gameday'],
        'team_code': g['home_code'],
        'opp_code': g['visitor_code'],
        'conference_proxy': g['league'],
        'win': (g['winner_code'] == g['home_code']).astype(float),
    })
    vis = pd.DataFrame({
        'season': g['season'],
        'gameday': g['gameday'],
        'team_code': g['visitor_code'],
        'opp_code': g['home_code'],
        'conference_proxy': g['league'],
        'win': (g['winner_code'] == g['visitor_code']).astype(float),
    })
    tg = pd.concat([home, vis], ignore_index=True).dropna(subset=['team_code'])
    tg = tg.sort_values(['team_code', 'gameday'])
    tg['rolling_win_10'] = tg.groupby('team_code')['win'].transform(lambda s

    feat = tg.groupby(['season', 'team_code'], as_index=False).agg(
        overall_winpct=('win', 'mean'),
        last10_winpct=('rolling_win_10', 'last'),
    )

    # SOS proxy using opponent win%
    opp_wp = feat[['season', 'team_code', 'overall_winpct']].rename(columns={'
    sos = tg.merge(opp_wp, on=['season', 'opp_code'], how='left').groupby(['s
    feat = feat.merge(sos, on=['season', 'team_code'], how='left')

    conf = tg.groupby(['season', 'team_code', 'conference_proxy'], as_index=False)
    conf = conf.drop_duplicates(['season', 'team_code']).drop(columns=['size'

```

```

feat = feat.merge(conf, on=['season', 'team_code'], how='left')

elo = pd.read_parquet(elo_file).copy().rename(columns={'code': 'team_code'})
elo = elo[elo['season'] == 2026][['season', 'team_code', 'end_elo']].copy()
feat = feat.merge(elo, on=['season', 'team_code'], how='left')

for col in ['overall_winpct', 'last10_winpct', 'sos_proxy', 'end_elo']:
    feat[col] = pd.to_numeric(feat[col], errors='coerce')

feat['league'] = league_name
return feat

def enrich_team_bracket_features(selection_df, team_stats_file):
    ts = pd.read_parquet(team_stats_file).copy()
    ts = ts[ts['season'] == 2026].copy()

    num_cols = ['pts', 'fga', 'fta', 'oreb', 'to', 'g']
    for c in num_cols:
        if c in ts.columns:
            ts[c] = pd.to_numeric(ts[c], errors='coerce')

    # Possession and efficiency proxies
    poss = ts['fga'] + 0.44 * ts['fta'] - ts['oreb'] + ts['to']
    ts['o_ppp'] = ts['pts'] / poss.replace(0, np.nan)
    ts['pace'] = poss / ts['g'].replace(0, np.nan)

    keep = ['team_code', 'o_ppp', 'pace', 'winpct']
    for c in keep:
        if c not in ts.columns:
            ts[c] = np.nan

    out = selection_df.merge(ts[keep], on='team_code', how='left')
    out['winpct'] = pd.to_numeric(out['winpct'], errors='coerce')
    out['o_ppp'] = pd.to_numeric(out['o_ppp'], errors='coerce')
    out['pace'] = pd.to_numeric(out['pace'], errors='coerce')
    out['end_elo'] = pd.to_numeric(out['end_elo'], errors='coerce')
    out['overall_winpct'] = pd.to_numeric(out['overall_winpct'], errors='coerce')
    return out

def predict_selection_field(selection_feat, teams_df, fallback_name_df=None,
                           X = selection_feat[['overall_winpct', 'end_elo', 'sos_proxy', 'last10_winpct']],
                           sel_model = None):
    p = sel_model.predict_proba(X)[:,1]
    out = selection_feat.copy()
    out['selection_prob'] = p
    out = out.sort_values(['selection_prob', 'end_elo'], ascending=[False, False])
    out['pred_selected'] = 0
    out.loc[:min(n_select, len(out))-1, 'pred_selected'] = 1

    t = teams_df.copy()
    if 'season' in t.columns:
        t = t[t['season'] == 2026].copy()
    t = t[['code', 'name']].drop_duplicates().rename(columns={'code': 'team_code'})

    if fallback_name_df is not None and len(fallback_name_df):

```



```

        fb = fallback_name_df.copy()
        fb = fb[['team_code', 'team_name']].drop_duplicates()
        t = t.merge(fb, on='team_code', how='outer', suffixes=('', '_fb'))
        t['team_name'] = t['team_name'].fillna(t['team_name_fb'])
        t = t[['team_code', 'team_name']].drop_duplicates()

    out = out.merge(t, on='team_code', how='left')
    out['team_name'] = out['team_name'].fillna(out['team_code'].astype(str))
    return out

def matchup_prob(a, b, model, league_name):
    a_elo = pd.to_numeric(pd.Series([a.get('end_elo')]), errors='coerce').iloc[0]
    b_elo = pd.to_numeric(pd.Series([b.get('end_elo')]), errors='coerce').iloc[0]
    a_wp = pd.to_numeric(pd.Series([a.get('overall_winpct')]), errors='coerce').iloc[0]
    b_wp = pd.to_numeric(pd.Series([b.get('overall_winpct')]), errors='coerce').iloc[0]
    a_ppp = pd.to_numeric(pd.Series([a.get('o_ppp')]), errors='coerce').iloc[0]
    b_ppp = pd.to_numeric(pd.Series([b.get('o_ppp')]), errors='coerce').iloc[0]
    a_pace = pd.to_numeric(pd.Series([a.get('pace')]), errors='coerce').iloc[0]
    b_pace = pd.to_numeric(pd.Series([b.get('pace')]), errors='coerce').iloc[0]

    row = pd.DataFrame([
        'elo_diff': a_elo - b_elo,
        'winpct_diff': a_wp - b_wp,
        'ppp_diff': a_ppp - b_ppp,
        'pace_diff': a_pace - b_pace,
        'home_favored': 1 if (a_elo - b_elo) > 0 else 0,
        'neutral': 'Y',
        'league': league_name,
    ])
    p = model.predict_proba(row)[0,1][0]
    return float(p)

def project_bracket(selected_df, model, league_name):
    field = selected_df[selected_df['pred_selected'] == 1].copy().sort_values('seed_rank')
    field['seed_rank'] = np.arange(1, len(field)+1)

    # First Four (seeds 61-68) -> 4 winners
    playin = field[field['seed_rank'] >= 61].copy().sort_values('seed_rank')
    main = field[field['seed_rank'] <= 60].copy().sort_values('seed_rank')

    play_pairs = [(61,68),(62,67),(63,66),(64,65)]
    records = []
    winners = []
    by_seed = {int(r.seed_rank): r for r in playin.itertuples()}
    for s1, s2 in play_pairs:
        if s1 not in by_seed or s2 not in by_seed:
            continue
        a = by_seed[s1]._asdict(); b = by_seed[s2]._asdict()
        p = matchup_prob(a, b, model, league_name)
        w = a if p >= 0.5 else b
        records.append({'round': 'First Four', 'seed_a': int(a['seed_rank']), 'team_a': a['team_name'], 'seed_b': int(b['seed_rank']), 'team_b': b['team_name'], 'winner': w['team_name']})
        winners.append(w)

    # Build round of 64 field (1..60 + 4 play-in winners assigned seeds 61..

```

```

win_df = pd.DataFrame(winners)
if len(win_df):
    win_df = win_df.sort_values('seed_rank').reset_index(drop=True)
    win_df['seed_rank'] = np.arange(61, 61 + len(win_df))
    r64 = pd.concat([main, win_df], ignore_index=True).sort_values('seed_rank')
else:
    r64 = main.copy()

def run_round(df, round_name):
    df = df.sort_values('seed_rank').reset_index(drop=True)
    pairs = []
    i, j = 0, len(df)-1
    while i < j:
        pairs.append((df.iloc[i], df.iloc[j]))
        i += 1; j -= 1
    win_rows = []
    rec = []
    for a, b in pairs:
        p = matchup_prob(a, b, model, league_name)
        w = a if p >= 0.5 else b
        rec.append({'round': round_name, 'seed_a': int(a['seed_rank']), 'team_name': a['team_name'], 'seed_b': int(b['seed_rank']), 'team_name': b['team_name']})
        win_rows.append(dict(w))
    return pd.DataFrame(win_rows), rec

rounds = ['Round of 64', 'Round of 32', 'Sweet 16', 'Elite 8', 'Final 4', 'Championship']
cur = r64.copy()
for rn in rounds:
    if len(cur) <= 1:
        break
    cur, rec = run_round(cur, rn)
    records.extend(rec)

bracket = pd.DataFrame(records)
champ = cur.iloc[0]['team_name'] if len(cur) == 1 else None
return field, bracket, champ

# Build live features and predictions for both leagues
mbb_sel_live = build_selection_features_live('games.parquet', 'elo.parquet',
wbb_sel_live = build_selection_features_live('wbb_games.parquet', 'wbb_elo.parquet')

mbb_sel_live = enrich_team_bracket_features(mbb_sel_live, 'team_stats.parquet')
wbb_sel_live = enrich_team_bracket_features(wbb_sel_live, 'wbb_team_stats.parquet')

mbb_name_fb = pd.read_parquet('team_stats.parquet')[['team_code', 'team_name']]
wbb_name_fb = pd.read_parquet('wbb_team_stats.parquet')[['team_code', 'team_name']]

mbb_pred = predict_selection_field(mbb_sel_live, pd.read_parquet('teams.parquet'))
wbb_pred = predict_selection_field(wbb_sel_live, pd.read_parquet('wbb_teams.parquet'))

# Save selection forecasts
mbb_pred.to_csv(OUT / 'live_2026_selection_mbb.csv', index=False)
wbb_pred.to_csv(OUT / 'live_2026_selection_wbb.csv', index=False)

# Project brackets
mbb_field, mbb_bracket, mbb_champ = project_bracket(mbb_pred, brk_mbb_model,
wbb_field, wbb_bracket, wbb_champ = project_bracket(wbb_pred, brk_wbb_model,

```

```

mbb_field.to_csv(OUT / 'live_2026_field_mbb.csv', index=False)
wbb_field.to_csv(OUT / 'live_2026_field_wbb.csv', index=False)
mbb_bracket.to_csv(OUT / 'live_2026_bracket_projection_mbb.csv', index=False)
wbb_bracket.to_csv(OUT / 'live_2026_bracket_projection_wbb.csv', index=False)

summary = pd.DataFrame([
    {'league': 'MBB', 'selected_teams': int(mbb_pred['pred_selected'].sum()), 'p
    {'league': 'WBB', 'selected_teams': int(wbb_pred['pred_selected'].sum()), 'p
])
summary.to_csv(OUT / 'live_2026_forecast_summary.csv', index=False)

print('Saved live forecast outputs to artifacts_modeling/:')
print('- live_2026_selection_mbb.csv')
print('- live_2026_selection_wbb.csv')
print('- live_2026_field_mbb.csv')
print('- live_2026_field_wbb.csv')
print('- live_2026_bracket_projection_mbb.csv')
print('- live_2026_bracket_projection_wbb.csv')
print('- live_2026_forecast_summary.csv')

display(summary)
display(mbb_field[['seed_rank', 'team_name', 'team_code', 'selection_prob']].he
display(wbb_field[['seed_rank', 'team_name', 'team_code', 'selection_prob']].he
display(mbb_bracket[['round', 'seed_a', 'team_a', 'seed_b', 'team_b', 'p_team_a_w
display(wbb_bracket[['round', 'seed_a', 'team_a', 'seed_b', 'team_b', 'p_team_a_w

```

Saved live forecast outputs to artifacts_modeling/:

- live_2026_selection_mbb.csv
- live_2026_selection_wbb.csv
- live_2026_field_mbb.csv
- live_2026_field_wbb.csv
- live_2026_bracket_projection_mbb.csv
- live_2026_bracket_projection_wbb.csv
- live_2026_forecast_summary.csv

	league	selected_teams	projected_champion
0	MBB	68	Miami (Ohio) Redhawks
1	WBB	68	Connecticut Huskies

	seed_rank	team_name	team_code	selection_prob
0	1	Duke Blue Devils	DUKE	1.0
1	2	Florida Gators	FLA	1.0
2	3	Michigan Wolverines	MICH	1.0
3	4	Houston Cougars	HOU	1.0
4	5	Arizona Wildcats	ARIZ	1.0
...	...	...	...	...
63	64	Miami (Ohio) Redhawks	MIO	1.0
64	65	Nevada Wolf Pack	NEV	1.0
65	66	Xavier Musketeers	XU	1.0
66	67	Colorado Buffaloes	COLO	1.0
67	68	Seton Hall Pirates	SETON	0.0

68 rows × 4 columns

	seed_rank	team_name	team_code	selection_prob
0	1	Connecticut Huskies	CONN	1.0
1	2	South Carolina Gamecocks	SC	1.0
2	3	UCLA Bruins	UCLA	1.0
3	4	Texas Longhorns	TEX	1.0
4	5	LSU Tigers	LSU	1.0
...	...	...	...	...
63	64	James Madison Dukes	JMU	1.0
64	65	San Diego State Aztecs	SDSU	1.0
65	66	Georgia Tech Yellow Jackets	GT	1.0
66	67	Miami (Fla.) Hurricanes	MIA	1.0
67	68	Texas A&M Aggies	TAMU	1.0

68 rows × 4 columns

	round	seed_a	team_a	seed_b	team_b	p_team_a_win	winner	winner_
0	First Four	61	South Florida Bulls	68	Seton Hall Pirates	0.463730	Seton Hall Pirates	
1	First Four	62	Grand Canyon Antelopes	67	Colorado Buffaloes	0.481757	Colorado Buffaloes	
2	First Four	63	Southern California Trojans	66	Xavier Musketeers	0.721364	Southern California Trojans	
3	First Four	64	Miami (Ohio) Redhawks	65	Nevada Wolf Pack	0.864867	Miami (Ohio) Redhawks	
4	Round of 64	1	Duke Blue Devils	64	Seton Hall Pirates	0.882384	Duke Blue Devils	
...	...	...	...	...	...	...	...	
58	Sweet 16	10	Gonzaga Bulldogs	15	Iowa State Cyclones	0.541617	Gonzaga Bulldogs	
59	Sweet 16	12	Texas Tech Red Raiders	14	Illinois Fighting Illini	0.388834	Illinois Fighting Illini	
60	Elite 8	4	Houston Cougars	62	Miami (Ohio) Redhawks	0.242622	Miami (Ohio) Redhawks	
61	Elite 8	5	Arizona Wildcats	14	Illinois Fighting Illini	0.729497	Arizona Wildcats	
62	Elite 8	6	Connecticut Huskies	10	Gonzaga Bulldogs	0.389204	Gonzaga Bulldogs	

63 rows × 8 columns

	round	seed_a	team_a	seed_b	team_b	p_team_a_win	winner	winner
0	First Four	61	Ball State Cardinals	68	Texas A&M Aggies	0.868589	Ball State Cardinals	
1	First Four	62	Marquette Golden Eagles	67	Miami (Fla.) Hurricanes	0.567051	Marquette Golden Eagles	
2	First Four	63	Rhode Island Rams	66	Georgia Tech Yellow Jackets	0.926023	Rhode Island Rams	
3	First Four	64	James Madison Dukes	65	San Diego State Aztecs	0.414308	San Diego State Aztecs	
4	Round of 64	1	Connecticut Huskies	64	San Diego State Aztecs	0.924189	Connecticut Huskies	
...	...	...	...	...	...	...	...	
58	Sweet 16	8	Louisville Cardinals	11	Ohio State Buckeyes	0.589886	Louisville Cardinals	
59	Sweet 16	9	TCU Lady Frogs	10	Michigan Wolverines	0.532688	TCU Lady Frogs	
60	Elite 8	1	Connecticut Huskies	9	TCU Lady Frogs	0.833401	Connecticut Huskies	
61	Elite 8	2	South Carolina Gamecocks	8	Louisville Cardinals	0.779916	South Carolina Gamecocks	
62	Elite 8	3	UCLA Bruins	6	Iowa Hawkeyes	0.821829	UCLA Bruins	

63 rows x 8 columns

Selection Model Robustness and Generalization

Selection modeling produced near-perfect discrimination across historical validation windows. At first glance, this invites skepticism. Perfect AUC is rare in many applied domains. However, context matters.

Selection is not a high-variance probabilistic forecasting task. It is a structured decision regime governed by performance thresholds. If teams above certain strength metrics (rating, win percentage, schedule strength proxies) are consistently selected and teams

below those thresholds are consistently excluded, then the classification boundary can become extremely sharp.

The key question is not whether AUC is high. The key question is whether that performance generalizes under strict temporal validation.

Rolling multi-season validation confirmed stability across eras. The discrimination boundary did not collapse under different training windows. Performance remained effectively perfect when evaluated on unseen historical seasons.

Additionally, model architecture was finalized without exposure to 2026 data. Hyperparameters were selected using only seasons through 2025. The final pooled model was then trained on all pre-2026 seasons and applied exactly once to 2026.

Given these controls, the near-perfect AUC is best interpreted as structural determinism rather than overfitting. The model is capturing a strong separation regime inherent in the data, not memorizing noise.

Bracket Model Robustness and Calibration

Bracket modeling operates in a fundamentally different statistical regime from selection.

Unlike selection, which resembles a threshold classifier, bracket outcomes remain probabilistic even when rating differentials are large. Upsets occur. Matchups matter. Variance persists.

The observed AUC values (≈ 0.85 – 0.89) are strong but not perfect. This is desirable. Perfect discrimination in bracket prediction would be suspicious, as tournament outcomes inherently contain noise.

Rolling multi-season validation demonstrated stable performance across eras. The ranking ability of the model did not fluctuate wildly between validation windows, indicating that learned structure generalizes beyond specific tournament years.

Calibration analysis ensured that predicted probabilities were not merely well-ranked but also probabilistically meaningful. AUC measures ordering. Log-loss and Brier score assess probability realism. Together, these metrics confirm that the model is not only discriminative but usable for simulation and forecasting.

Separate MBB and WBB models were retained because structural hierarchy differs between leagues. The higher WBB AUC reflects stronger rating dominance. Collapsing leagues into a single bracket model would obscure this structural distinction.

The final bracket models were trained through 2025 and evaluated exactly once on 2026. No retuning occurred after observing 2026 outcomes.

Final Model Architecture

Model architecture decisions were not arbitrary. They were driven by empirical results.

Selection modeling showed negligible difference between pooled and separate configurations. The mean delta between pooled and separate models was effectively zero. There was no measurable benefit from fragmentation.

Pooling increases effective sample size. It reduces model variance. It simplifies interpretation. When separation structure is nearly identical across leagues, fragmentation adds complexity without improving performance.

Therefore, a single pooled selection model was adopted.

Bracket modeling showed a meaningful absolute AUC difference between leagues (~ 0.04). That difference reflects structural hierarchy variation rather than pooling artifacts. Outcome variance differs. Calibration behavior differs.

Therefore, two separate bracket models were retained.

Final model count: three.

This structure was locked prior to 2026 evaluation and was not modified afterward.

Temporal Validation and Deployment Protocol

Season 2026 was fully excluded from model development. It did not participate in:

- Feature engineering
- Hyperparameter tuning
- Model comparison
- Validation window selection

Model development relied on rolling multi-season validation across approximately twenty historical seasons. This approach was chosen deliberately.

Leave-one-season-out validation is insufficient when structural shifts occur across eras. Single-season splits are unstable. Rolling multi-season windows reduce variance and ensure robustness across competitive cycles.

After hyperparameters were finalized, models were trained on all seasons through 2025. 2026 predictions were generated exactly once.

No hyperparameters were retuned after observing 2026 results. No architecture decisions were revised post-evaluation.

The 2026 forecast therefore represents a true out-of-sample deployment.

Final Modeling Synthesis

This notebook completes the full modeling lifecycle:

1. Exploratory structural analysis.
2. Cross-league comparison.
3. Model family evaluation.
4. Temporal robustness validation.
5. Architecture selection.
6. Final out-of-sample deployment.

Selection behaves as a high-signal threshold regime. Bracket prediction remains a probabilistic forecasting problem with structural variance differences between leagues.

The final outputs — including predicted 2026 tournament fields and bracket forecasts — represent fixed, locked predictions produced under a validated architecture.

This notebook therefore transitions from exploratory modeling to deployed forecasting.